

Software Design Basics

Ravi Suppiah
Lecturer, NUS SoC

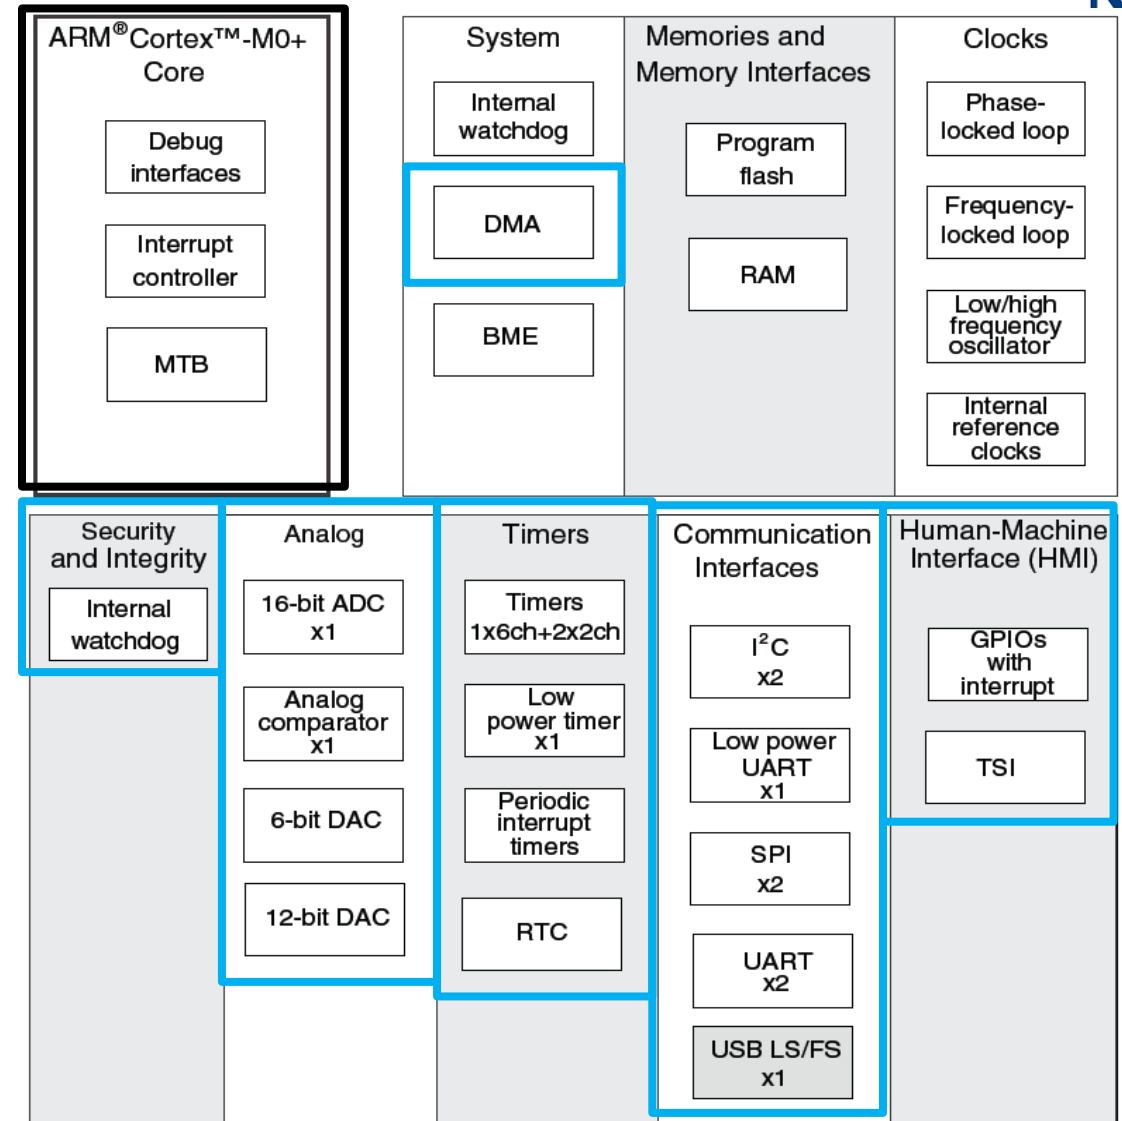
Overview

- Concurrency
 - How do we make things happen at the right time?
- Software Engineering for Embedded Systems
 - How do we develop working code quickly?

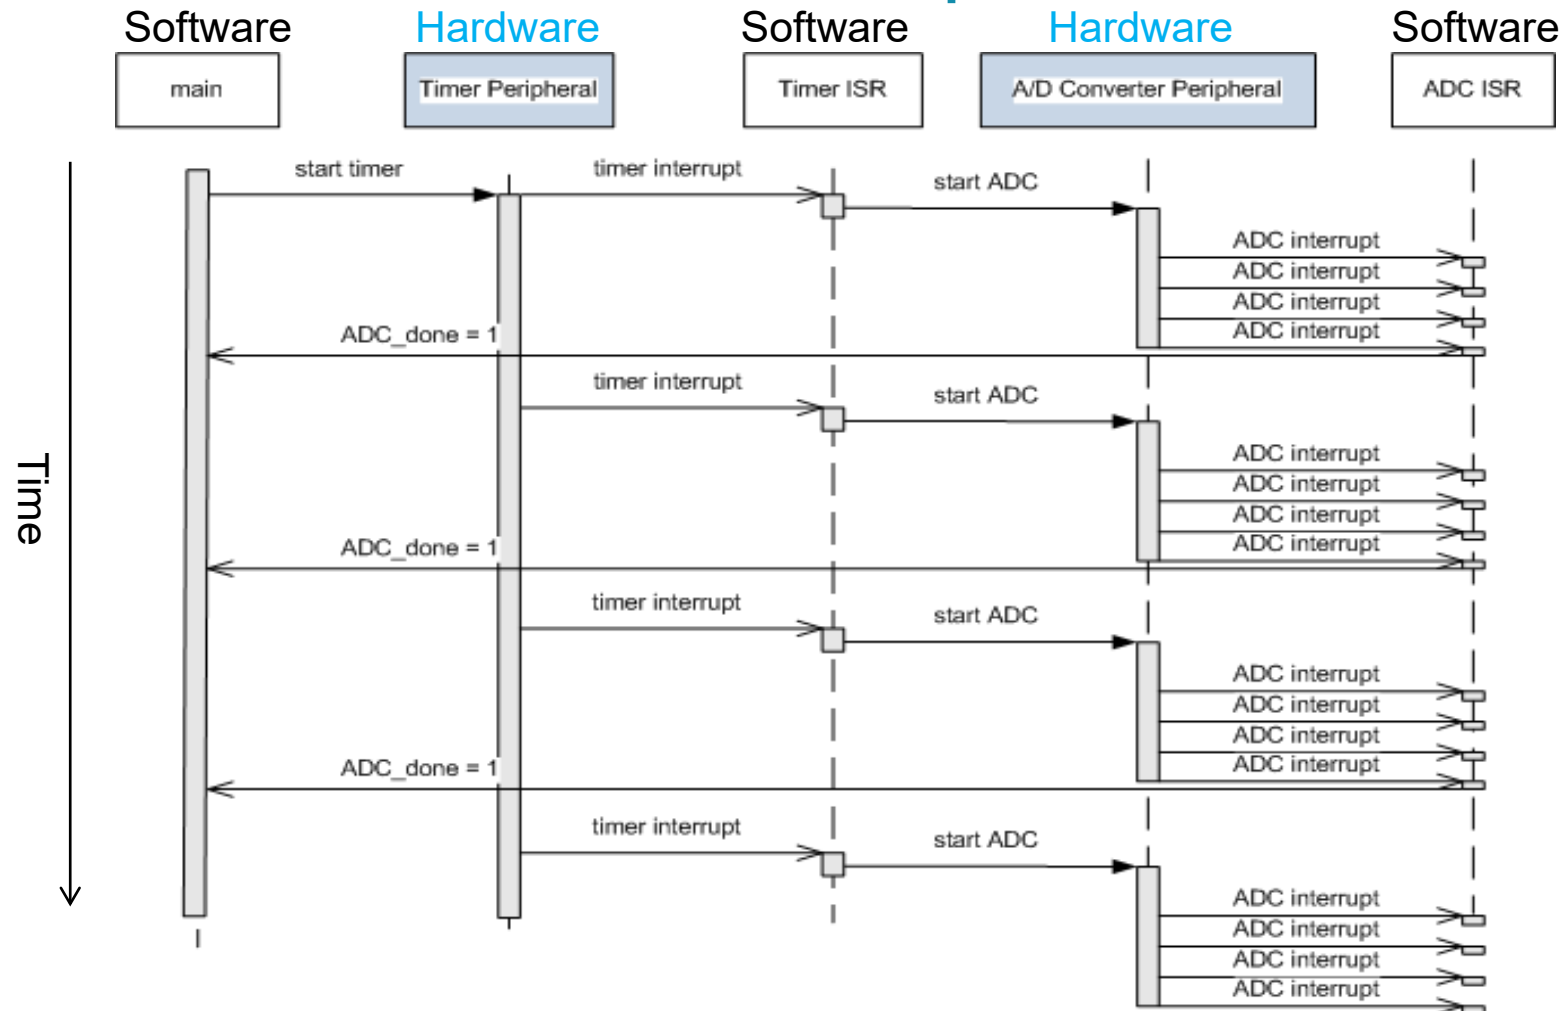
CONCURRENCY

MCU Hardware & Software for Concurrency

- CPU executes instructions from one or more threads of execution
- Specialized hardware peripherals add dedicated concurrent processing
 - DMA - transferring data between memory and peripherals
 - Watchdog timer
 - Analog interfacing
 - Timers
 - Communications with other devices
 - Detecting external signal events
- Peripherals use ***interrupts*** to notify CPU of events



Concurrent Hardware & Software Operation

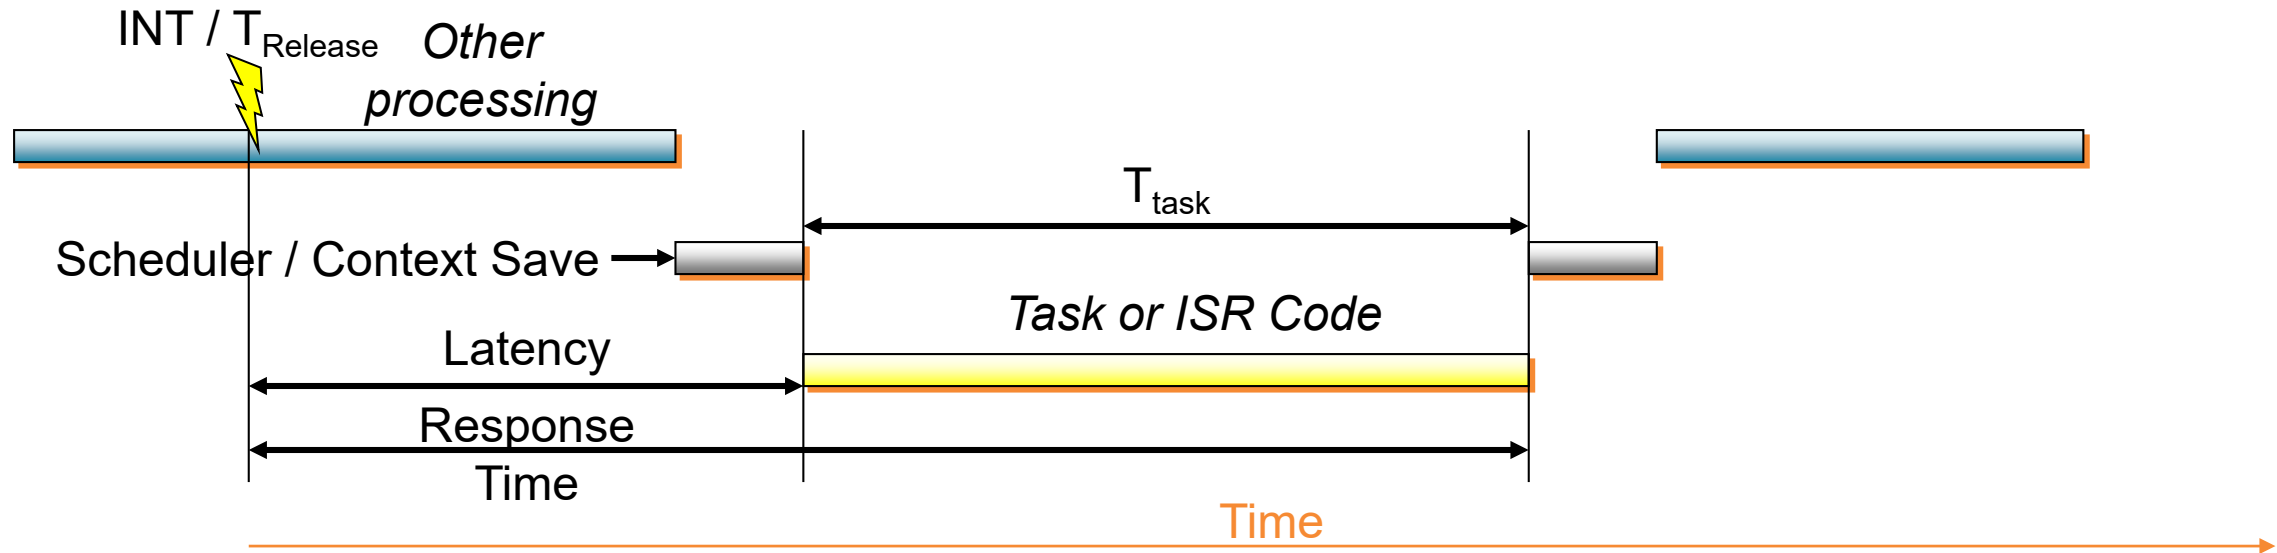


- Embedded systems rely on both MCU **hardware peripherals** and **software** to get everything done on time

CPU Scheduling

- MCU's Interrupt system provides a basic scheduling approach for CPU
 - “Run this subroutine every time this hardware event occurs”
 - Is adequate for simple systems
- More complex systems need to support multiple concurrent independent threads of execution
 - Use task scheduler to share CPU
 - Different approaches to task scheduling
- How do we make the processor responsive? (How do we make it do the right things at the right times?)
 - If we have more software threads than hardware threads, we need to share the processor.

Definitions



- $T_{\text{Release}}(i)$ = Time at which task (or interrupt) i requests service/is released/is ready to run
- $T_{\text{Latency}}(i)$ = Delay between release and *start of service* for task i
- $T_{\text{Response}}(i)$ = Delay between request for service and *completion of service* for task i
- $T_{\text{Task}}(i)$ = Time needed to perform computations for task i
- $T_{\text{ISR}}(i)$ = Time needed to perform interrupt service routine i

Scheduling Approaches

- Rely on MCU's hardware interrupt system to run right code
 - Event-triggered scheduling with interrupts
 - Works well for many simple systems

- Use software to schedule CPU's time
 - Static cyclic executive
 - Dynamic priority
 - Without task-level preemption
 - With task-level preemption

Event-Triggered Scheduling using Interrupts

- Basic architecture, useful for simple low-power devices
 - Very little code or time overhead
- Leverages built-in task dispatching of interrupt system
 - Can trigger ISRs with input changes, timer expiration, UART data reception, analog input level crossing comparator threshold
- Function types
 - Main function configures system and then goes to sleep
 - If interrupted, it goes right back to sleep
 - Only interrupts are used for normal program operation
- Example: bike performance monitor
 - Int1: wheel rotation
 - Int2: mode key
 - Int3: clock
 - Output: Liquid Crystal Display



Bike Performance Monitor Functions

Reset

```
Configure timer
inputs and
outputs

cur_time = 0;
rotations = 0;
tenth_miles = 0;

while (1) {
    sleep;
}
```

ISR 1: Wheel rotation

```
rotations++;
if (rotations >
    R_PER_MILE/10) {
    tenth_miles++;
    rotations = 0;
}
speed =
    circumference/
    (cur_time - prev_time);
compute avg_speed;
prev_time = cur_time;
return from interrupt
```

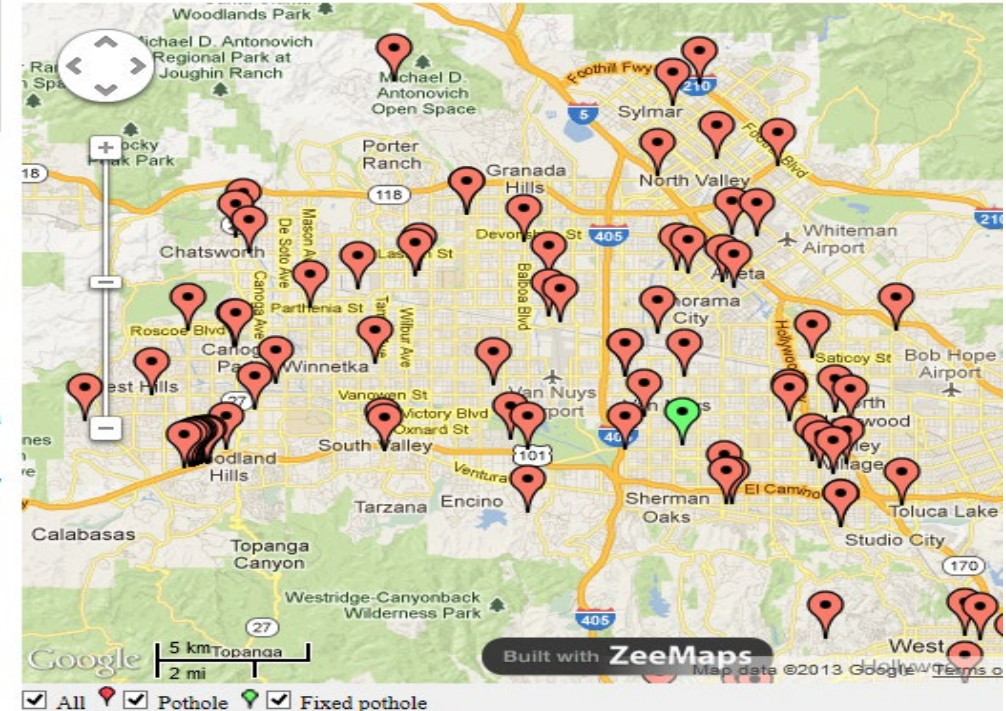
ISR 2: Mode Key

```
mode++;
mode = mode %
    NUM_MODES;
return from interrupt;
```

ISR 3: Time of Day Timer

```
cur_time ++;
lcd_refresh--;
if (lcd_refresh == 0) {
    convert tenth_miles
    and display
    convert speed
    and display
    if (mode == 0)
        convert cur_time
        and display
    else
        convert avg_speed
        and display
    lcd_refresh =
        LCD_REF_PERIOD
}
```

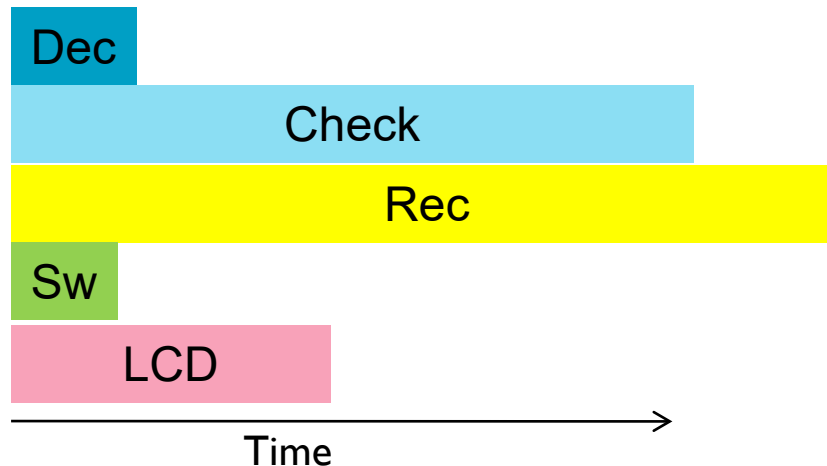
A More Complex Application



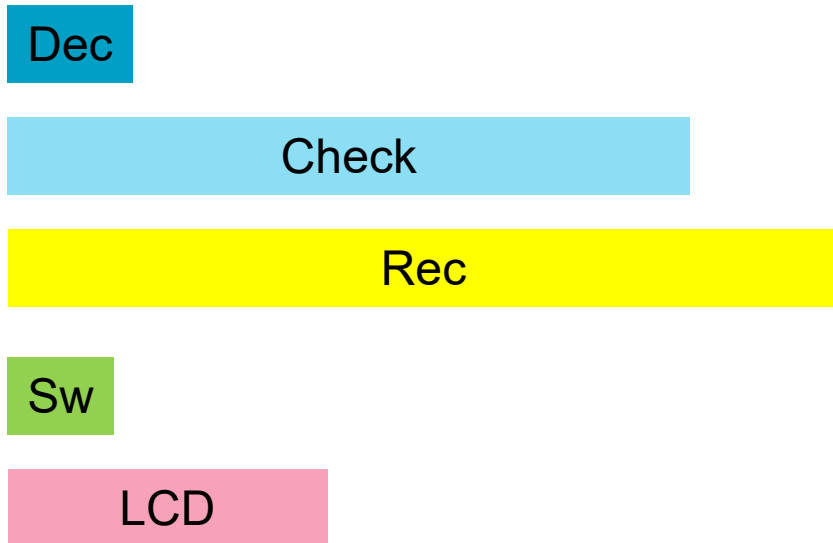
- GPS-based Pothole Alarm and Moving Map
 - Sounds alarm when approaching a pothole
 - Display's vehicle position on LCD
 - Also logs driver's position information
 - Hardware: GPS, user switches, speaker, LCD, flash memory

Application Software Tasks

- Dec: Decode GPS sentence to find current vehicle position.
- Check: Check to see if approaching any pothole locations. Takes longer as the number of potholes in database increases.
- Rec: Record position to flash memory. Takes a long time if erasing a block.
- Sw: Read user input switches. Run 10 times per second
- LCD: Update LCD with map. Run 4 times per second



How do we schedule these tasks?



- Task scheduling: Deciding which task should be run now
- Two fundamental questions
 - **Do we run tasks in the same order every time?**
 - Yes: Static schedule (cyclic executive, round-robin)
 - No: Dynamic, prioritized schedule
 - **Can one task preempt another, or must it wait for completion?**
 - Yes: Preemptive
 - No: Non-preemptive (cooperative, run-to-completion)

Static Schedule (Cyclic Executive)



- Pros

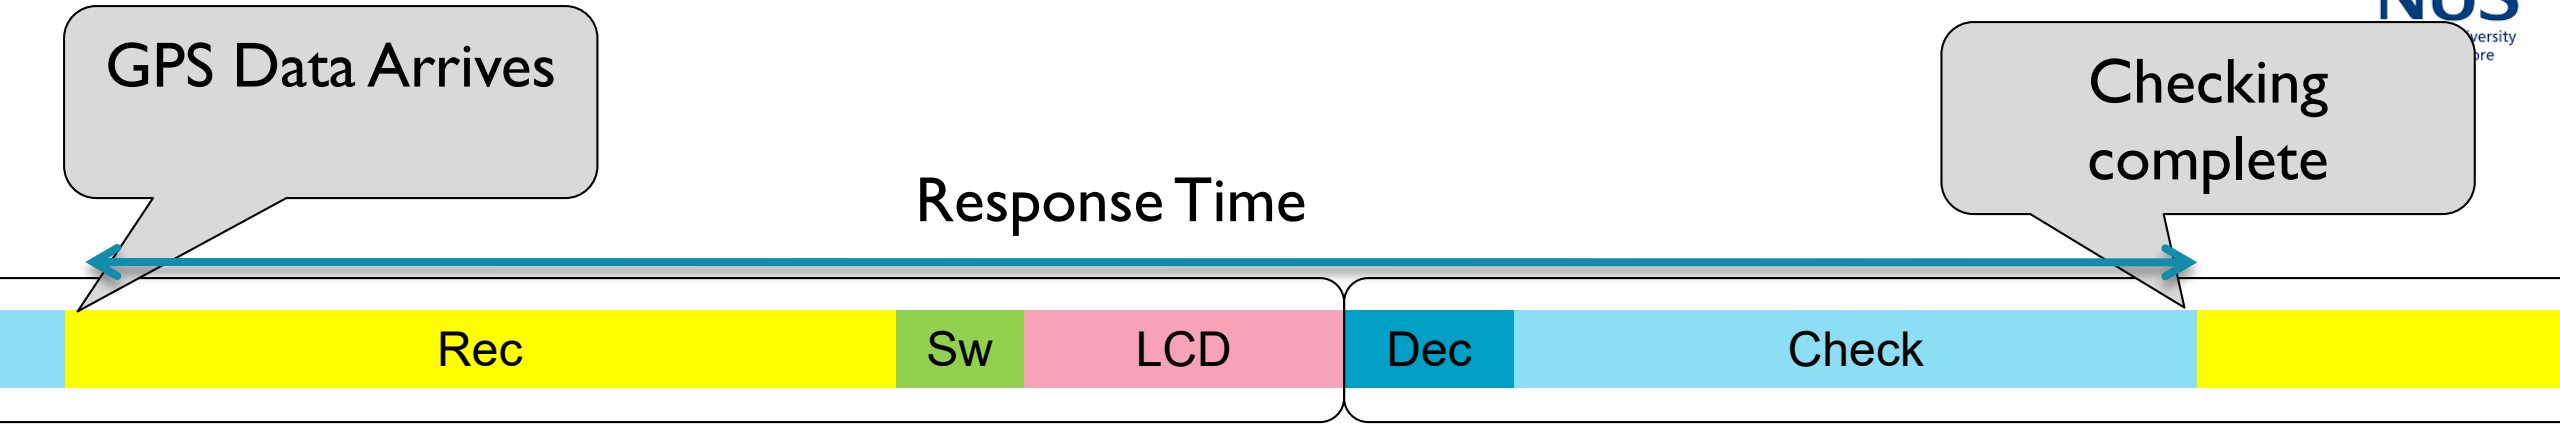
- Very simple

- Cons

- Always run the same schedule, regardless of changing conditions and relative importance of tasks.
- All tasks run at same rate. Changing rates requires adding extra calls to the function.
- Maximum delay is sum of all task run times. Polling/execution rate is $1/\text{maximum delay}$.

```
while (1) {  
    Dec () ;  
    Check () ;  
    Rec () ;  
    Sw () ;  
    LCD () ;  
}
```

Static Schedule Example



- What if we receive GPS position right after Rec starts running?
- Delays
 - Have to wait for Rec, Sw, LCD before we start decoding position with Dec.
 - Have to wait for Rec, Sw, LCD, Dec, Check before we know if we are approaching a pothole!

Dynamic Scheduling

- Allow schedule to be computed on-the-fly
 - Based on importance or something else
 - Simplifies creating multi-rate systems
- Schedule based on importance
 - Prioritization means that less important tasks don't delay more important ones
- How often do we decide what to run?
 - Coarse grain – After a task finishes. Called *Run-to-Completion* (RTC) or non-preemptive
 - Fine grain – Any time. Called *Preemptive*, since one task can preempt another.

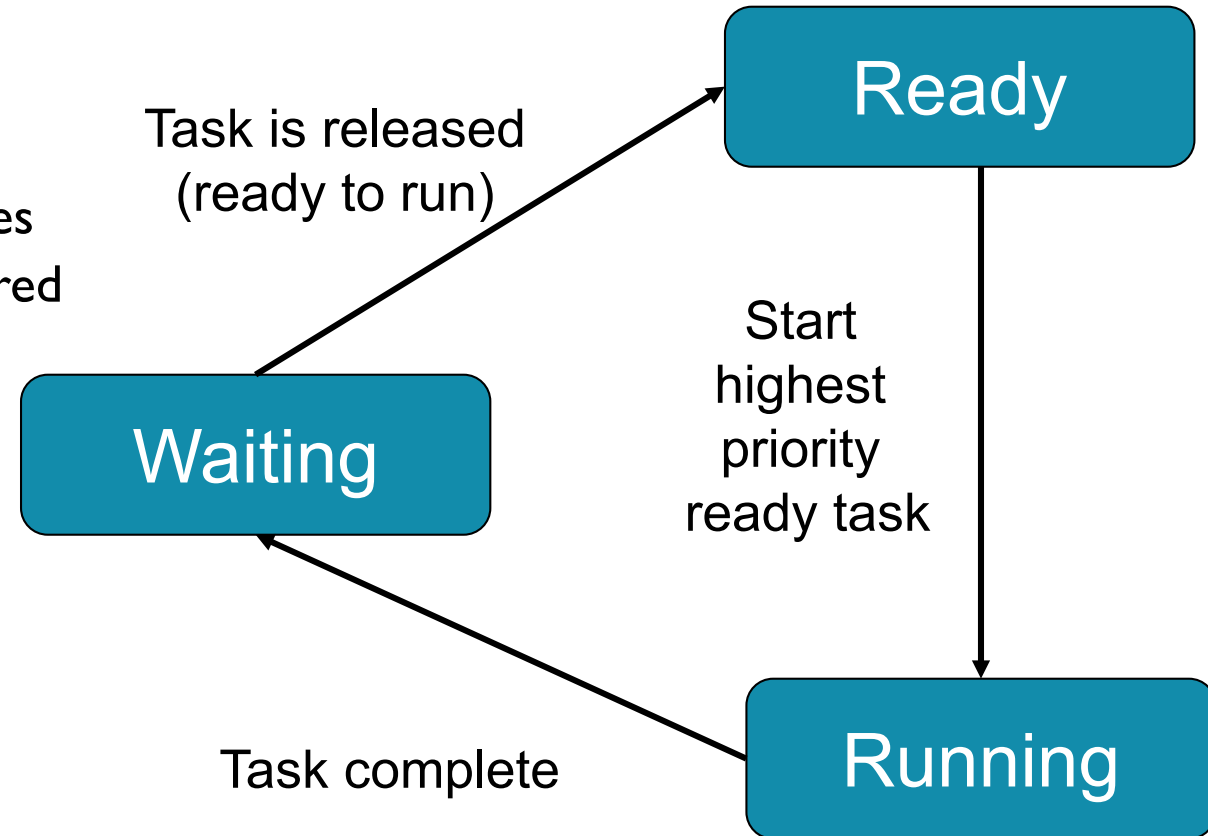
Dynamic RTC Schedule



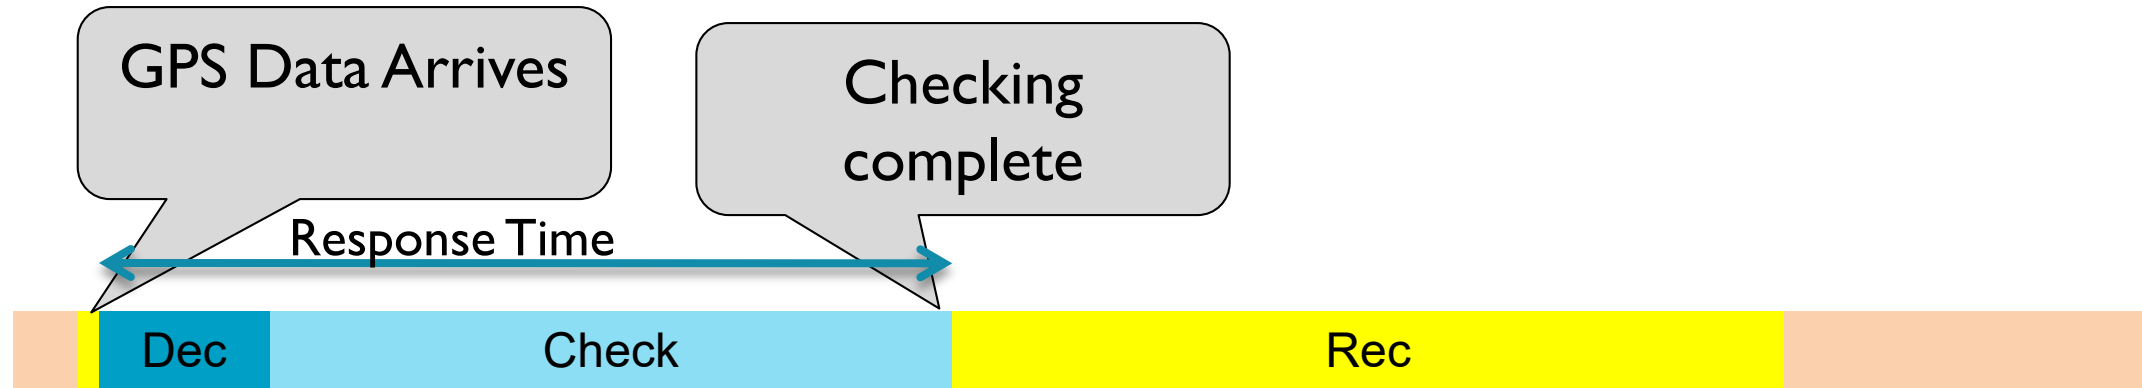
- What if we receive GPS position right after Rec starts running?
- Delays
 - Have to wait for Rec to finish before we start decoding position with Dec.
 - Have to wait for Rec, Dec, Check before we know if we are approaching a pothole

Task State and Scheduling Rules

- Scheduler chooses among *Ready* tasks for execution based on priority
- Scheduling Rules
 - If no task is running, scheduler starts the highest priority ready task
 - Once started, a task runs until it completes
 - Tasks then enter waiting state until triggered or released again





Dynamic Preemptive Schedule



- What if we receive GPS position right after Rec starts running?
- Delays
 - Scheduler switches out Rec so we can start decoding position with Dec immediately
 - Have to wait for Dec, Check to complete before we know if we are approaching a pothole

Comparison of Response Times



■ Pros

- Preemption offers best response time
 - Can do more processing (support more potholes, or higher vehicle speed)
 - Or can lower processor speed, saving money, power

■ Cons

- Requires more complicated programming, more memory
- Introduces vulnerability to data race conditions

Common Schedulers

- Cyclic executive - non-preemptive and static
- Run-to-completion - non-preemptive and dynamic
- Preemptive and dynamic

Cyclic Executive with Interrupts

- Two priority levels
 - `main` code – background
 - Interrupts – foreground
- Example of a **foreground / background system**
- Interrupt routines run in foreground (high priority)
 - Run when triggered
 - Handle most urgent work
 - Set flags to request processing by main loop
- Main user code runs in background
 - Uses “round-robin” approach to pick tasks, takes turns
 - Tasks do not preempt each other

```
BOOL DeviceARequest, DeviceBRequest,
DeviceCRequest;
void interrupt HandleDeviceA() {
    /* do A's urgent work */
    ...
    DeviceARequest = TRUE;
}
void main(void) {
    while (TRUE) {
        if (DeviceARequest) {
            FinishDeviceA();
        }
        if (DeviceBRequest) {
            FinishDeviceB();
        }
        if (DeviceCRequest) {
            FinishDeviceC();
        }
    }
}
```

Run-To-Completion Scheduler

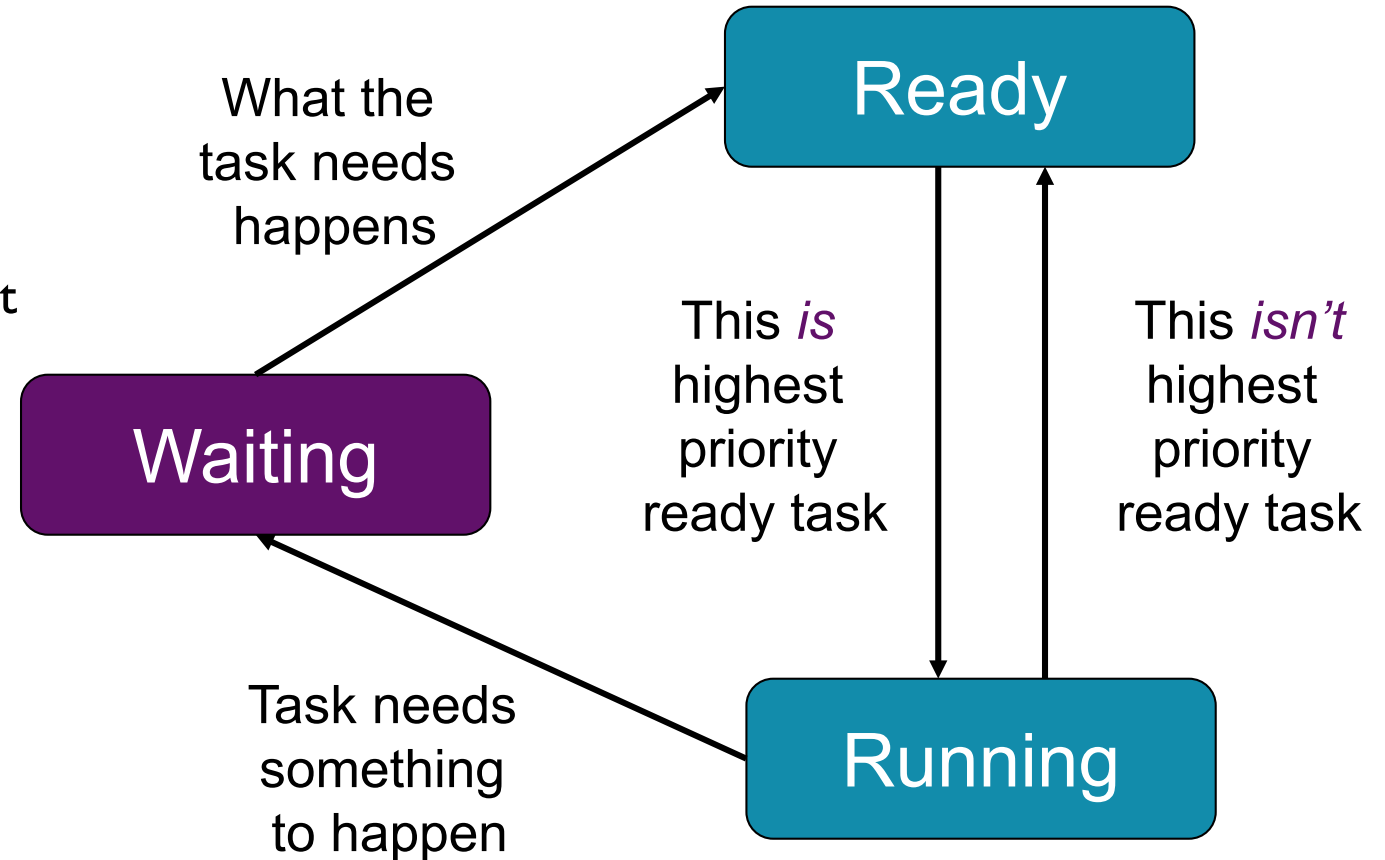
- Use a **scheduler** function to run task functions at the right rates
 - Table stores information per task
 - Period: How many ticks between each task release
 - Release Time: how long until task is ready to run
 - ReadyToRun: task is ready to run immediately
 - Scheduler runs forever, examining schedule table which indicates tasks which are ready to run (have been “released”)
 - A periodic timer interrupt triggers an ISR, which updates the schedule table
 - Decrements “time until next release”
 - If this time reaches 0, set the task’s Run flag and reload its time with the period
- Follows a “run-to-completion” model
 - A task’s execution is **not interleaved** with any other task
 - Only ISRs can interrupt a task
 - After ISR completes, the previously-running task resumes
- Priority is typically static, so can use a table with highest priority tasks first for a fast, simple scheduler implementation.

Preemptive Scheduler

- Task functions need not run to completion, but can be interleaved with each other
 - Simplifies writing software
 - Improves response time
 - Introduces new potential problems
- Worst case response time for highest priority task does not depend on other tasks, only ISRs and scheduler
 - Lower priority tasks depend only on higher priority tasks

Task State and Scheduling Rules

- Scheduler chooses among *Ready* tasks for execution based on priority
- Scheduling Rules
 - A task's activities may lead it to *waiting* (*blocked*)
 - A *waiting* task never gets the CPU. It must be signaled by an ISR or another task.
 - Only the scheduler moves tasks between *ready* and *running*




What's an RTOS?

- What does Real-Time mean?
 - Can calculate and guarantee the *maximum response time* for each task and interrupt service routine
 - This “bounding” of response times allows use in hard-real-time systems (which have deadlines which must be met)


- What's in the RTOS
 - Task Scheduler
 - Preemptive, prioritized to minimize response times
 - Interrupt support
 - Core Integrated RTOS services
 - Inter-process communication and synchronization (safe data sharing)
 - Time management
 - Optional Integrated RTOS services
 - *I/O abstractions?*
 - *memory management?*
 - *file system?*
 - *networking support?*
 - *GUI??*

Comparison of Timing Dependence


Non-preemptive Static



Non-preemptive Dynamic




Preemptive Dynamic




- Code can be delayed by everything at same level (in oval) or above

Comparison of RAM Requirements


*Non-preemptive
Static*



*Non-preemptive
Dynamic*



*Preemptive
Dynamic*



- Preemption requires space for each stack
- Need space for all static variables (including globals)

SOFTWARE ENGINEERING FOR EMBEDDED SYSTEMS

THIS SLIDE ONWARDS IS FOR YOUR OWN READING...

Good Enough Software, Soon Enough

- How do we make software *correct enough* without going bankrupt?
 - Need to be able to develop (and test) software efficiently
- Follow a good plan
 - Start with customer requirements
 - Design architectures to define the building blocks of the systems (tasks, modules, etc.)
 - Add missing requirements
 - Fault detection, management and logging
 - Real-time issues
 - Compliance to a firmware standards manual
 - Fail-safes
 - Create detailed design
 - Implement the code, following a good development process
 - Perform frequent design and code reviews
 - Perform frequent testing (unit and system testing, preferably automated)
 - Use revision control to manage changes
 - Perform post-mortems to improve development process

What happens when the plan meets reality?

- We want a robust plan which considers likely risks
 - What if the code turns out to be a lot more complex than we expected?
 - What if there is a bug in our code (or a library)?
 - What if the system doesn't have enough memory or throughput?
 - What if the system is too expensive?
 - What if the lead developer quits?
 - What if the lead developer is incompetent, lazy, or both (and *won't* quit!)?
 - What if the rest of the team gets sick?
 - What if the customer adds new requirements?
 - What if the customer wants the product two months early?
- Successful software engineering depends on balancing many factors, many of which are non-technical!

Risk Reduction

- Plan to the work to accommodate risks
- Identify likely risks up front
 - Historical problem areas
 - New implementation technologies
 - New product features
 - New product line
- Severity of risk is a combination of likelihood and impact of failure

Software Lifecycle Concepts

- Coding is the most visible part of a software development process but is not the only one
- Before we can code, we must know
 - What must the code do? *Requirements specification*
 - How will the code be structured? *Design specification*
 - (only at this point can we start writing code)
- How will we know if the code works? *Test plan*
 - Best performed when defining requirements
- The software will likely be enhanced over time - *Extensive downstream modification and maintenance!*
 - Corrections, adaptations, enhancements & preventive maintenance

Product Development Lifecycle

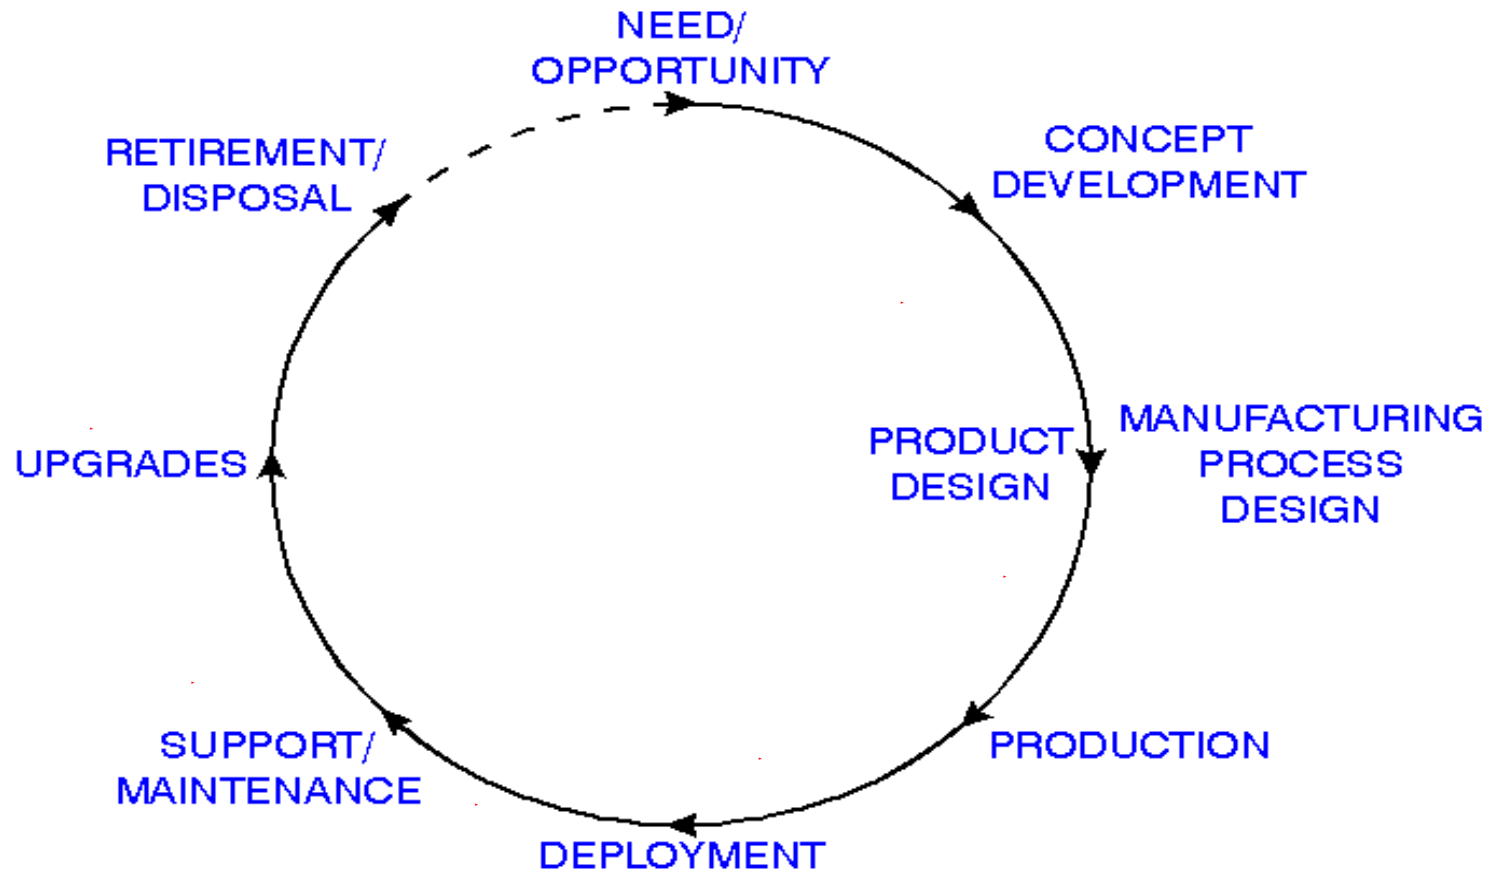



Diagram: Phil Koopman,
Carnegie Mellon
University

- With all this code development and modification, it is worth putting extra effort into simplifying code development activities: understanding, maintaining, enhancing, testing

Requirements

- Why embedded projects fail: **#Vague Requirements**
- Types of requirements
 - Functional - what the system needs to do
 - Nonfunctional - emergent system behaviors such as response time, reliability, energy efficiency, safety, etc.
 - Constraints - limit design choices
- Representations
 - Text – Liable to be incomplete, bloated, ambiguous, even contradictory
 - Diagrams (state charts, flow charts, message sequence charts)
 - Concise
 - Can often be used as design documents
- Traceability
 - Each requirement should be verifiable with a test
- Stability
 - Requirements churn leads to inefficiency and often “recency” problems (most recent requirement change is assumed to be most important)










Design Before Coding



- Why embedded projects fail: #: **Starting coding too soon**
- Underestimating the complexity of the needed software is a very common risk
- Writing code locks you in to specific implementations
 - Starting too early may paint you into a corner
- Benefits of **designing** system before **coding**
 - Get early insight into system's complexity, allowing more accurate effort estimation and scheduling
 - Can use design diagrams rather than code to discuss what system should do and how.
 - Can use design diagrams in documentation to simplify code maintenance and reduce risks of staff turnover

Design Before Coding

- How much of the system do you design before coding?



Development Models

- How do we schedule these pieces?
- Consider amount of development risk
 - New MCU?
 - Exceptional requirements (throughput, power, safety certification, etc.)
 - New product?
 - New customer?
 - Changing requirements?
- Choose model based on risk
 - Low: Can create detailed plan. Big-up-front design, waterfall
 - High: Use iterative or agile development method. Prototype high-risk parts first

Waterfall (Idealized)

- Plan the work, and then work the plan
- BUFD: Big Up-Front Design
- Model implies that we and the customers know
 - All of the requirements up front
 - All of the interactions between components, etc.
 - How long it will take to write the software and debug it

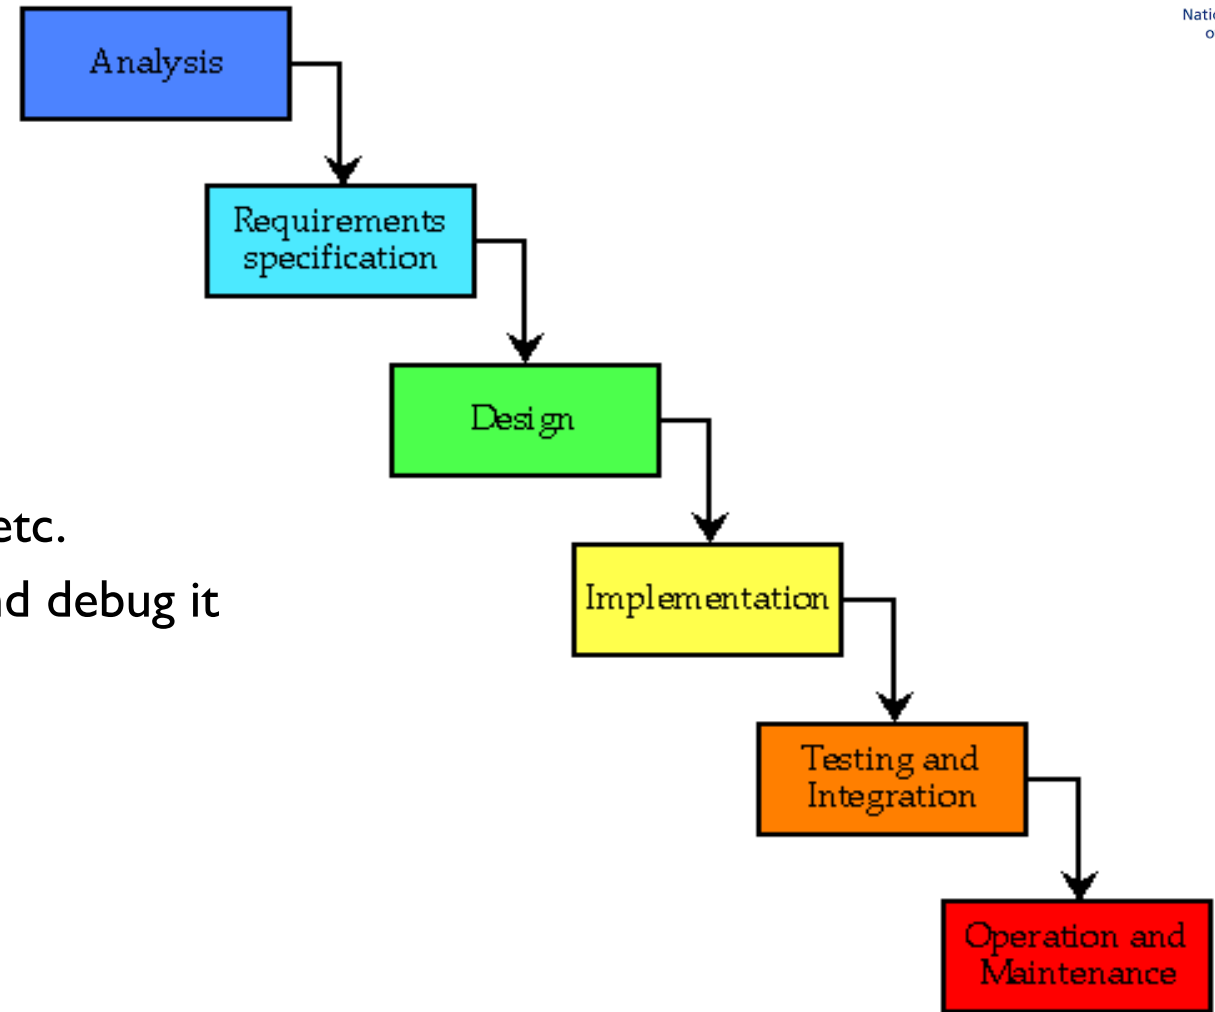


Diagram: Jon McCormack, Monash University

Waterfall (As Implemented)

- Reality: We are not omniscient, so there is plenty of backtracking

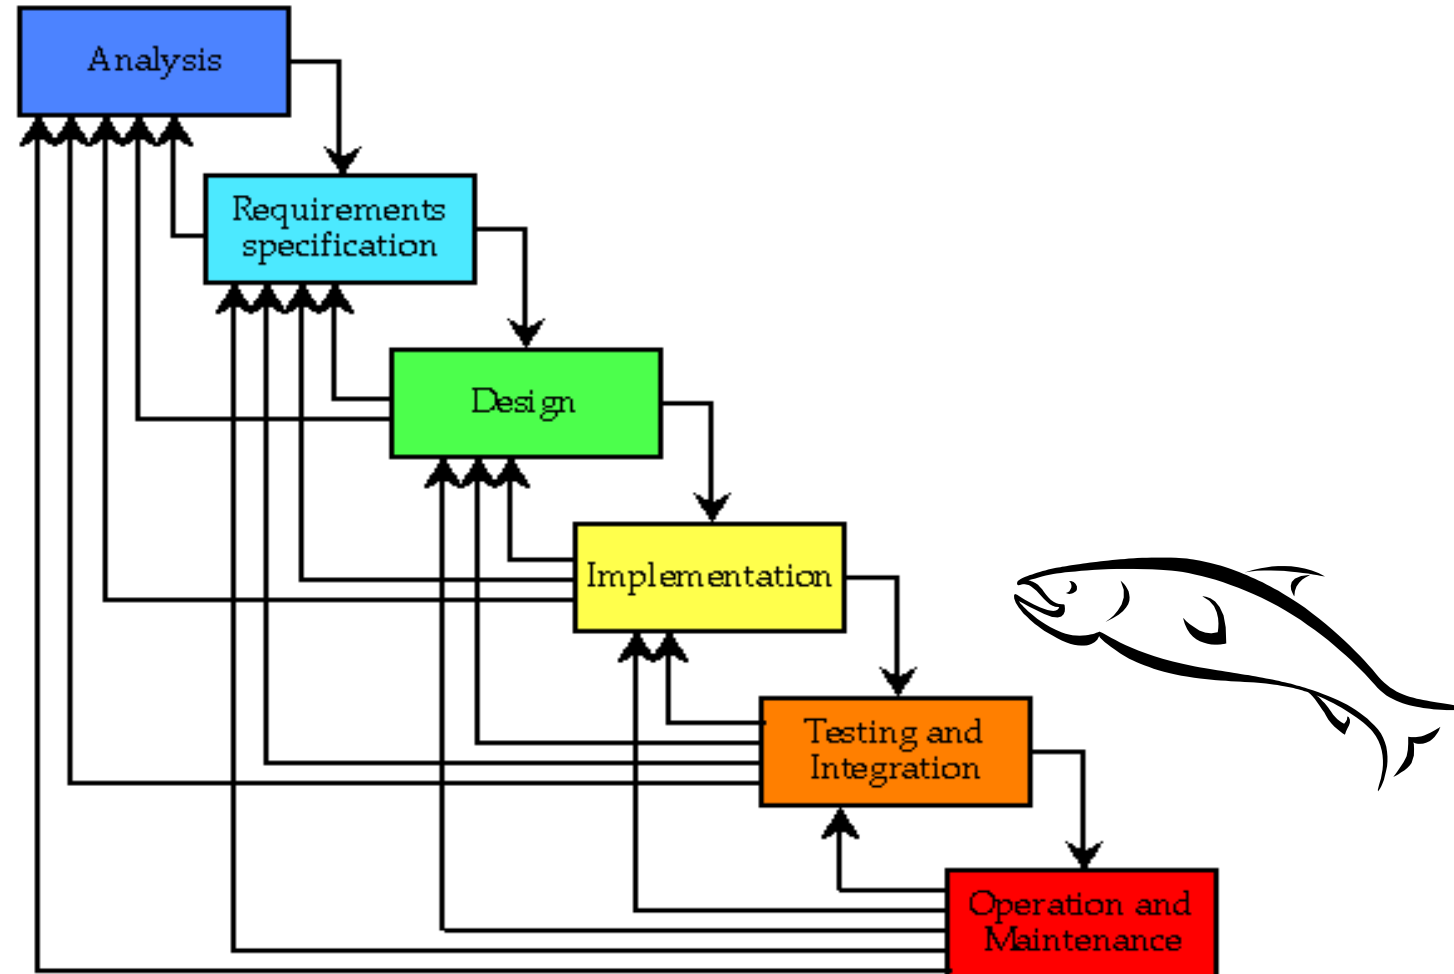


Diagram: Jon McCormack, Monash University

I. Requirements Specification and Validation Plan

- Result of Requirements Analysis
- Should contain:
 - *Introduction* with goals and objectives of system
 - *Description of problem* to solve
 - *Functional description*
 - provides a “processing narrative” per function
 - lists and justifies design constraints
 - explains performance requirements
 - *Behavioral description* shows how system reacts to internal or external events and situations
 - State-based behavior
 - General control flow
 - General data flow
 - *Validation criteria*
 - tell us how we can decide that a system is acceptable. (*Are we done yet?*)
 - is the foundation for a validation test plan
 - *Bibliography and Appendix* refer to all documents related to project and provide supplementary information

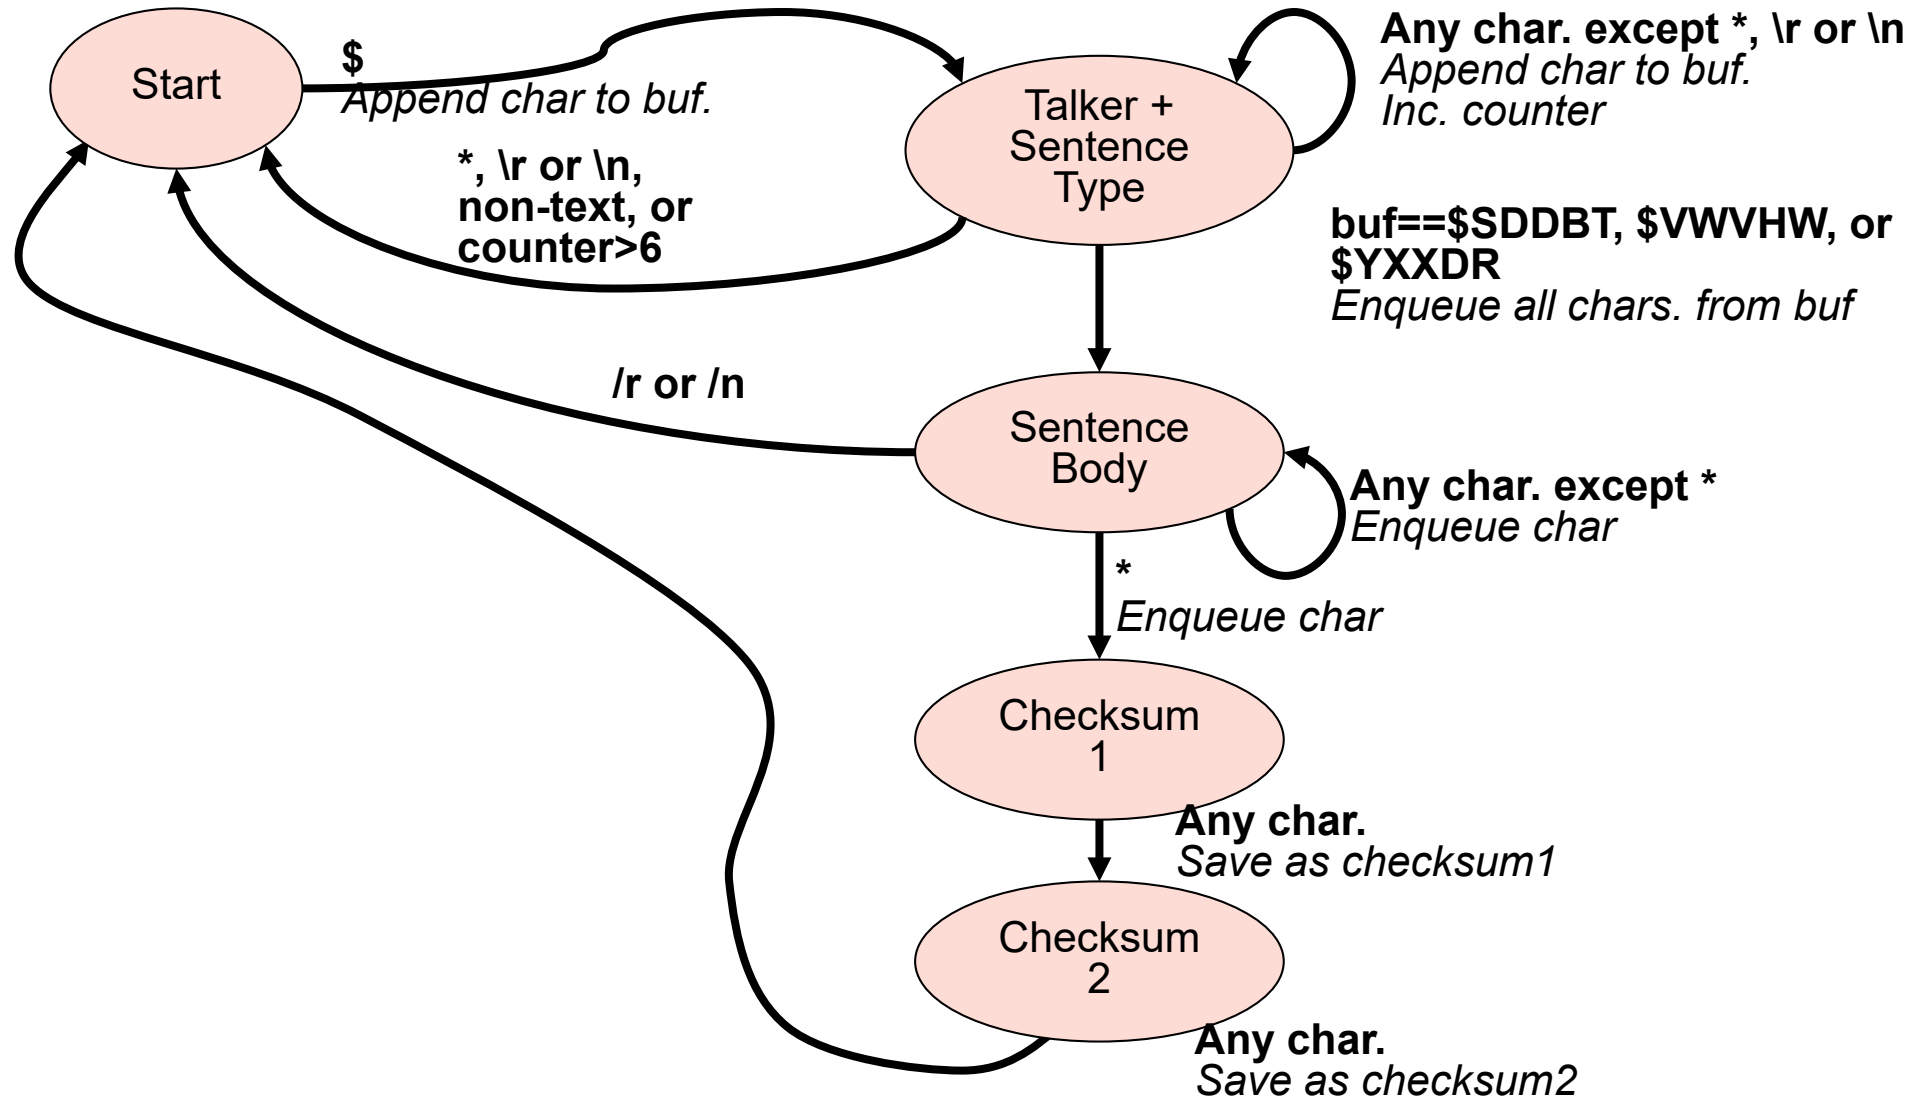
2. Architectural (High-Level) Design

- Architecture defines the structure of the system
 - Components
 - Externally visible properties of components
 - Relationships among components
- Architecture is a representation which lets the designer...
 - Analyze the design's effectiveness in meeting requirements
 - Consider alternative architectures early
 - Reduce down-stream implementation risks
- Architecture matters because...
 - It's small and simple enough to fit into a single person's brain (as opposed to comprehending the entire program's source code)
 - It gives stakeholders a way to describe and therefore discuss the system

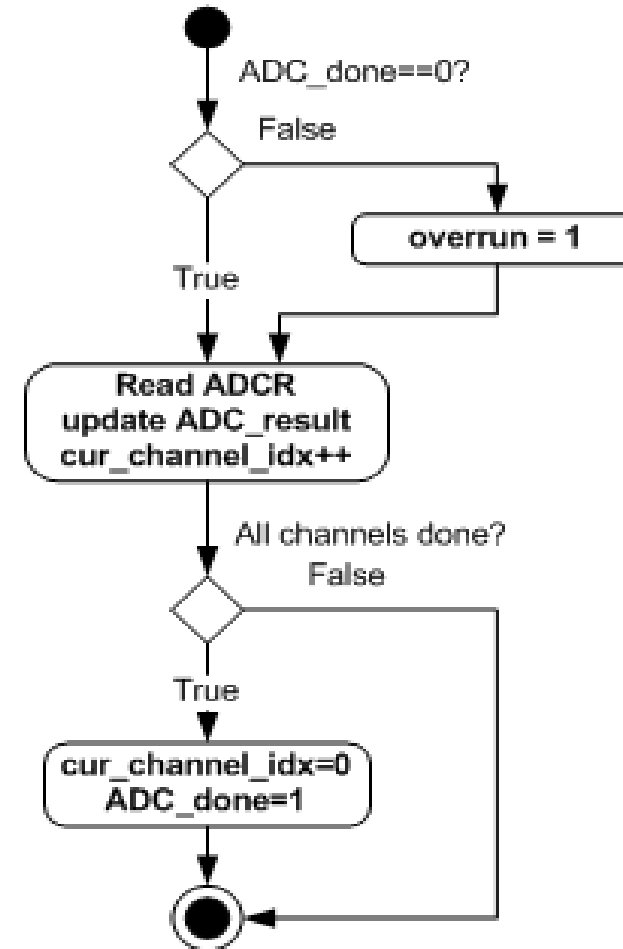
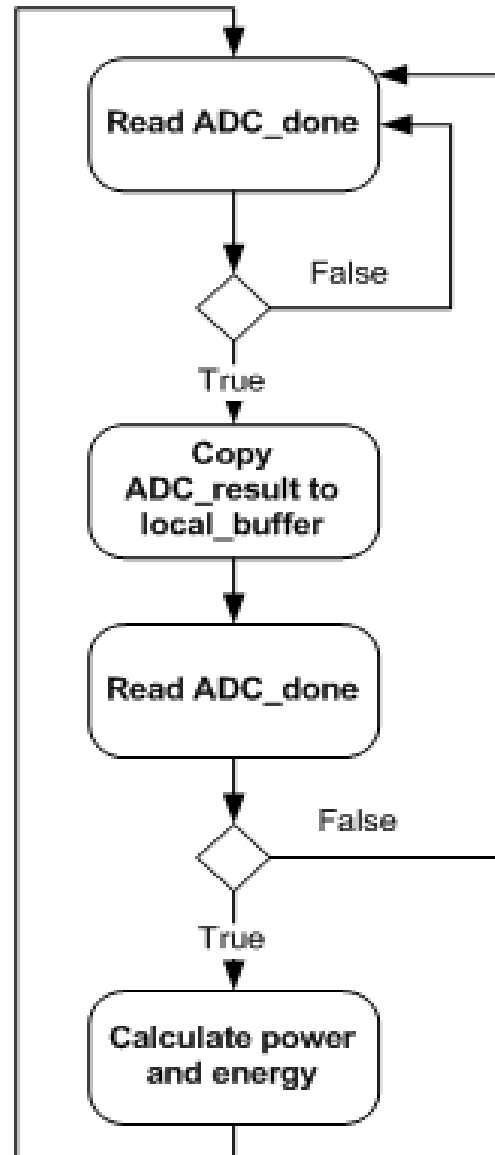
3. Detailed Design

- Describe aspects of how system behaves
 - Flow charts for control or data
 - State machine diagram
 - Event sequences
- Graphical representations very helpful
 - Can provide clear, single-page visualization of what system component should do
- Unified Modeling Language (UML)
 - Provides many types of diagrams
 - Some are useful for embedded system design to describe structure or behavior

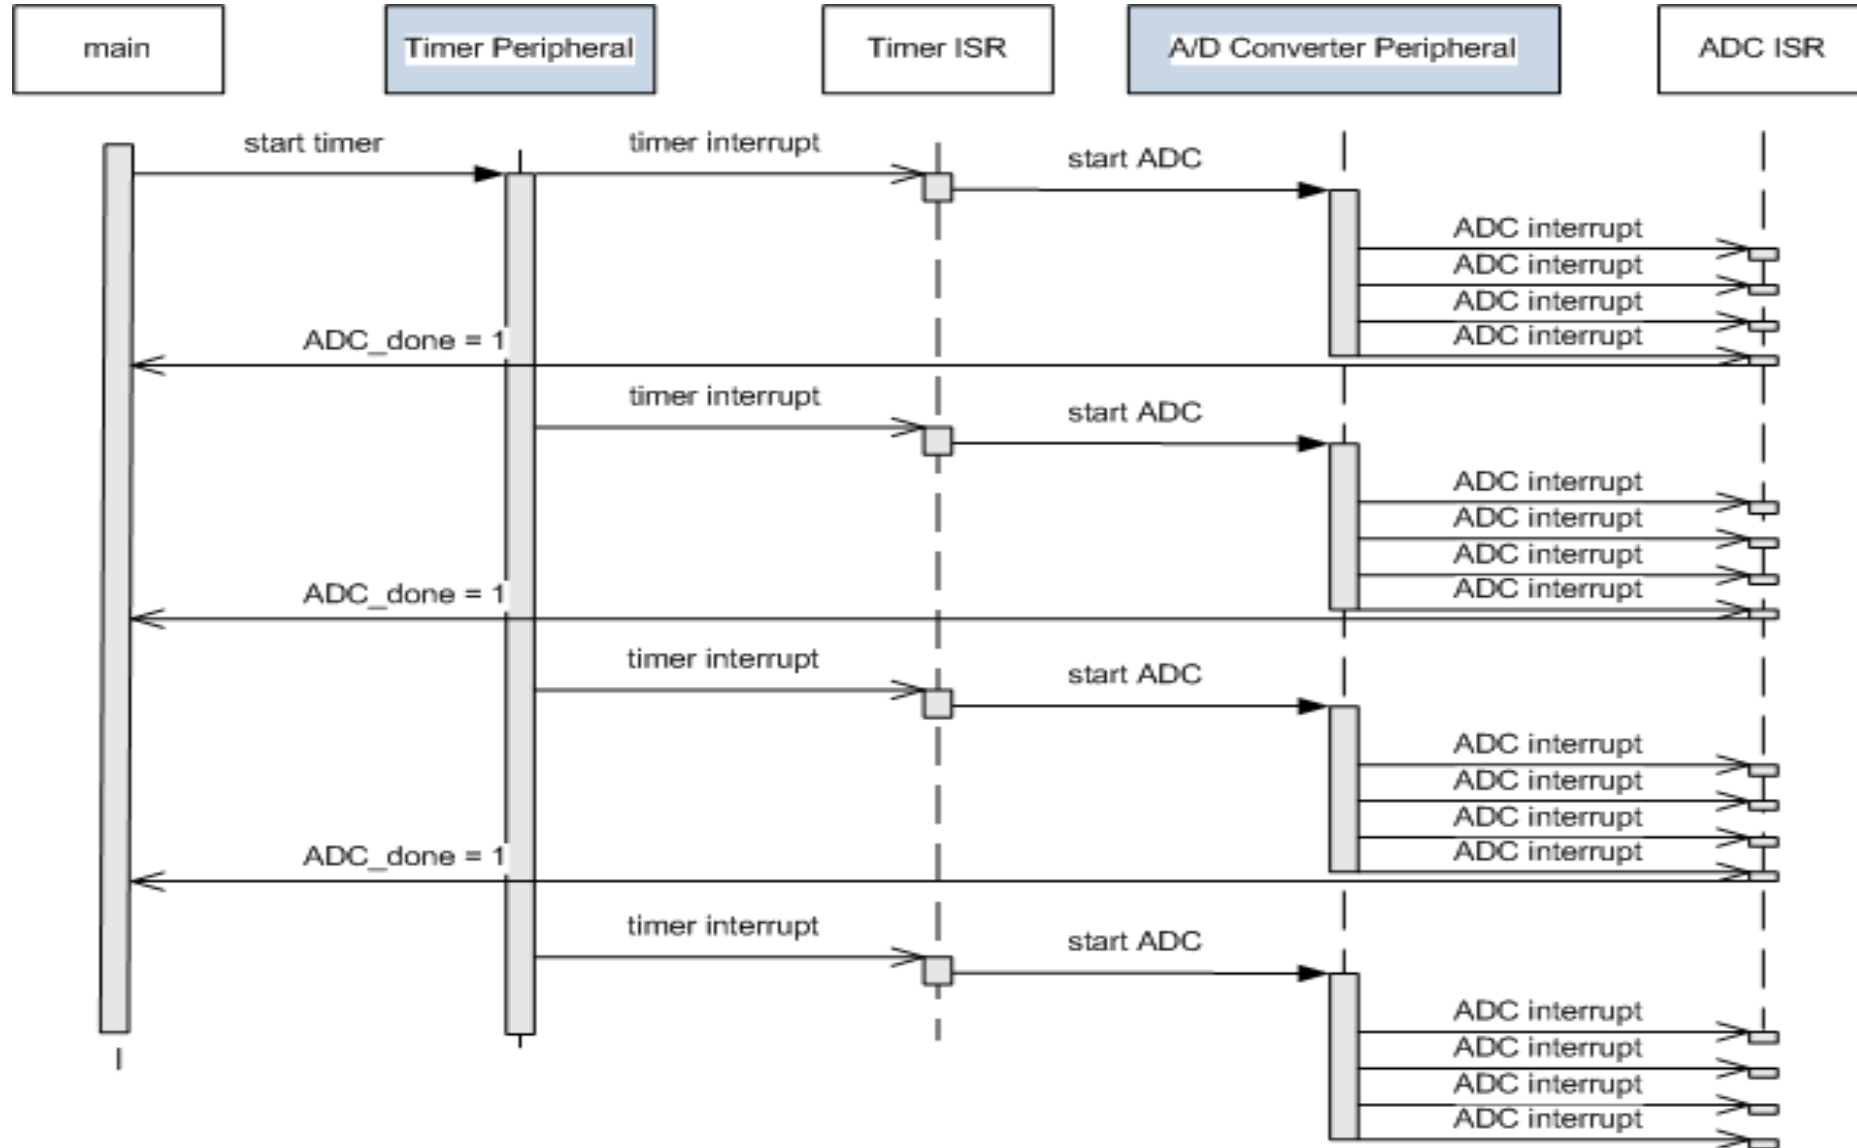
State Machine for Parsing NMEA-0183



Flowcharts



Sequence of Interactions between Components



4. Coding and Code Inspections

- Coding driven directly by Detailed Design Specification
- Use a version control system while developing the code
- Follow a coding standard
 - Eliminate stylistic variations which make understanding code more difficult
 - Avoid known questionable practices
 - Spell out best practices to make them easier to follow
- Perform code reviews
- Perform unit testing on modules as appropriate

Peer Code Review

- Inspect the code before testing it
- Extensive positive industry results from code inspections
 - IBM removed 82% of bugs
 - 9 hours saved by finding each defect
 - For AT&T quality rose by 1000% and productivity by 14%
- Finds bugs which testing often misses
 - 80% of the errors detected by HP's inspections were unlikely to be caught by testing
 - HP, Shell Research, Bell Northern, AT&T: inspections 20-30x more efficient than testing

5. Software Testing

- Testing IS NOT “the process of verifying the program works correctly”
 - The program probably won’t work correctly in all possible cases
 - Professional programmers have 1-3 bugs per 100 lines of code after it is “done”
 - Testers shouldn’t try to prove the program works correctly (impossible)
 - If you want and expect your program to work, you’ll unconsciously miss failure because human beings are inherently biased
- The purpose of testing is to find problems quickly
 - Does the software violate the specifications?
 - Does the software violate unstated requirements?
- The purpose of finding problems is to fix the ones which matter
 - Fix the most important problems, as there isn’t enough time to fix all of them
 - The *Pareto Principle* defines “the vital few, the trivial many”
 - Bugs are uneven in frequency – a vital few contribute the majority of the program failures. Fix these first.

Approaches to Testing

■ Incremental Testing

- Code a function and then test it (*module/unit/element testing*)
- Then test a few working functions together (*integration testing*)
 - Continue enlarging the scope of tests as you write new functions
- Incremental testing requires extra code for the *test harness*
 - A *driver* function calls the function to be tested
 - A *stub* function might be needed to simulate a function called by the function under test, and which returns or modifies data.
 - The test harness can *automate* the testing of individual functions to detect later bugs

■ Big Bang Testing

- Code up all of the functions to create the system
- Test the complete system
 - Plug and pray

Why Test Incrementally?

- Finding out what failed is much easier
 - With Big Bang, since no function has been thoroughly tested, most probably have bugs
 - Now the question is “Which bug in which module causes the failure I see?”
 - Errors in one module can make it difficult to test another module
 - Errors in fundamental modules (e.g. kernel) can appear as bugs in other many other dependent modules
- Less finger pointing = happier SW development team
 - It's clear who made the mistake, and it's clear who needs to fix it
- Better automation
 - Drivers and stubs initially require time to develop, but save time for future testing

6. Perform Project Retrospectives

- Goals – improve your engineering processes
 - Extract all useful information learned from the just-completed project – provide “virtual experience” to others
 - Provide positive non-confrontational feedback
 - Document problems and solutions clearly and concisely for future use
- Basic rule: problems need solutions
- Often small changes improve performance, but are easy to forget

Example Postmortem Structure

- Product
 - Bugs
 - Software design
 - Hardware design
- Process
 - Code standards
 - Code interfacing
 - Change control
 - How we did it
 - Team coordination
- Support
 - Tools
 - Team burnout
 - Change orders
 - Personnel availability

The End!

- Thank You!
- Lets go onto more exciting stuff! 😊