

Memory


Ravi Suppiah
Lecturer, NUS SoC

Current Presentation Content

- Different components of the process memory
 - text, data, BSS, heap, stack
- How is memory allocated to processes and the heap within a process
 - Memory allocation and deallocation
 - External and internal fragmentation
 - Memory allocation bitmap, Free list
 - Allocation policies: First fit, next fit, best fit, worst fit
 - Compaction
 - Multiple free list
 - Buddy system

Process Memory: Code

- Executable code
 - Program binary and any other libraries it loads
- OS knows everything in advance
 - Knows amount of space needed
 - Knows the contents of the memory
- Known as the “text” segment




Process Memory: “Static” Data

- Variables that exist for the entire program
 - Global variables and “static” local variables
 - Amount of space required is known in advance
- Data: Initialized in the code
 - Initial value specified by the programmer
 - `int x = 97;`
 - Memory is initialized with this value
- BSS: not initialized in the code
 - Initial value is not specified
 - `int x;`
 - All memory initialized to 0
 - BSS stands for “Block Stated by Symbol”




Process Memory: Dynamic Memory

- Memory allocated while program is running
 - Allocated using malloc () function
 - Deallocated using free () function
- OS knows nothing in advance
 - Doesn't know the amount of space
 - Doesn't know the contents
- So need to allow room to grow
 - Known as the “heap”




Process Memory: Temporary Variables

- Temporary memory during lifetime of a function
 - Storage for function parameters and local variables
- Need to support nested function calls
 - One function calls another
 - Store the variables of calling function
 - Know where to return when done
- So must allow room to grow
 - Known as the “stack”
 - Push on the stack as new function is called
 - Pop off the stack as the function ends




Memory Layout: Summary

- **Text:** Code
- **Data:** Initialized global and static variables
- **BSS:** Uninitialized global and static variables
- **Heap:** Dynamic Memory
- **Stack:** Local variables



More about Stack

```
void main (void ) {  
    int a=10, b=20, c;  
    c = f (a, b);  
}  
int f (int x, int y){  
    int temp;  
    temp = x + y;  
    return temp;  
}
```



Memory Layout Example

```
int string_length = 8;
```

```
int iSize;
```

```
char * f (void)
```

```
{
```

```
    char *p;
```

```
    static int count = 0;
```


```
        iSize = string_length ;
```

```
    p = malloc(iSize);
```

```
    count ++;
```

```
    return p;
```

```
}
```



Memory Layout Example: Text

```
int string_length = 8;
```

```
int iSize;
```

```
char * f (void)
```

```
{
```

```
    char *p;
```

```
    static int count = 0;
```


```
        iSize = string_length ;
```

```
    p = malloc(iSize);
```

```
    count ++;
```

```
    return p;
```

```
}
```



Memory Layout Example: Data

```
int string_length = 8;
```

```
int iSize;
```

```
char * f (void)
```

```
{
```

```
    char *p;
```

```
    static int count = 0;
```

```
        iSize = string_length ;
```


```
    p = malloc(iSize);
```

```
    count ++;
```

```
    return p;
```

```
}
```

```
||
```



Memory Layout Example: BSS

```
int string_length = 8;
```

```
int iSize;
```

```
char * f (void)
```

```
{
```

```
    char *p;
```

```
    static int count = 0;
```


```
        iSize = string_length ;
```

```
    p = malloc(iSize);
```

```
    count ++;
```

```
    return p;
```

```
}
```



Memory Layout Example: Heap

```
int string_length = 8;
```

```
int iSize;
```

```
char * f (void)
```

```
{
```

```
    char *p;
```

```
    static int count = 0;
```


```
        iSize = string_length ;
```

```
    p = malloc(iSize);
```

```
    count ++;
```

```
    return p;
```

```
}
```



Memory Layout Example: Stack

```
int string_length = 8;
```

```
int iSize;
```

```
char * f (void)
```

```
{
```

```
    char *p;
```

```
    static int count = 0;
```


```
        iSize = string_length ;
```

```
    p = malloc(iSize);
```

```
    count++;
```


```
    return p;
```

```
}
```



Memory Allocation & Deallocation

- How and when is memory allocated?
 - Global and static variable @ program startup
 - Local variables @ function call
 - Dynamic memory @ malloc
- How is memory deallocated?
 - Global and static variables @ program finish
 - Local variable @ function return
 - Dynamic memory @ free
- All memory deallocated when program ends
 - Good programming practice to free allocated memory




Heap: Dynamic Memory

```
#include <stdlib.h>
void *malloc(size_t size);
void free(void *ptr);
```

➔


```
char *p1 = malloc(3);
char *p2 = malloc(1);
char *p3 = malloc(4);
free(p2);
char *p4 = malloc(6);
free(p3);
char *p5 = malloc(2);
free(p1);
free(p4);
free(p5);
```



Heap: Dynamic Memory

```
#include <stdlib.h>
void *malloc(size_t size);
void free(void *ptr);
```


```
➔ char *p1 = malloc(3);
   char *p2 = malloc(1);
   char *p3 = malloc(4);
   free(p2);
   char *p4 = malloc(6);
   free(p3);
   char *p5 = malloc(2);
   free(p1);
   free(p4);
   free(p5);
```



Heap: Dynamic Memory

```
#include <stdlib.h>
void *malloc(size_t size);
void free(void *ptr);
```


```
char *p1 = malloc(3);
char *p2 = malloc(1);
➔ char *p3 = malloc(4);
free(p2);
char *p4 = malloc(6);
free(p3);
char *p5 = malloc(2);
free(p1);
free(p4);
free(p5);
```



Heap: Dynamic Memory

```
#include <stdlib.h>
void *malloc(size_t size);
void free(void *ptr);
```

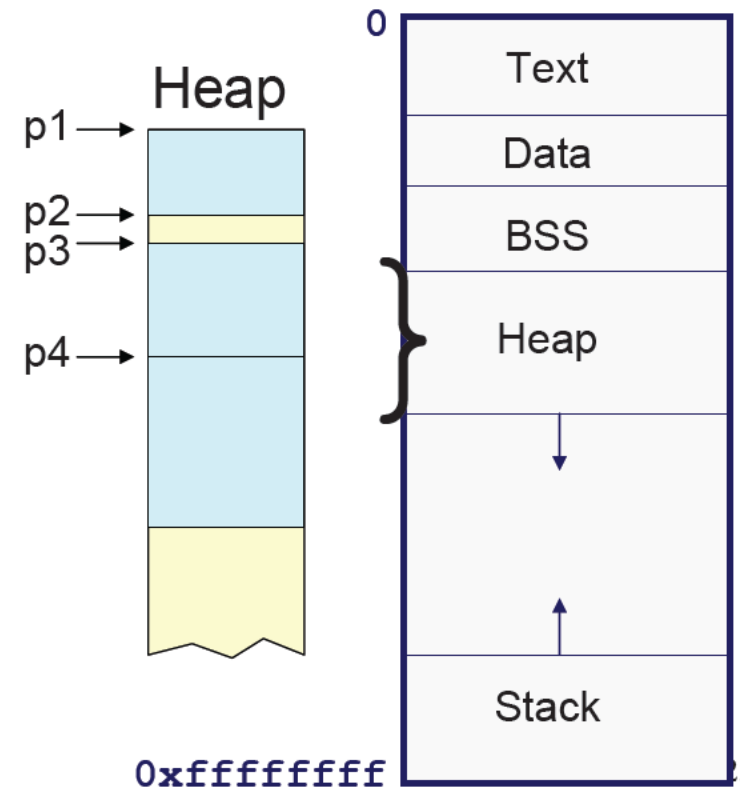
```
char *p1 = malloc(3);
char *p2 = malloc(1);
char *p3 = malloc(4);
➔ free(p2);
char *p4 = malloc(6);
free(p3);
char *p5 = malloc(2);
free(p1);
free(p4);
free(p5);
```



Heap: Dynamic Memory

```
#include <stdlib.h>
void *malloc(size_t size);
void free(void *ptr);
```


```
char *p1 = malloc(3);
char *p2 = malloc(1);
char *p3 = malloc(4);
free(p2);
➔ char *p4 = malloc(6);
free(p3);
char *p5 = malloc(2);
free(p1);
free(p4);
free(p5);
```



Heap: Dynamic Memory

```
#include <stdlib.h>
void *malloc(size_t size);
void free(void *ptr);
```


```
char *p1 = malloc(3);
char *p2 = malloc(1);
char *p3 = malloc(4);
free(p2);
char *p4 = malloc(6);
➔ free(p3);
char *p5 = malloc(2);
free(p1);
free(p4);
free(p5);
```



Heap: Dynamic Memory

```
#include <stdlib.h>
void *malloc(size_t size);
void free(void *ptr);
```


```
char *p1 = malloc(3);
char *p2 = malloc(1);
char *p3 = malloc(4);
free(p2);
char *p4 = malloc(6);
free(p3);
➔ char *p5 = malloc(2);
free(p1);
free(p4);
free(p5);
```



Heap: Dynamic Memory

```
#include <stdlib.h>
void *malloc(size_t size);
void free(void *ptr);
```


```
char *p1 = malloc(3);
char *p2 = malloc(1);
char *p3 = malloc(4);
free(p2);
char *p4 = malloc(6);
free(p3);
char *p5 = malloc(2);
➔ free(p1);
   free(p4);
   free(p5);
```



Heap: Dynamic Memory

```
#include <stdlib.h>
void *malloc(size_t size);
void free(void *ptr);
```


```
char *p1 = malloc(3);
char *p2 = malloc(1);
char *p3 = malloc(4);
free(p2);
char *p4 = malloc(6);
free(p3);
char *p5 = malloc(2);
free(p1);
➔ free(p4);
free(p5);
```



Heap: Dynamic Memory

```
#include <stdlib.h>
void *malloc(size_t size);
void free(void *ptr);
```

```
char *p1 = malloc(3);
char *p2 = malloc(1);
char *p3 = malloc(4);
free(p2);
char *p4 = malloc(6);
free(p3);
char *p5 = malloc(2);
free(p1);
free(p4);
➔ free(p5);
```



Avoid Leaking Memory

- Memory leaks “lose” references to dynamic memory

```
int f (void)
{
    char * p;
    p = (char *) malloc (8 * sizeof(char));
    ....
    return 0;
}

int main (void) {
    f ( );
    ...
}
```

Avoid Dangling Pointers

- Dangling pointers point to data that's not there anymore


```
char *f (void)
{
    char p[8];
    ....
    return p;
}
int main(void) {
    char *res = f ( );
}
```

Memory Allocation

- Memory is requested and granted in contiguous blocks
 - Think of memory as one huge array of bytes
 - *malloc* library call
 - used to allocate memory
 - finds sufficient contiguous memory
 - reserves that memory
 - returns the address of the first byte of the memory
 - *free* library call
 - give address of the first byte of memory to free
 - memory becomes available for reallocation
- Same for allocating and freeing memory for processes in OS

Example of Memory Allocation

```
char* ptr = malloc(4096);    // char* is address of a  
    single byte
```



Fragmentation



- Segments of memory can become unusable after allocation and de-allocation

- **External fragmentation**

- Variable allocation sizes
- Memory remains unallocated

- **Internal fragmentation**


- Fixed allocation sizes
- Memory is allocated but unused



External Fragmentation

- A series of *malloc* and *free*

```
char* first = malloc(100);  
char* second = malloc(100);  
char* third = malloc(100);  
char* fourth = malloc(100);  
free(second);  
free(fourth);  
char* problem = malloc(200);
```




- 250 bytes of free memory, only 150 contiguous
 - unable to satisfy final malloc request `malloc(200)`

Internal Fragmentation

- A series of *malloc*

- Assume allocation unit is 100 bytes

```
char* first = malloc(90);  
char* second = malloc(120);  
char* third = malloc(10);  
char* problem = malloc(50);
```




- All of memory has been allocated but only a fraction of it is used (220 bytes)
 - unable to handle final memory request `malloc(50)`

Internal vs. External Fragmentation

- Externally fragmented memory can be compacted (later)
- Fixed size allocation may lead to internal fragmentation, but less overhead to keep track of free memory
 - 8192 byte area of free memory
 - request for 8190 bytes
 - if exact size allocated: 2 bytes left to keep track of
 - if fixed size of 8192 used: 0 bytes to keep track of

Free List: External Fragmentation

- Need to keep track of available memory
 - Contiguous block of free mem is called a “hole”
 - Contiguous block of allocated mem is a “chunk”
- Keep a doubly linked list of free space
 - Build the pointers directly into the holes



Free List

- Prefer holes to be as large as possible
 - Large hole can satisfy a small request
 - The opposite is not true
 - Less overhead in tracking memory
 - Fewer holes; so faster search for available memory

Deallocating Memory

- When memory is freed
 - Place memory in free list; set next and previous pointers
 - Merge with hole before and/or after if possible



Cannot merge




Merge with hole before



Merge with hole after




Merge with two holes




Free List: Deallocating Memory

- Place memory in free list; set next and previous pointers
- Merge with hole before and/or after if possible




Allocation Algorithms: Best Fit

- Best Fit
 - Pick smallest hole that will satisfy the request
- Comments
 - Have to search entire list every time
 - Tends to leave lots of small holes
 - External fragmentation




Allocation Algorithms: Worst Fit

- Worst fit
 - Pick the largest hole to satisfy request
- Comments
 - Have to search entire list
 - Still leads to fragmentation issues



Allocation Algorithms: First Fit, Next Fit


- First fit
 - Pick the first hole large enough to satisfy the request
- Comments
 - Much faster than best or worst fit
 - Has fragmentation issues similar to best fit
- Next fit
 - Exactly like first fit except start search from where last search left off



Compaction

- To deal with internal fragmentation
 - Use paging or segmentation
 - More on this in next week's lecture
- To deal with external fragmentation
 - Can do compaction

Compaction




Compaction

- Simple concept
 - Move all allocated memory locations to one end
 - Combine all holes on the other end to form one large hole
- Major problems
 - Tremendous overhead to copy all data
 - Must find and change all pointer values
 - This can be very difficult


Multiple Free Lists

- Keep multiple lists of different hole sizes
- Take hole from a list that most closely matches size of request
- Leads to internal fragmentation
 - ~50% of memory in the common case



Multiple Free Lists

- Start out with single large hole
 - One entry in one list; Hole size is usually a power of 2
- Upon a request for memory, keep dividing hole by 2 until appropriate size memory is reached
 - At every division, a new hole is added to a different free list
- Once a hole is created, it cannot merge with another hole




Multiple Free Lists

```
char* ptr = malloc(100 K)
```

- Initially, only one entry in first list (1 M)
- In the end, one entry in each list except 1 M

list (1 M)



list (512 K)



list (256 K)




list (128 K)



allocate this hole

Buddy System

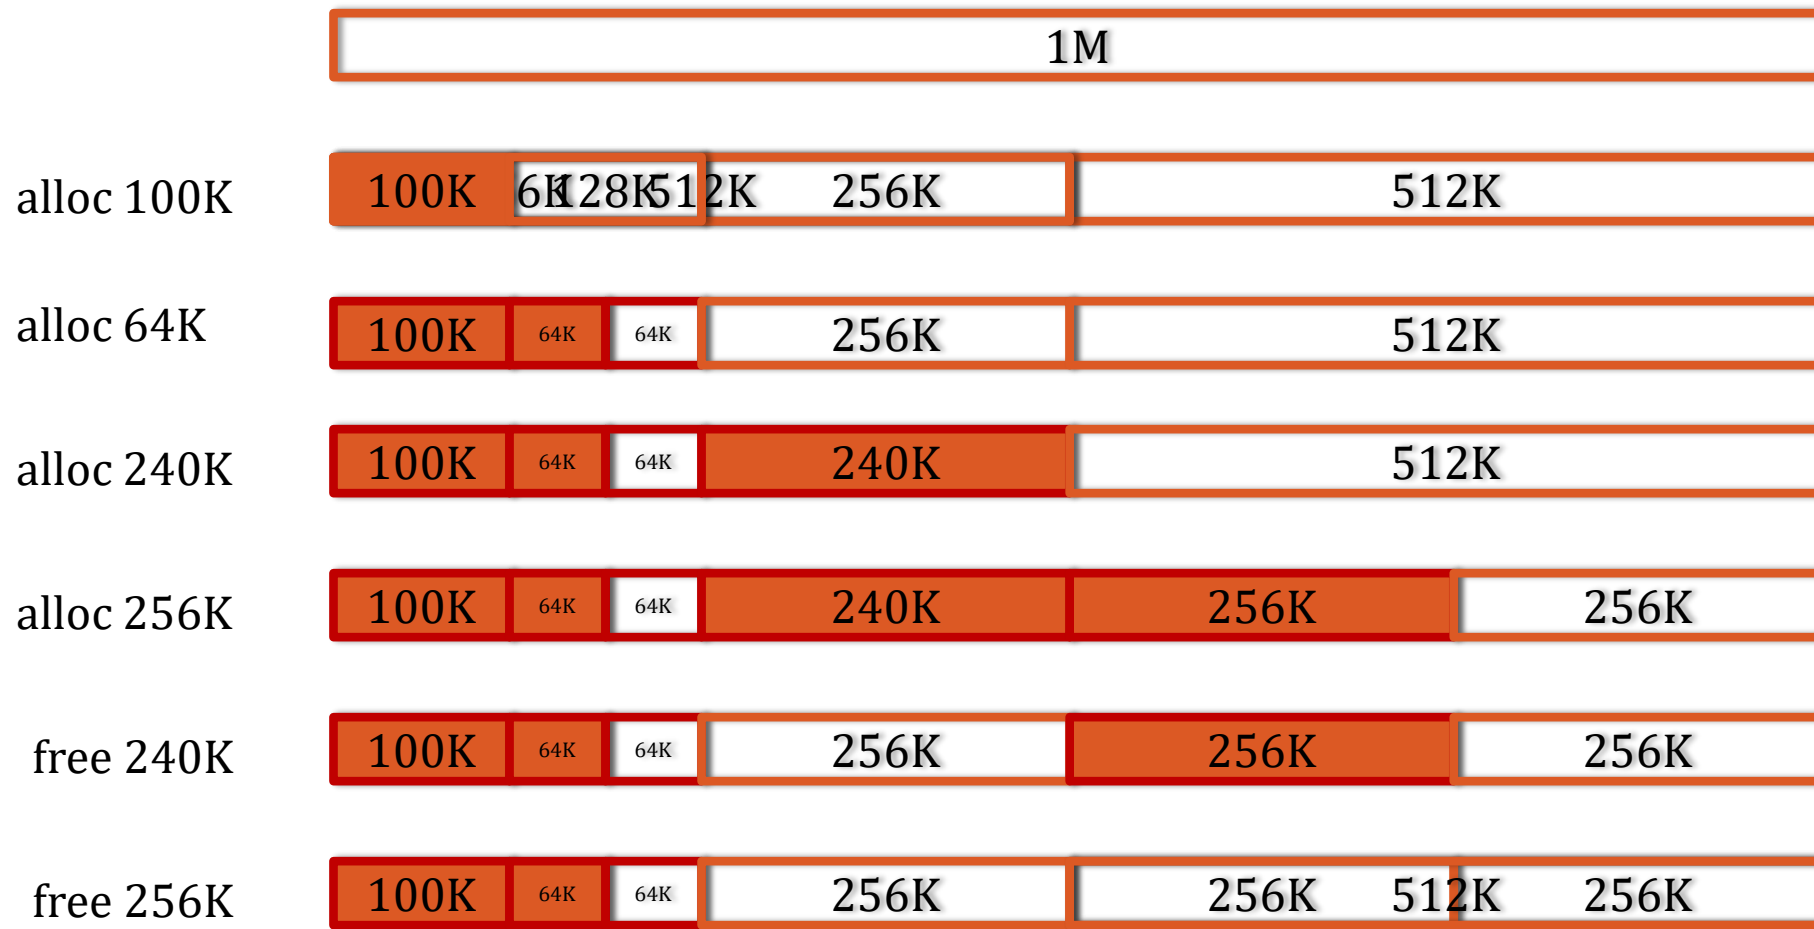
- Each hole in a free list has a *buddy*
 - If a hole and its buddy are combined, they create a new hole
 - New hole is twice as big and is aligned on proper boundary
- Example
 - A hole is of size 4
 - Starting location of each hole: 0, 4, 8, 12, 16, 20, ...
 - Buddies are the following: (0,4), (8, 12), ...
 - If buddies are combined, get holes of size 8
 - starting location of these holes: 0, 8, 16, ...



Buddy System

- When allocating memory
 - If list is empty, go up one level, take a hole and break it in 2
 - These 2 new holes are buddies
 - Now give one of these holes to the user
- When freeing memory
 - If chunk just returned and its buddy are in the free list, merge them and move the new hole up one level

Buddy System Example



Summary

- In this lecture, we learnt about different components of process memory
- We also learnt simple memory allocation and deallocation schemes