

# Virtual Memory & DMA

Ravi Suppiah  
Lecturer, NUS SoC

# Learning Objectives

- Motivation for virtual memory from  
Technology perspective  
Software perspective
- Memory Hierarchy:  
Cache, Main Memory, Disk  
Spatial and Temporal Locality
- Virtual Memory  
Virtual address space  
Physical address space  
Mapping between virtual and physical memory  
Virtual/Physical Page  
Page Table  
Page Fault  
TLB

# Memory Technology: 1950s



1948: Maurice Wilkes examining  
EDSAC's delay line memory tubes  
16-tubes each storing 32 17-bit words






Figure 10. IBM 2361 Core Storage

1952: IBM 2361 16KB magnetic core memory



Maurice Wilkes: 2005

# Memory Technology today: DRAM




Infineon Technologies: stores data equivalent to 640 books, 32,000 standard newspaper pages, and 1,600 still pictures or 64 hours of sound

# Recent DRAM chip

- 128GB DRAM chip from Samsung




# DRAM Capacity Growth



4X capacity increase almost every 3 years, i.e., 60% increase per year for 20 years

# DRAM Bandwidth Growth



DRAM Bandwidth grows by 1.3x per year



# Issue with DRAM

- 1.6x capacity growth of DRAM per year
- 1.3x bandwidth growth of DRAM per year
- DRAM latency improvement is far behind



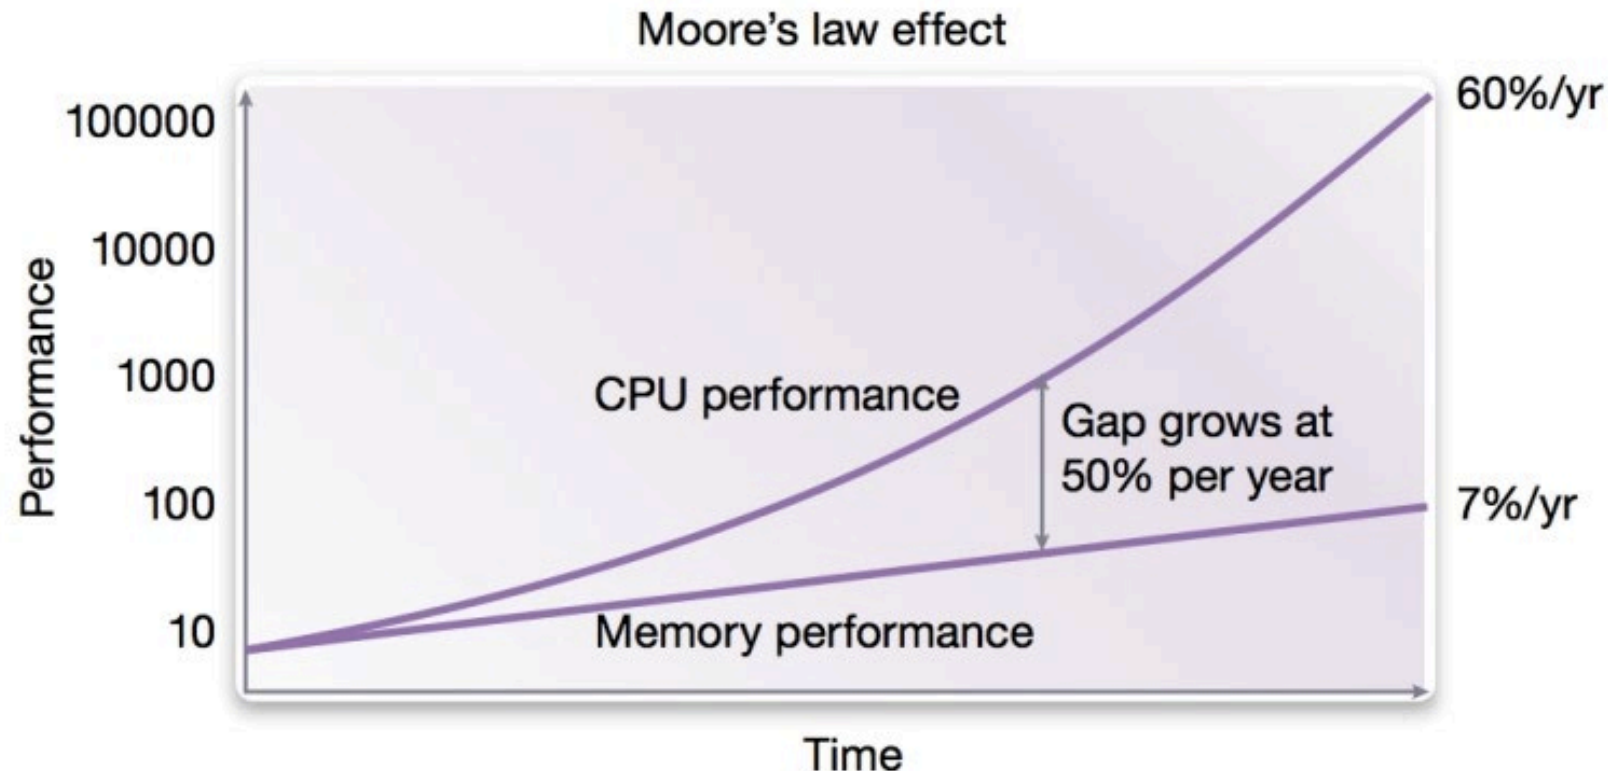
# Processor-DRAM Performance Gap

Memory Wall:

DRAM access latency: 50 ~ 70 ns

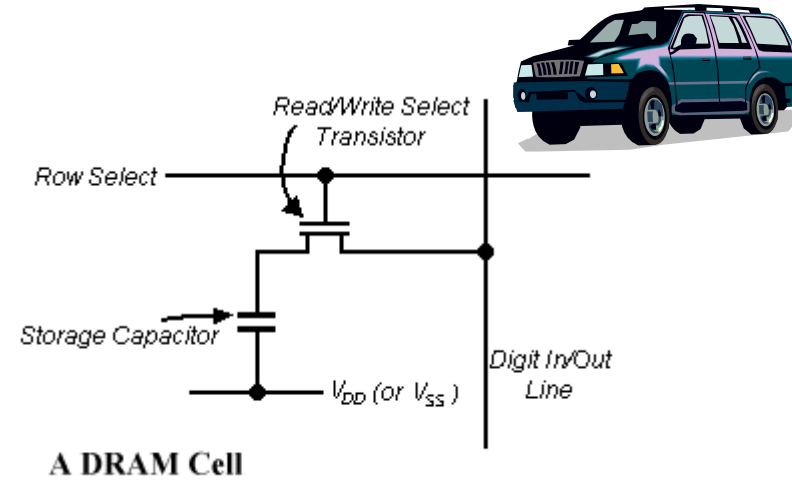
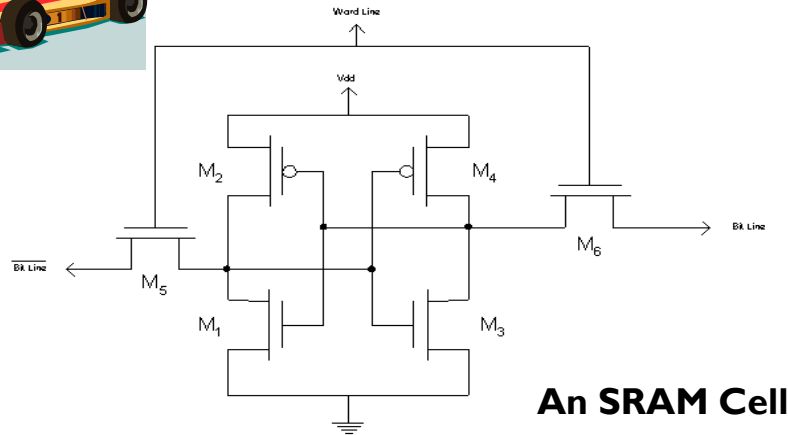
1 GHz Processor → 1 ns per clock cycle but 50 ns to access DRAM

50 processor clock cycles per memory access !!



Do we have any faster memory technology?

# Fast Memory Technologies: SRAM



## SRAM

6 transistor per memory cell

→ Low density

Fast access latency of 0.5 – 5 ns

## DRAM

1 transistor per memory cell


→ High density

Slow access latency of 50-70ns

# Slow Memory Technologies: Magnetic Disk



Drive Physical and Logical Organization




High-end hard disk:

Average latency: 5-20ms

Capacity: Seagate 4TB



# Quality vs. Quantity

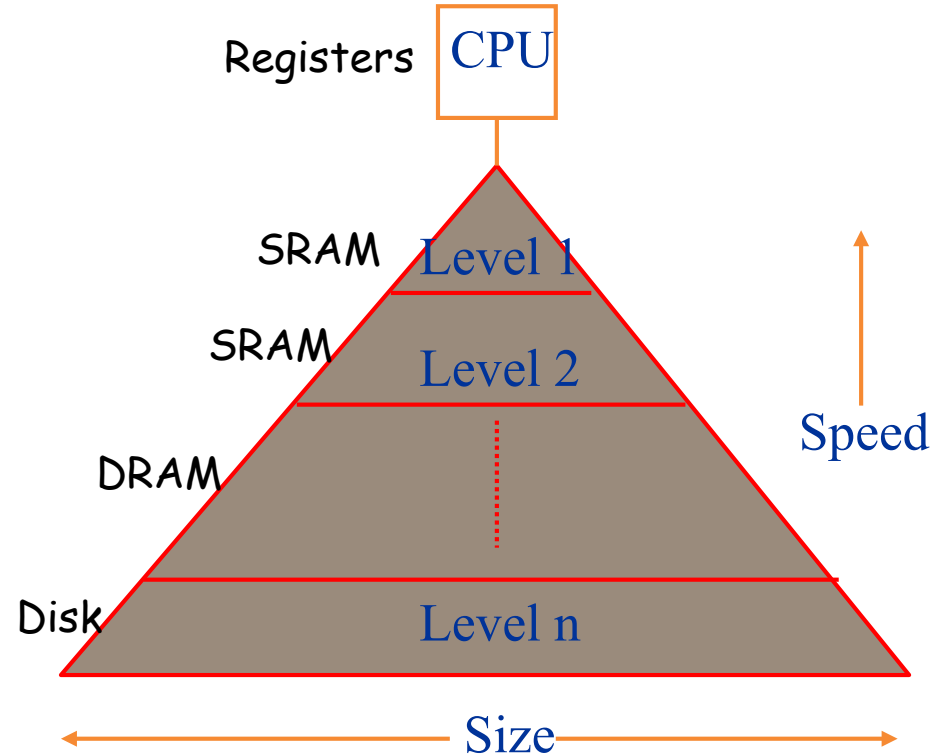


	Capacity	Latency	Cost/GB
Register	100s Bytes	20 ps	\$\$\$\$\$
SRAM	100s KB	0.5-5 ns	\$\$\$
DRAM	10s GB	50-70 ns	\$
Hard Disk	100s GB	5-20 ms	Cents
<b>Ideal</b>	<b>1 GB</b>	<b>1 ns</b>	<b>Cheap</b>

# Best of Both Worlds

- What we want: A BIG and FAST memory
- Memory system should perform like 1 GB of SRAM (1 ns access time) but cost like 1 GB of slow memory
- Key concept: Use a hierarchy of memory technologies
  - Small but fast memory near CPU
  - Large but slow memory farther away from CPU

# Memory Hierarchy




# The Basic Idea

- Keep the frequently and recently used data in smaller but faster memory
- Refer to bigger and slower memory only when you cannot find data/instruction in the faster memory
- Why does it work? Principle of Locality

Program accesses only a small portion of the memory address space within a small time interval

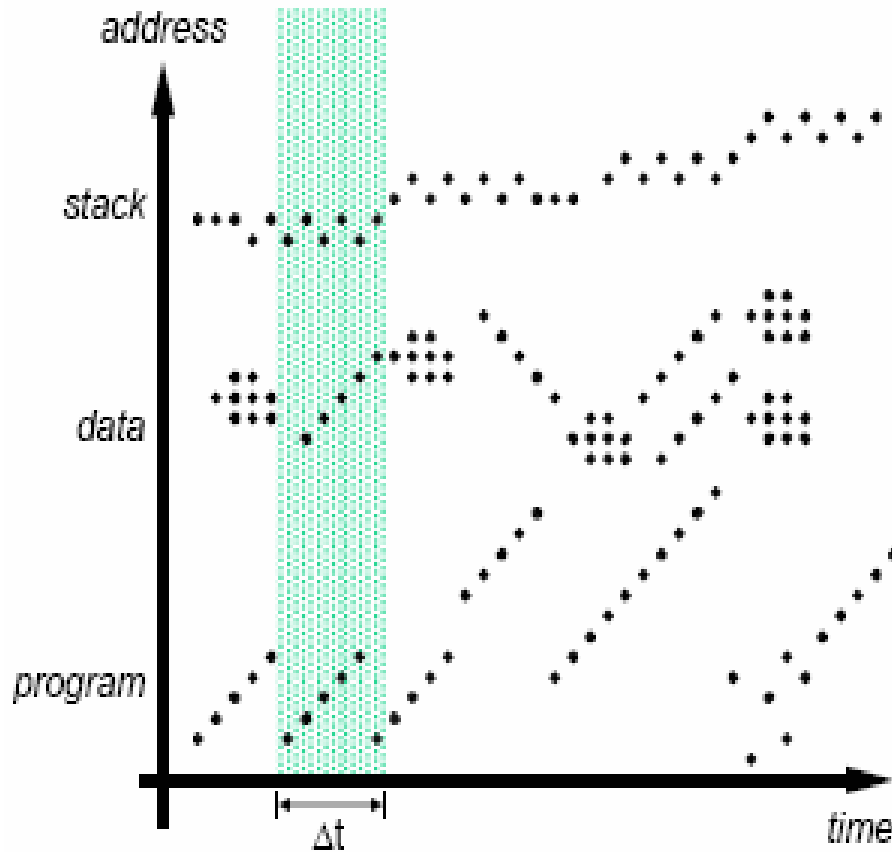


# Locality



- Temporal locality
  - If an item is referenced, it will tend to be referenced again soon
- Spatial locality
  - If an item is referenced, nearby items will tend to be referenced soon
- Locality for instruction
- Locality for data

# Working Set




- Set of locations accessed during  $\Delta t$
- Different phases of execution may use different working sets
- Want to capture the working set and keep it in the memory closest to CPU

# Exploiting Memory Hierarchy (1/2)

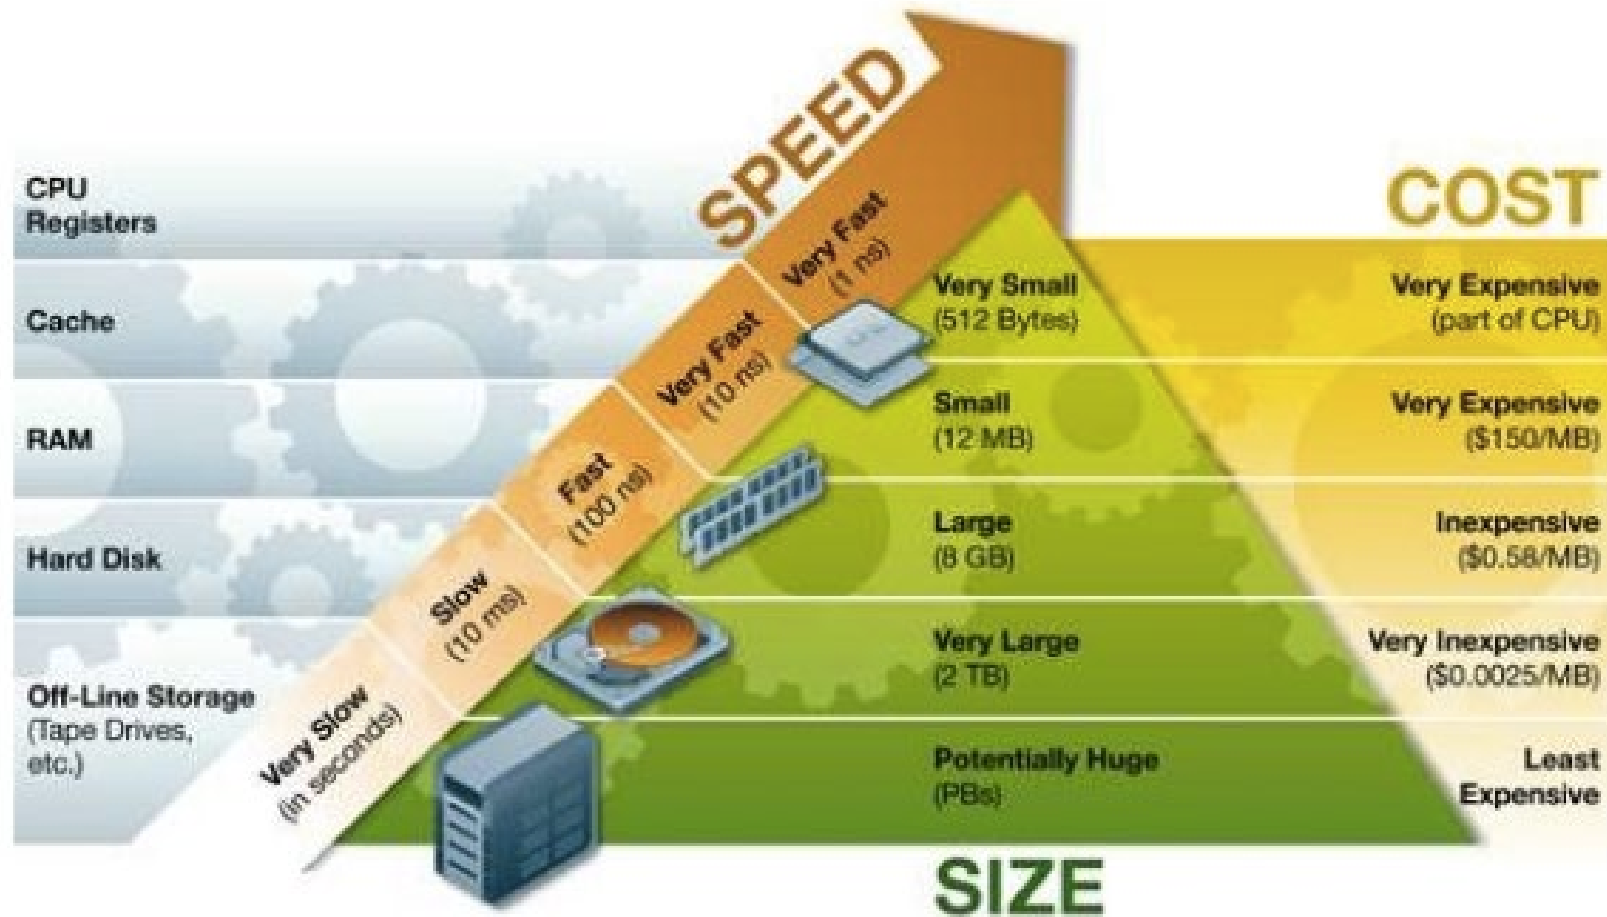
- **Visible to Programmer**
  - Various storage alternatives, e.g., register, main memory, hard disk
  - Tell programmer to use them cleverly
- **Transparent to Programmer (except for performance)**
  - Single address space for programmer
  - Processor automatically assigns locations to fast or slow memory depending on usage patterns

# Exploiting Memory Hierarchy (2/2)

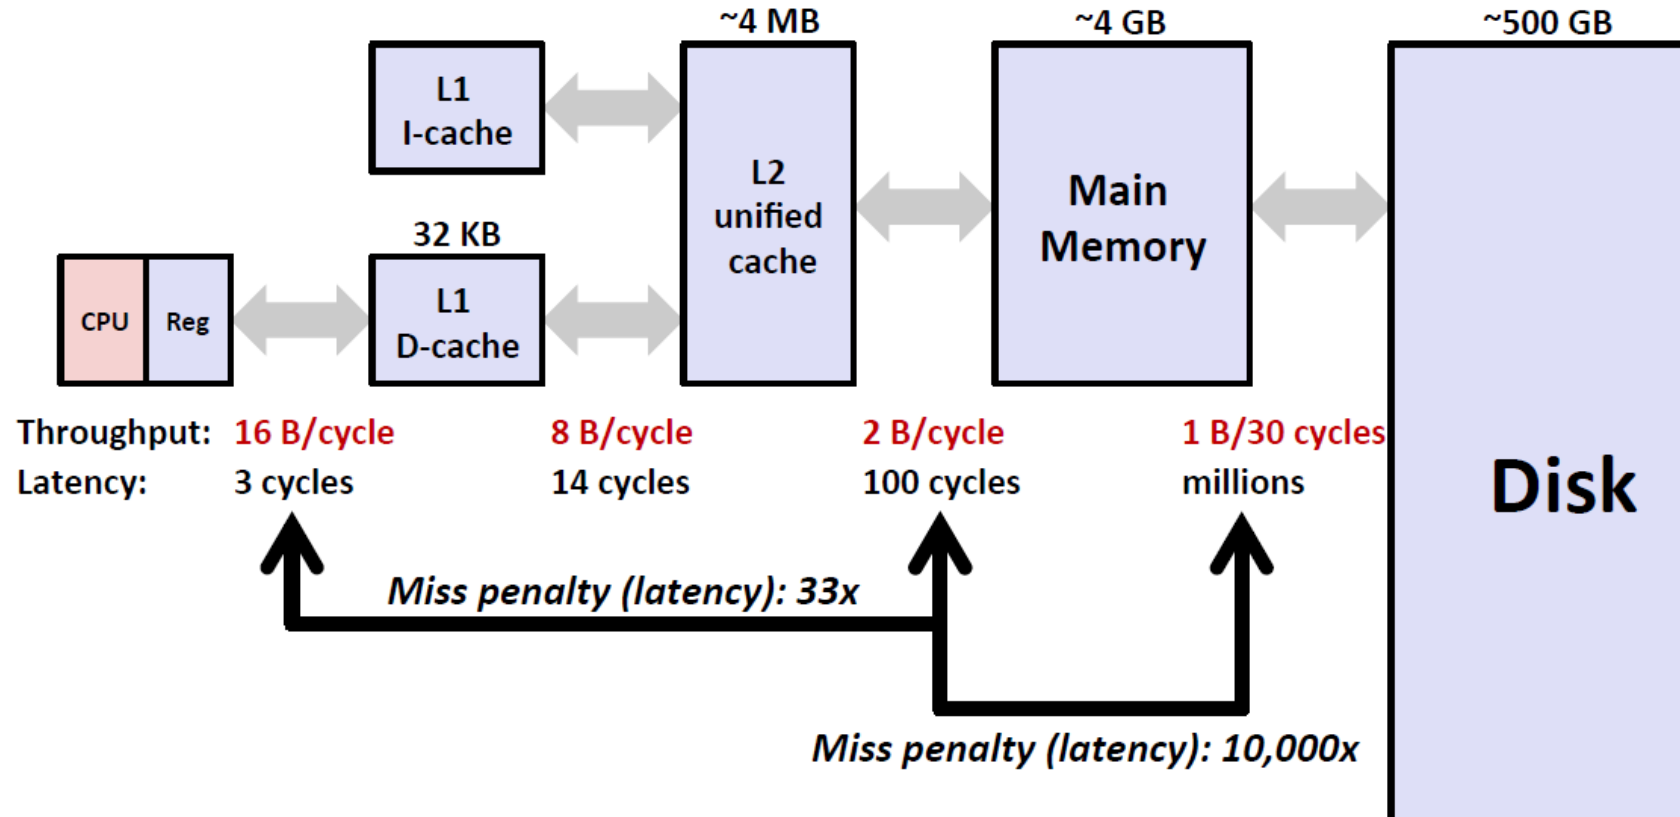


- How to make **SLOW** main memory appear faster?
- **Cache** – a small but fast **SRAM** near **CPU**
  - Often in the same chip as CPU
  - Introduced in 1960; almost every processor uses it today
  - Hardware managed: Transparent to programmer
- How to make **SMALL** main memory appear bigger than it is?  
Virtual memory
  - OS managed: Again transparent to programmer

# Memory Hierarchy Summary



# Memory Hierarchy: Core 2 Duo



# Virtual Memory Motivation Re-iterated

- Hard disk is cheap
- Make a portion of the hard disk *look like* memory to the CPU, so that we can fool the CPU into thinking there's more memory that we actually paid for!
- We can do this by having instructions and data that the CPU is currently interested in, in main memory. Everything else stays on disk
- In this way we can squeeze **MUCH MORE** instructions and data than otherwise possible!




# Processes

- **A process is an instance of a running program**
- **Process provides each program with two key abstraction**
  - Independent control flow
  - Each process seems to have exclusive use of the CPU
  - Private virtual address space
  - Each process seems to have exclusive use of main memory
- **How are these illusions maintained?**
  - Process executions are interleaved through scheduling
  - Address space managed by virtual memory system


# Process Memory

- **Text:** Code
- **Data:** Initialized global and static variables
- **BSS:** Uninitialized global and static variables
- **Heap:** Dynamic Memory
- **Stack:** Local variables



# Virtual Memory


- **Program refers to virtual memory addresses**
  - Conceptually memory is just a very large array of bytes
  - Each byte has its own address
  - System provides address space “private” to process
- **Allocation: Compiler and run-time system**
  - Where different program objects should be stored
  - All allocation within single virtual address space
- **What problems does virtual memory solve?**



# Problem 1: How does Everything Fit?

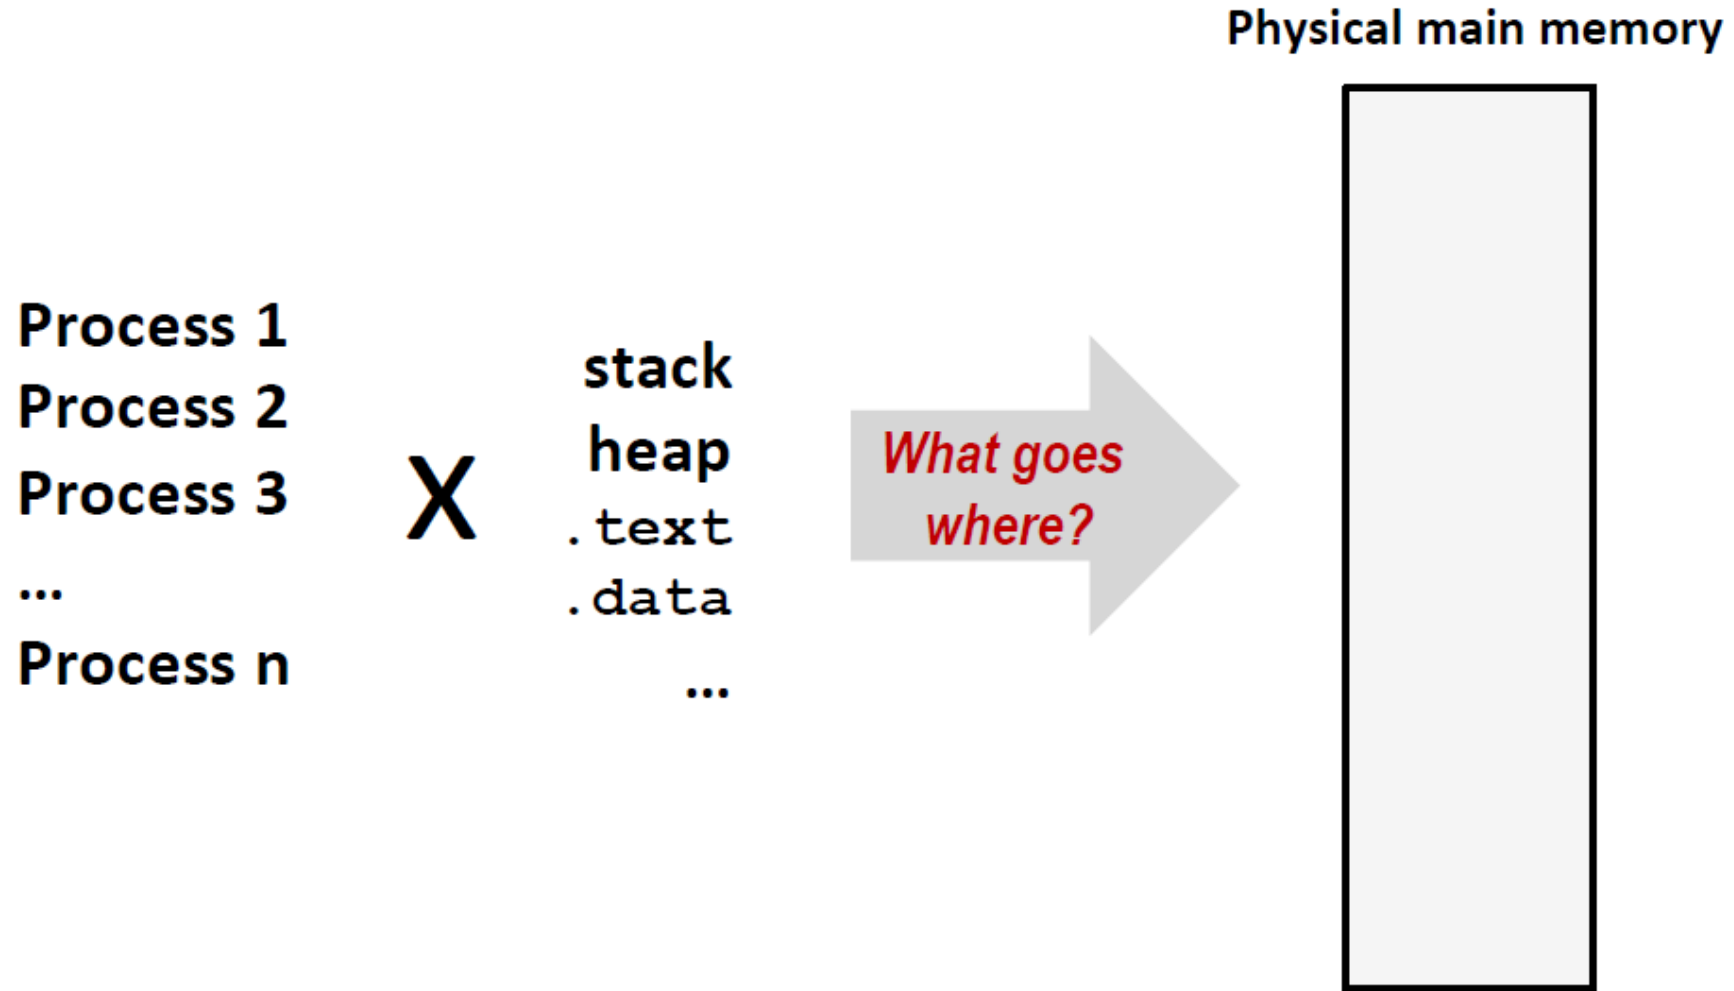
**64-bit addresses:  
16 Exabyte**

**Physical main memory:  
Few Gigabytes**




**And there are many processes ....**


# Problem 2: Memory Management



# Problem 3: How to Protect?



# Problem 4: How to Share?






# How would you solve those problems?

- Fitting a huge memory into a tiny physical memory
- Managing the memory space of multiple processes
- Protecting processes from stepping on each other's memory
- Allowing processes to share common parts of memory


# Indirection

- Any problem in Computer Science can be solved by adding another level of indirection
- The ability to reference something using a name, reference, or container instead of the value itself. A flexible mapping between a name and a thing allows changing the thing without notifying the holder of the name



# Virtual Memory Solution: Indirection

- Each process gets its own private virtual address space
- Virtual addresses are mapped to physical memory

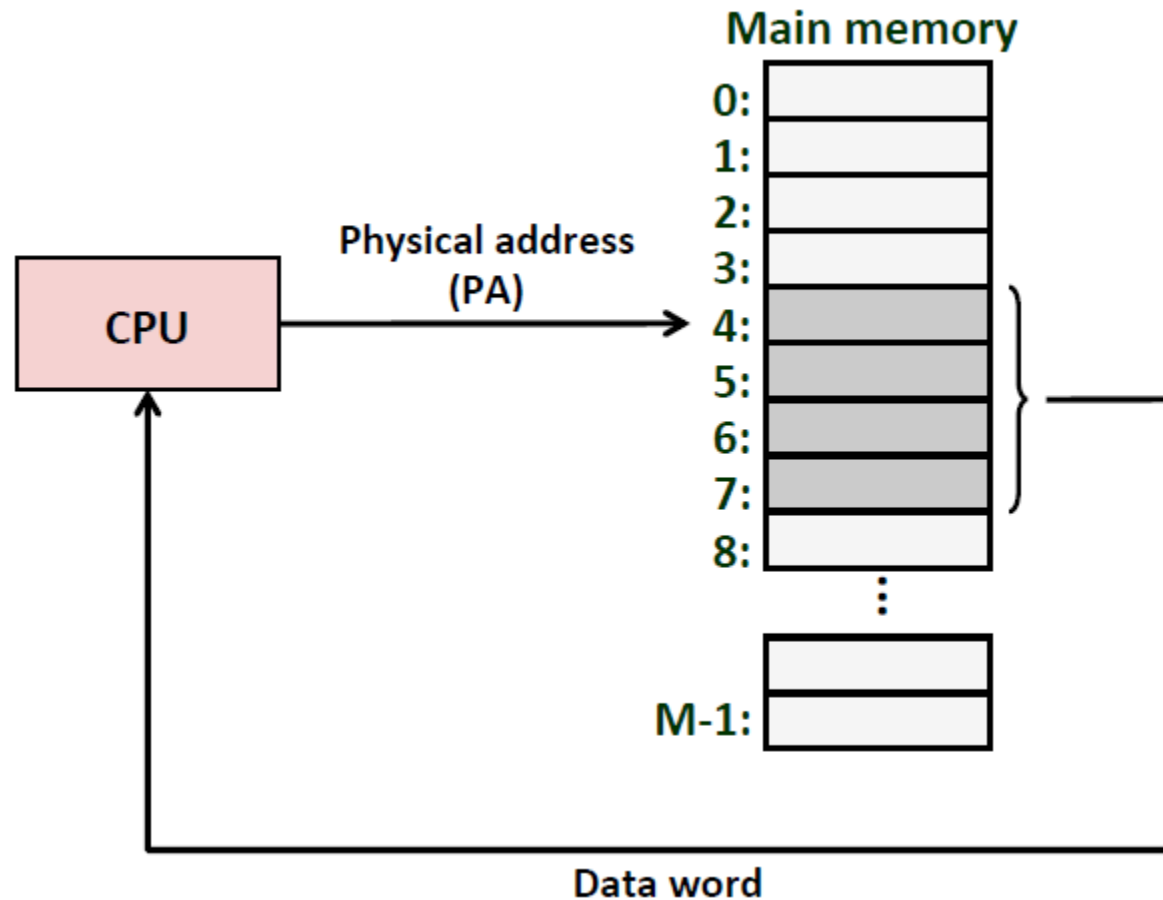


# Address Spaces

- **Virtual Address Space:** Set of  $N = 2^n$  virtual addresses  
 $\{0, 1, 2, 3, \dots, N-1\}$
- **Physical Address Space:** Set of  $M = 2^m$  physical addresses ( $n > m$ )  
 $\{0, 1, 2, 3, \dots, M-1\}$
- **Each byte in physical main memory:**
  - One physical address
  - Zero, one, or more virtual addresses

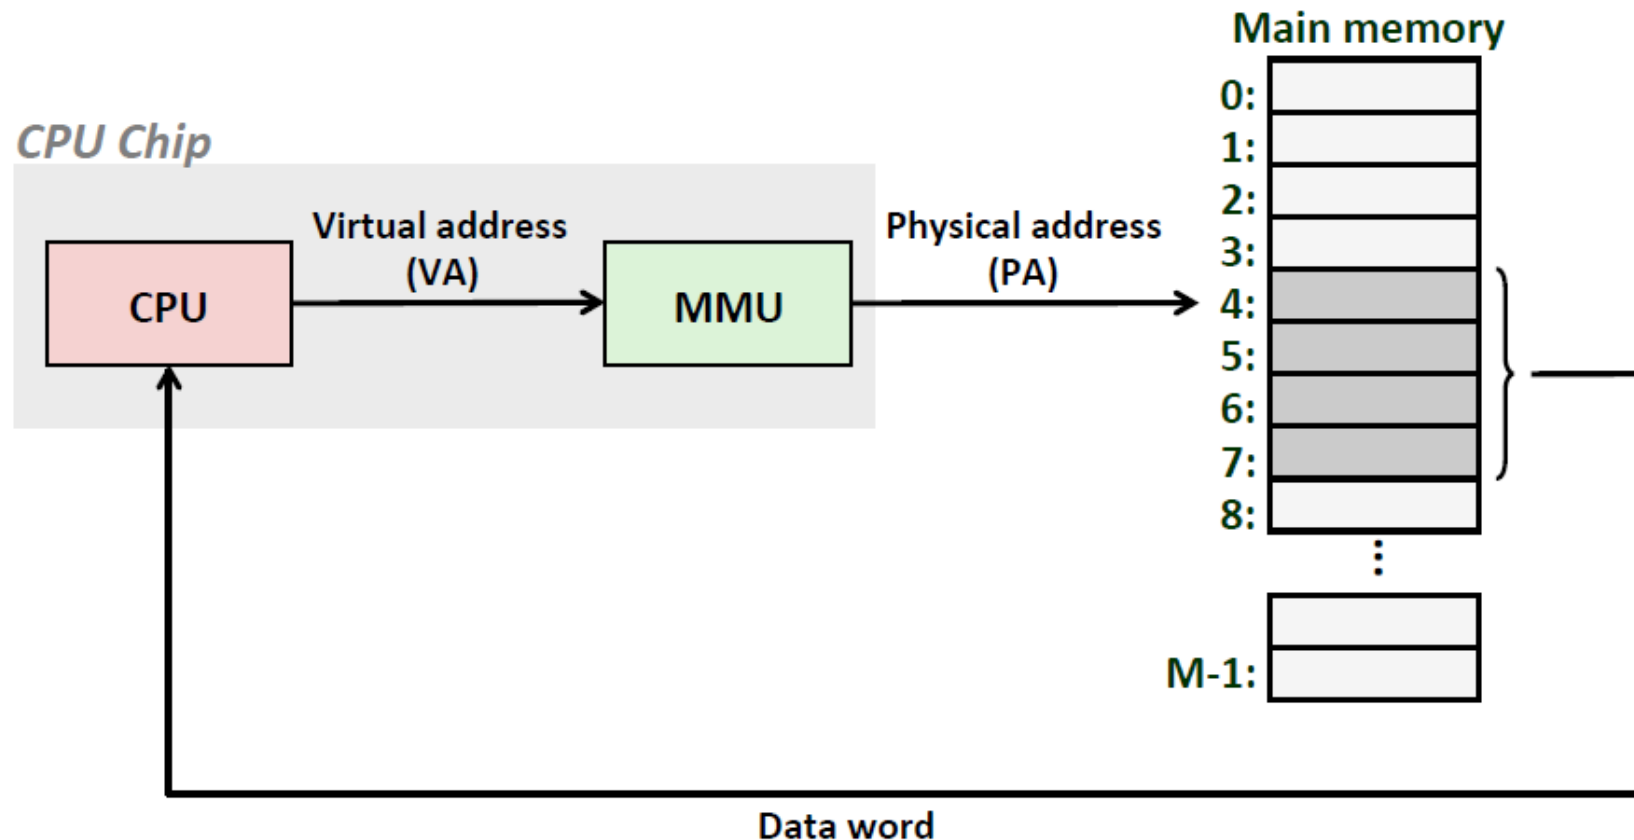
# A System using Physical Addressing

- Uses in simple systems like embedded micro-controllers



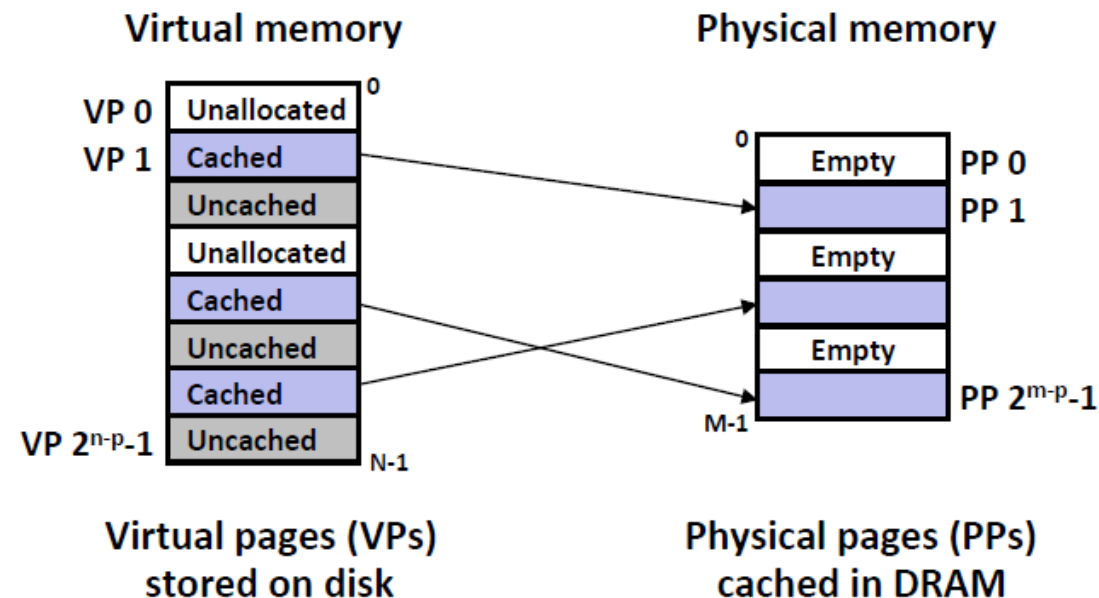
# A System using Virtual Addressing

- Used in all modern desktops, laptops, servers, smartphones
- One of the great ideas in Computer Science
- 



# Virtual Memory Pages

- Virtual memory can be thought of as an array of  $N = 2^n$  contiguous bytes stored **on a disk**
- Virtual memory is partitioned into number of blocks called **pages** (size of page  $P = 2^p$  bytes)
- Then physical main memory (**DRAM**) is used as a **cache** (temporary store) for some pages of the virtual memory array

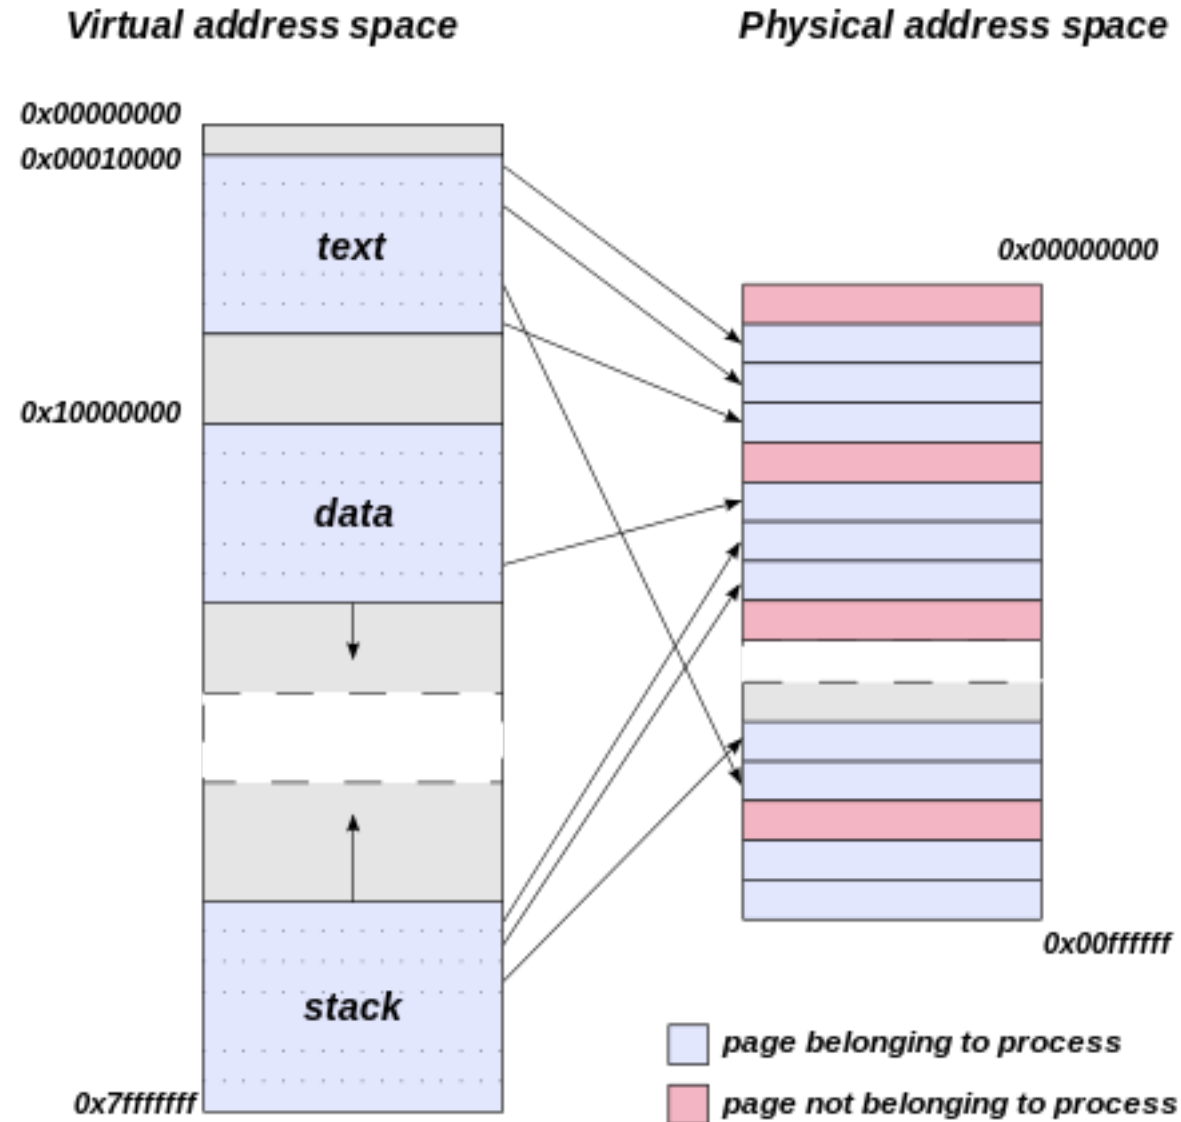




# Status of a Virtual Memory Page


- **Unallocated:** the page is not being used by the process (e.g., the pages between heap and stack)
- **Uncached:** the page is being used by the process but it is not cached in DRAM
- **Cached:** the page is being used by the process and it is cached in DRAM

# Virtual & Physical Memory Pages



# Virtual Address Mapping

- A virtual address can be mapped to either physical memory or disk



# Virtual Memory Example (1/3)



- **16KB virtual address space: 14-bit virtual address**
- **4KB physical address space: 12-bit physical address**
- **Page size = 64 bytes**
  
- **How many virtual pages?  $16\text{KB}/64\text{B} = 256$**
  
- **How many physical pages?  $4\text{KB}/64\text{B} = 64$**

# Virtual Memory Example (2/3)

- **16KB virtual address space: 14-bit virtual address**
- **4KB physical address space: 12-bit physical address**
- **Page size = 64 bytes**
  
- **How many virtual pages?  $16\text{KB}/64\text{B} = 256$**
- **How many bits for virtual page ID? 256 pages =  $2^8$  require 8 bits**
  
- **How many physical pages?  $4\text{KB}/64\text{B} = 64$**
- **How many bits for physical page ID? 64 pages =  $2^6$  require 6 bits**
  
- **How many bits to access an address within a page (page offset)?  
 $64\text{ bytes} = 2^6$  require 6 bits**

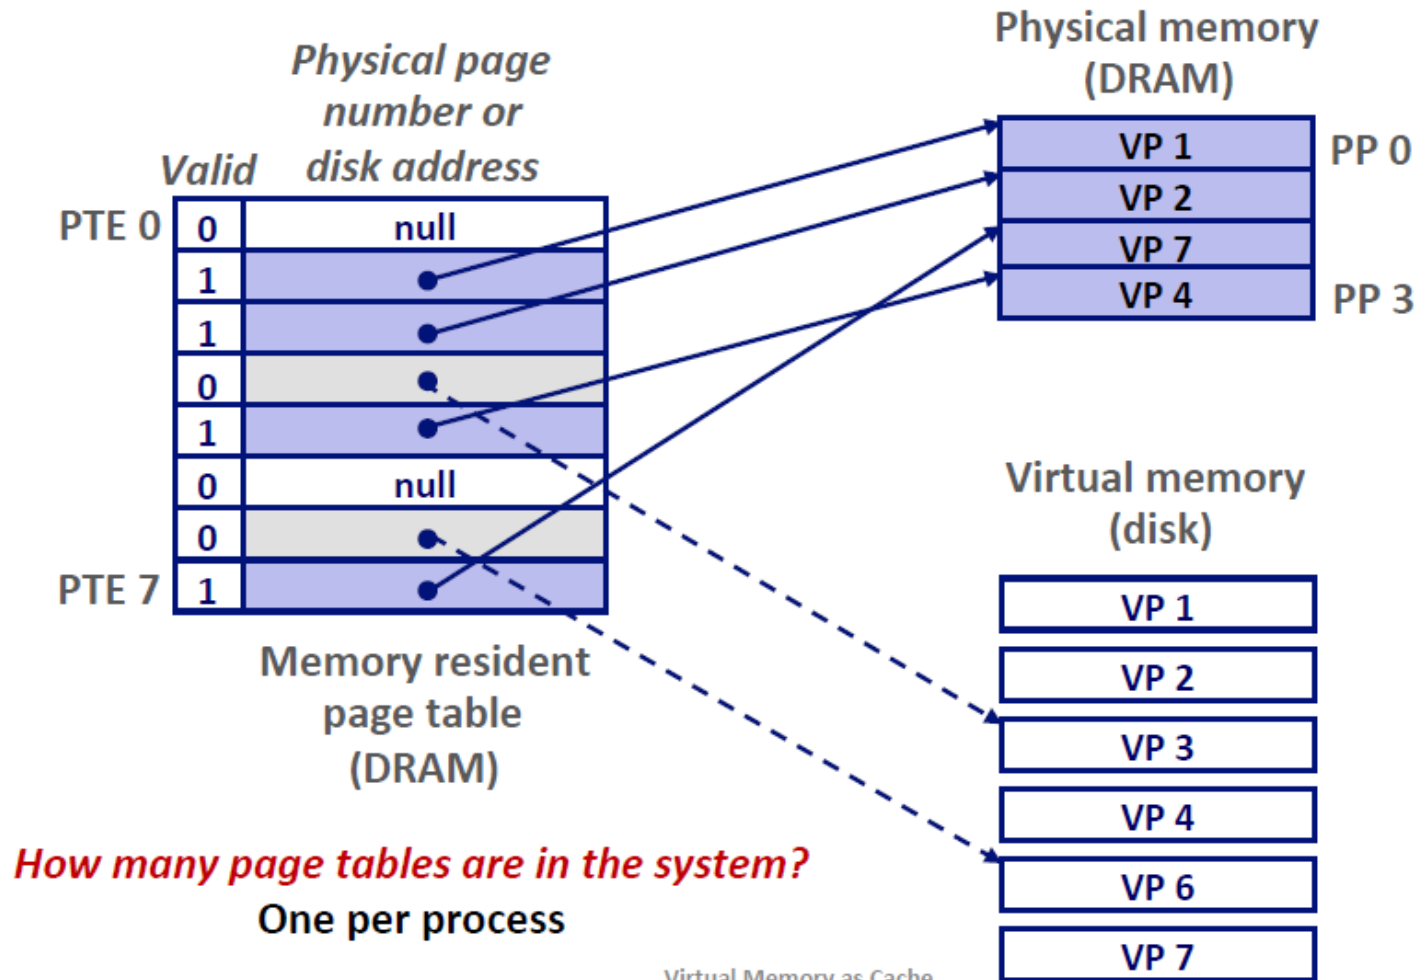
# Virtual Memory Example (3/3)

- 14-bit virtual addresses
- 12-bit physical addresses
- Page size = 64 bytes



# Address Translation: Page Tables

- A **page table** is an array of page table entries (PTE) that maps virtual pages to physical pages



# More about Page Table

- **One page table per process**
- **Page tables are stored in main memory (DRAM)**
- **Page table should have one page table entry (PTE) corresponding to each virtual page**
  - 256 virtual pages will need 256 PTE
- **Each PTE stores the status of the corresponding virtual page**
  - 0 implies uncached and 1 implies cached
- **If a virtual page is cached in DRAM, then PTE contains the physical page number**



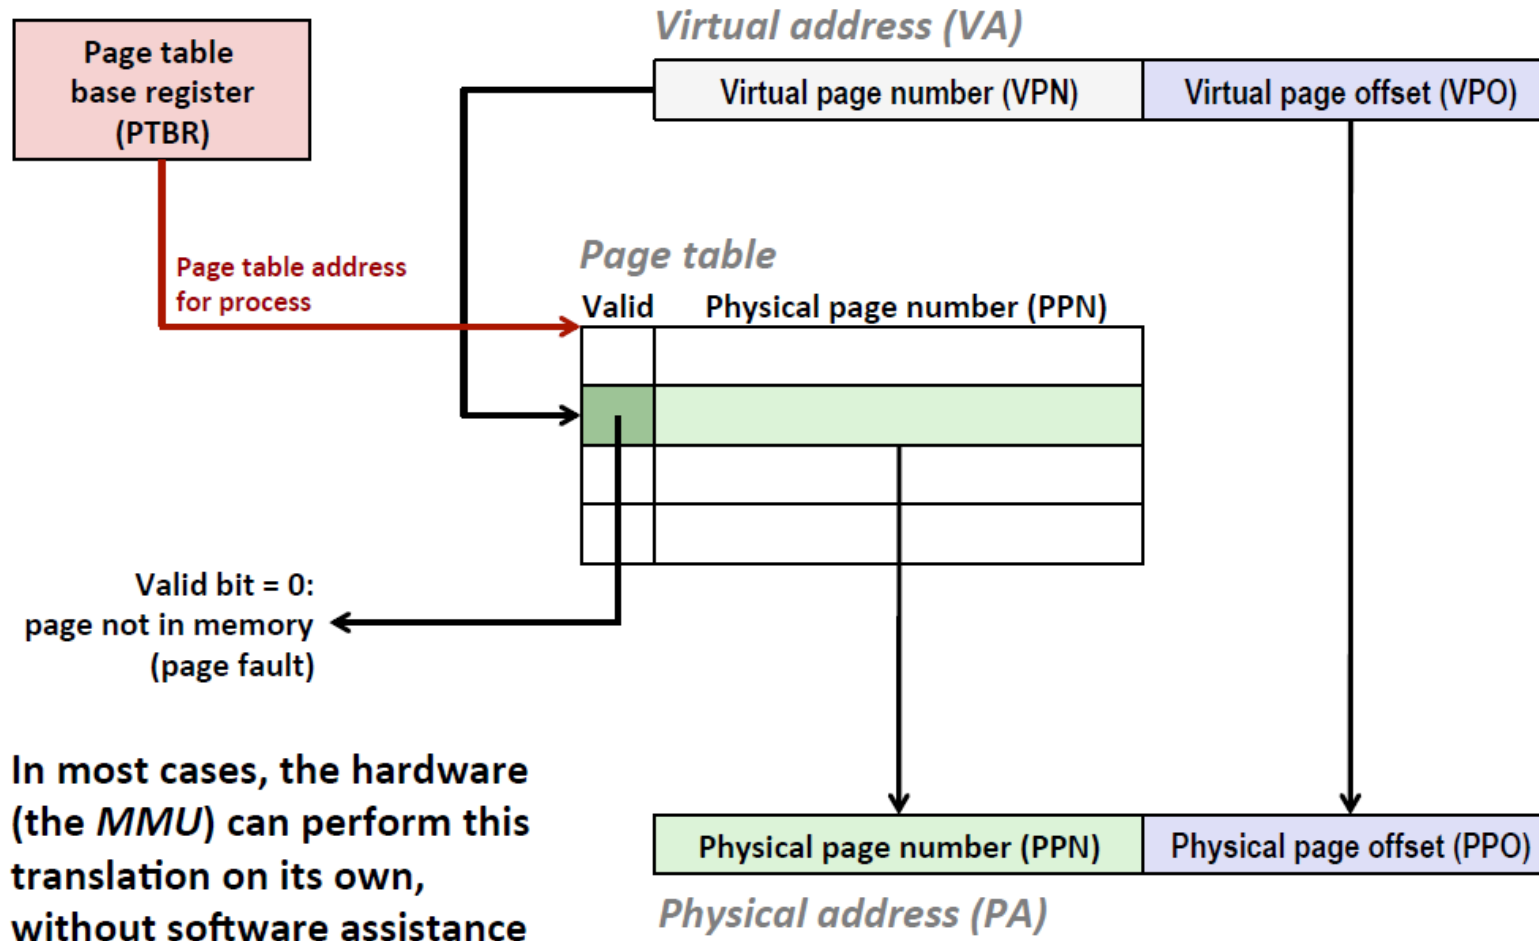
# Page Table Example

	Physical Page Number (PP)	Valid Bit
0	0x10	1
1	DISK	0
2	0x23	1
3	0x15	1
4	DISK	0
5	0x1B	1
6	0x2E	1
...	...	...
255	0x2A	1

# Page Table Size

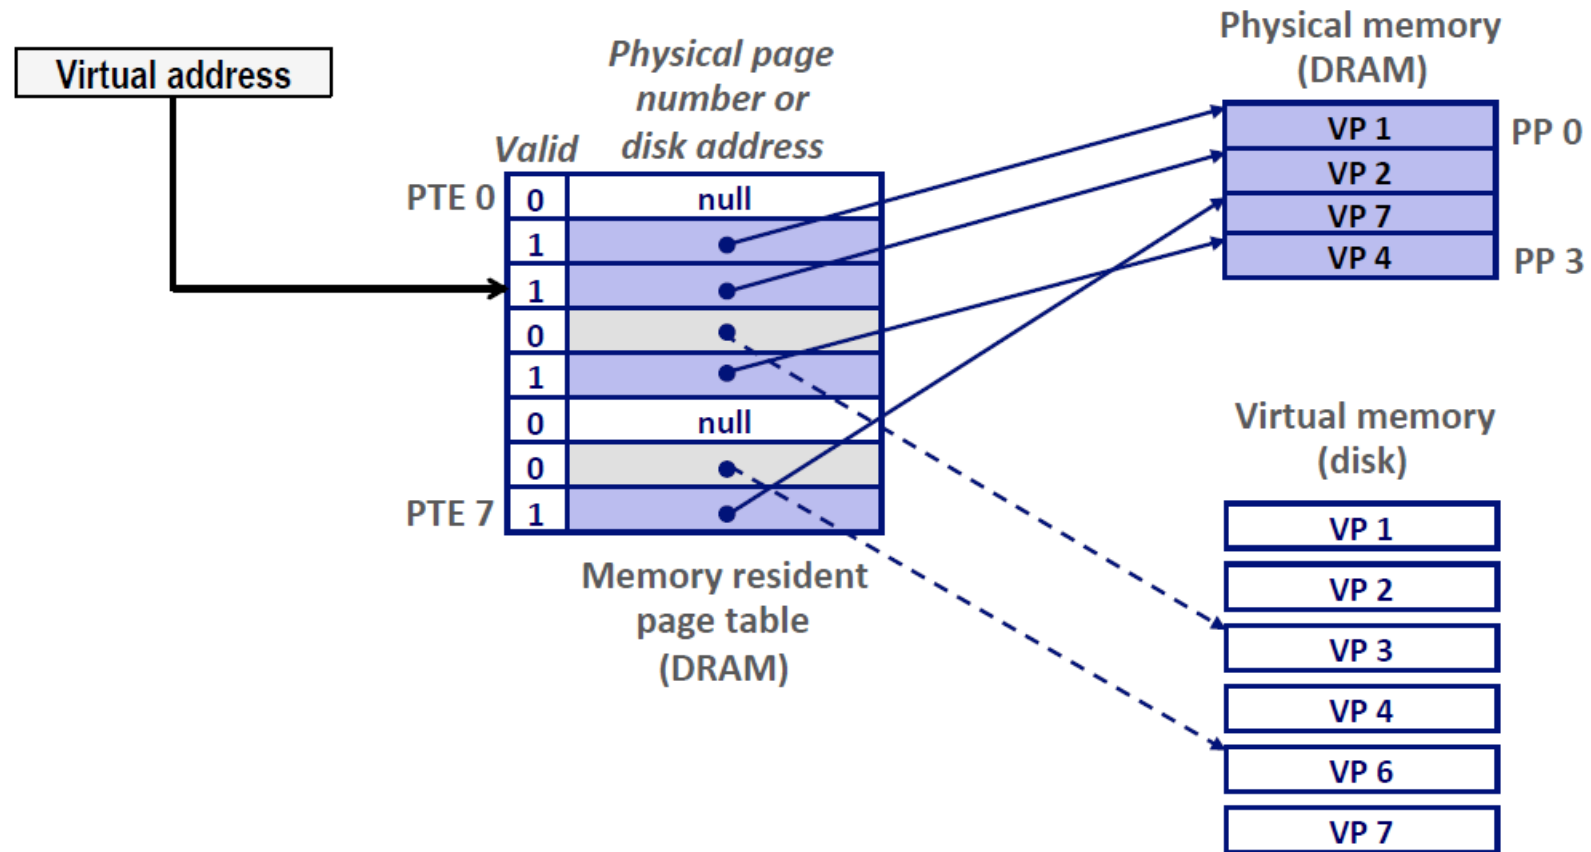
- **16KB virtual address space: 14-bit virtual address**
- **4KB physical address space: 12-bit physical address**
- **Page size = 64 bytes**
  
- **How many virtual pages?  $16\text{KB}/64\text{B} = 256$**
- **How many physical pages?  $4\text{KB}/64\text{B} = 64$**
- **How many bits for physical page ID? 64 pages =  $2^6$  require 6 bits**
  
- **Each page table entry: 1 valid bit + 6 physical page no bits = 7 bits**
- **256 virtual pages → 256 page table entries**
- **Total page table size =  $256 \times 7 \text{ bits} = 1,792 \text{ bits} = 224 \text{ bytes}$**

# Address Translation with a Page Table



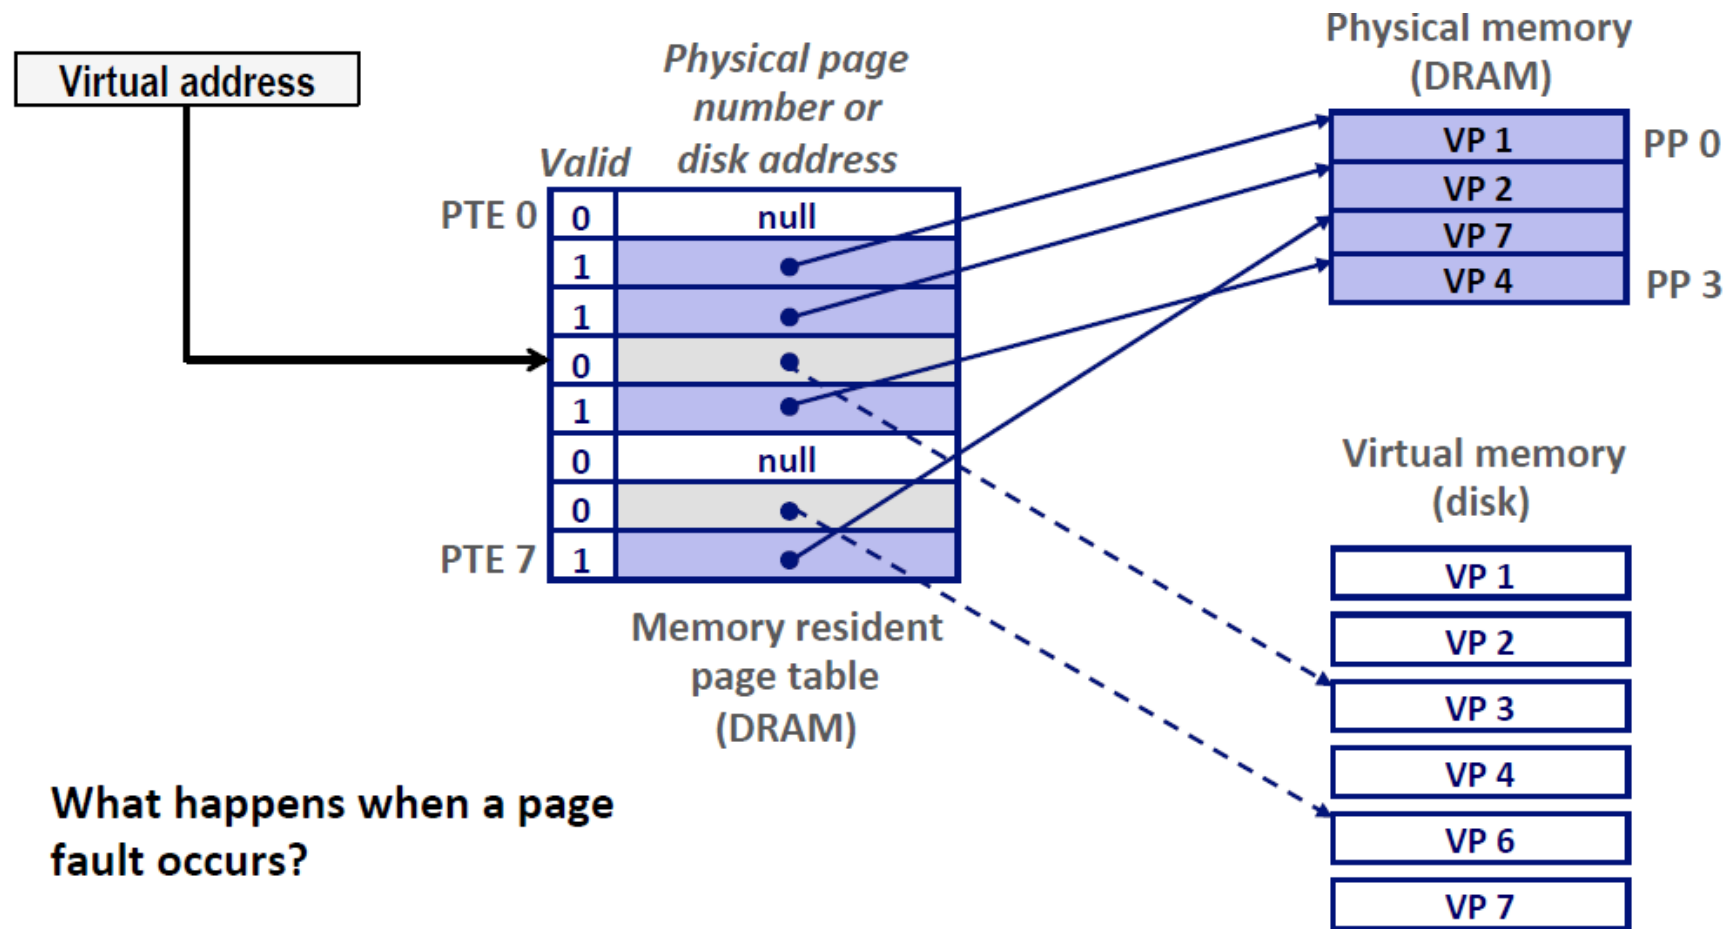
# Page Hit

- Reference to virtual memory byte that is in physical memory



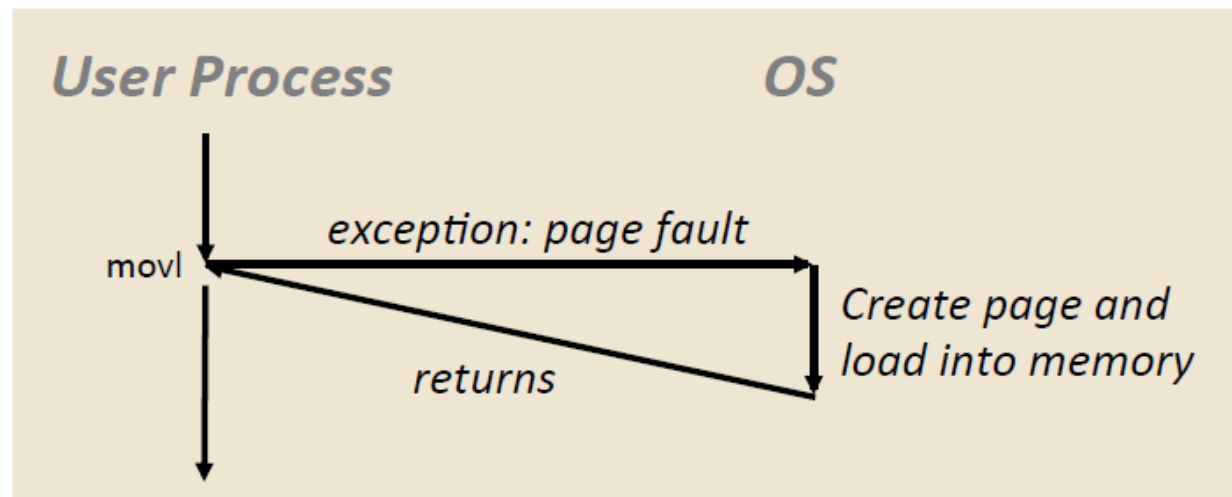
# Page Fault

- Reference to virtual memory byte that is **NOT** in physical memory



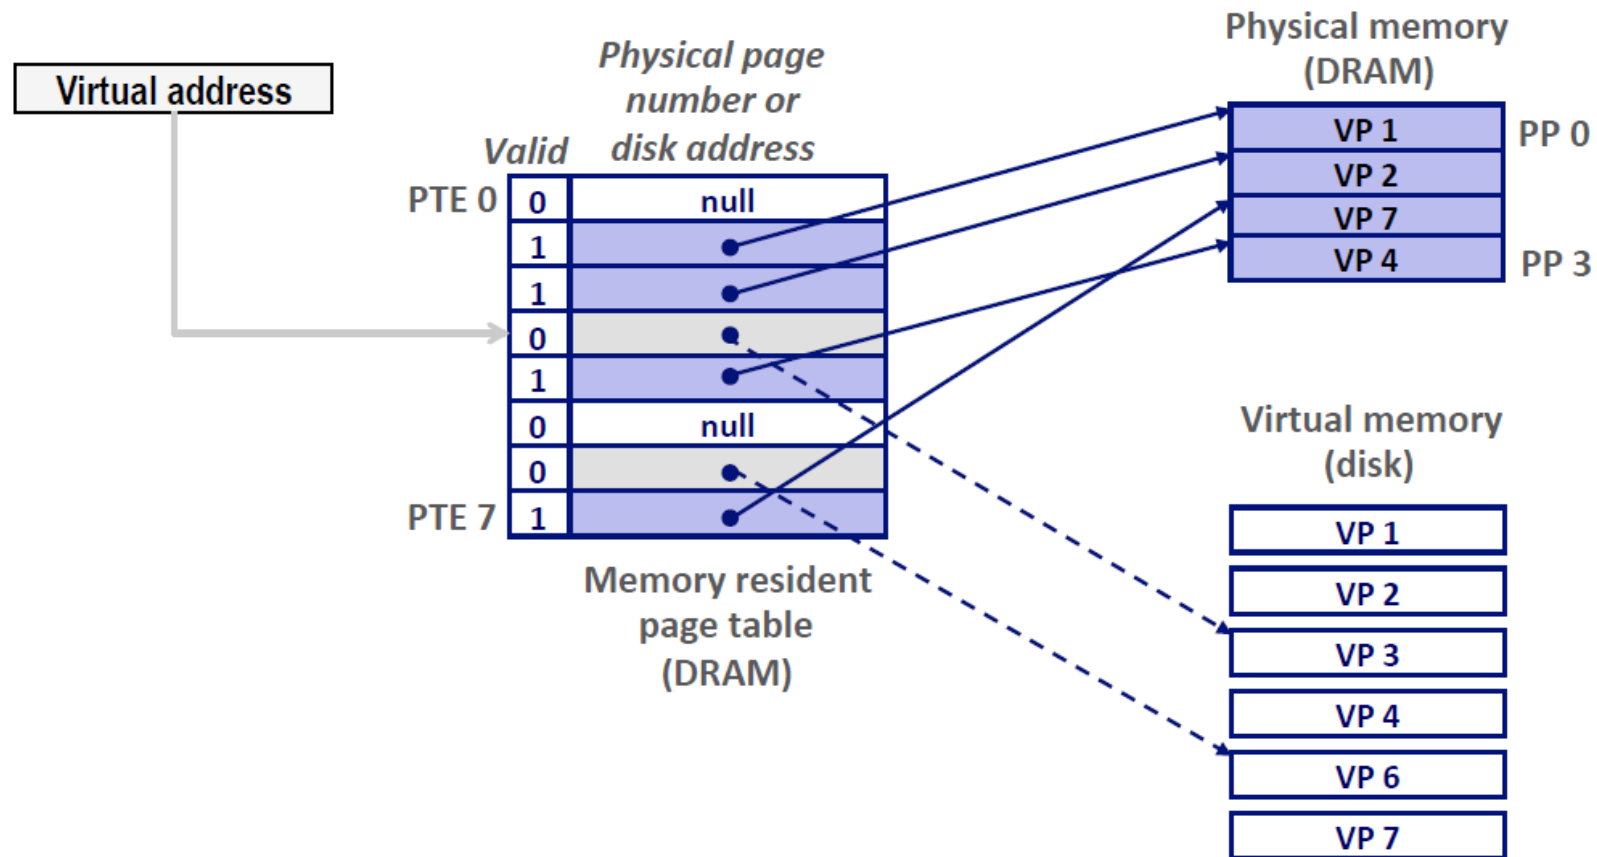
# Handling Page Fault

- **User accesses a memory location**
- **That portion (page) of user's memory is currently on disk**
- **Page fault (page miss) raises an exception (similar to interrupt)**
- **OS page handler must load page from disk to physical memory**
- **Modify page table entry corresponding to the virtual page**
- **Returns to faulting instruction**
- **Successful memory access on second try**



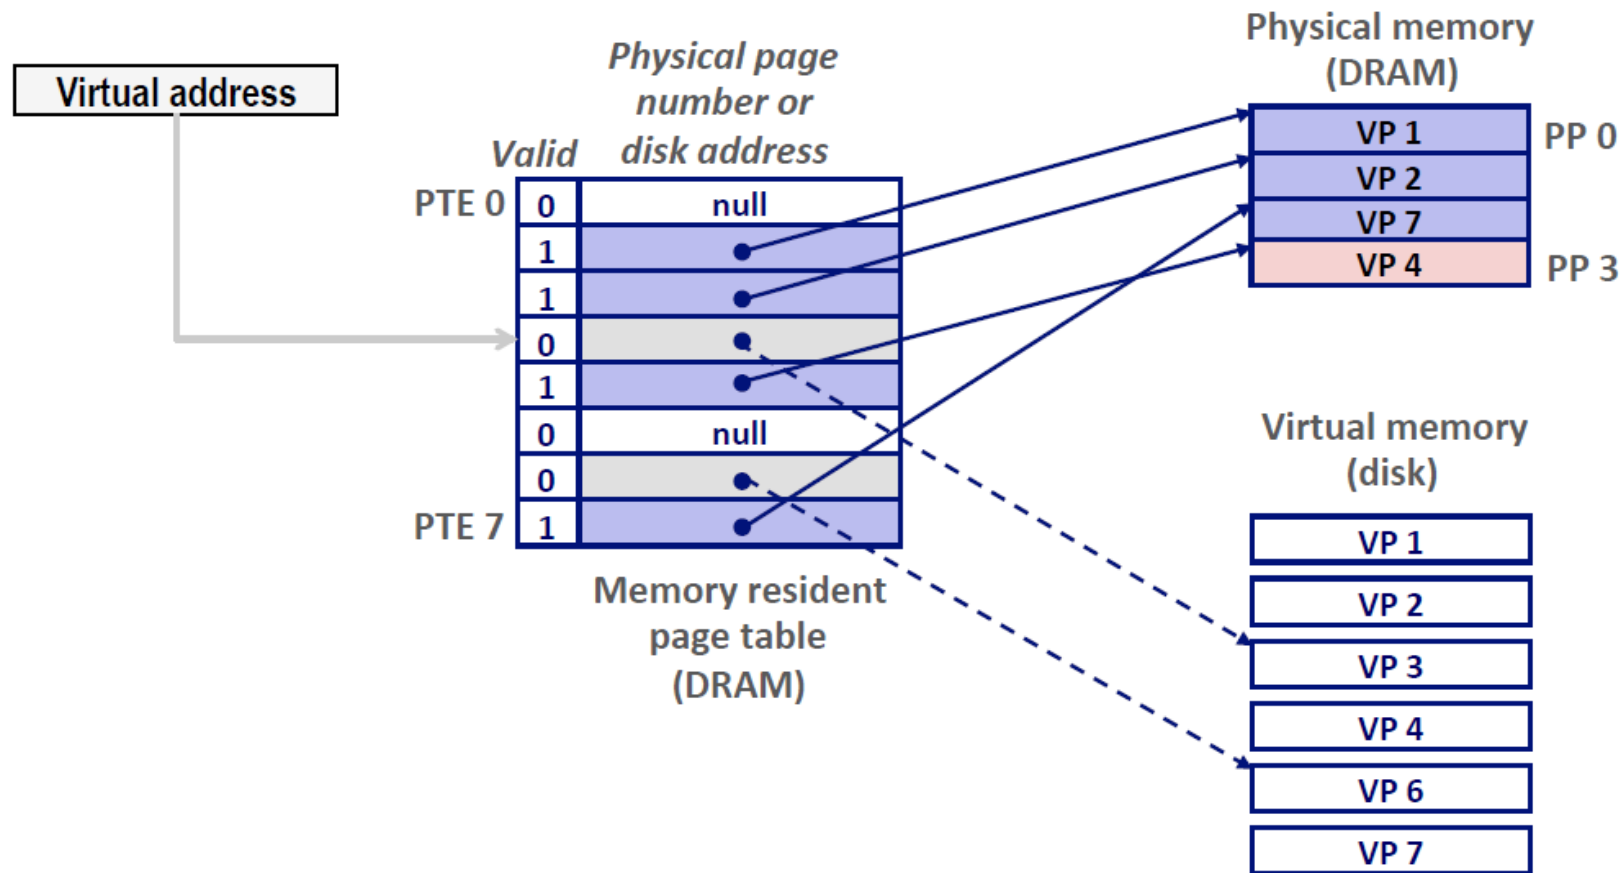
# Page Fault Example

- Page fault causes an exception for **Virtual Page Number (VP) 3**



# Page Fault Example

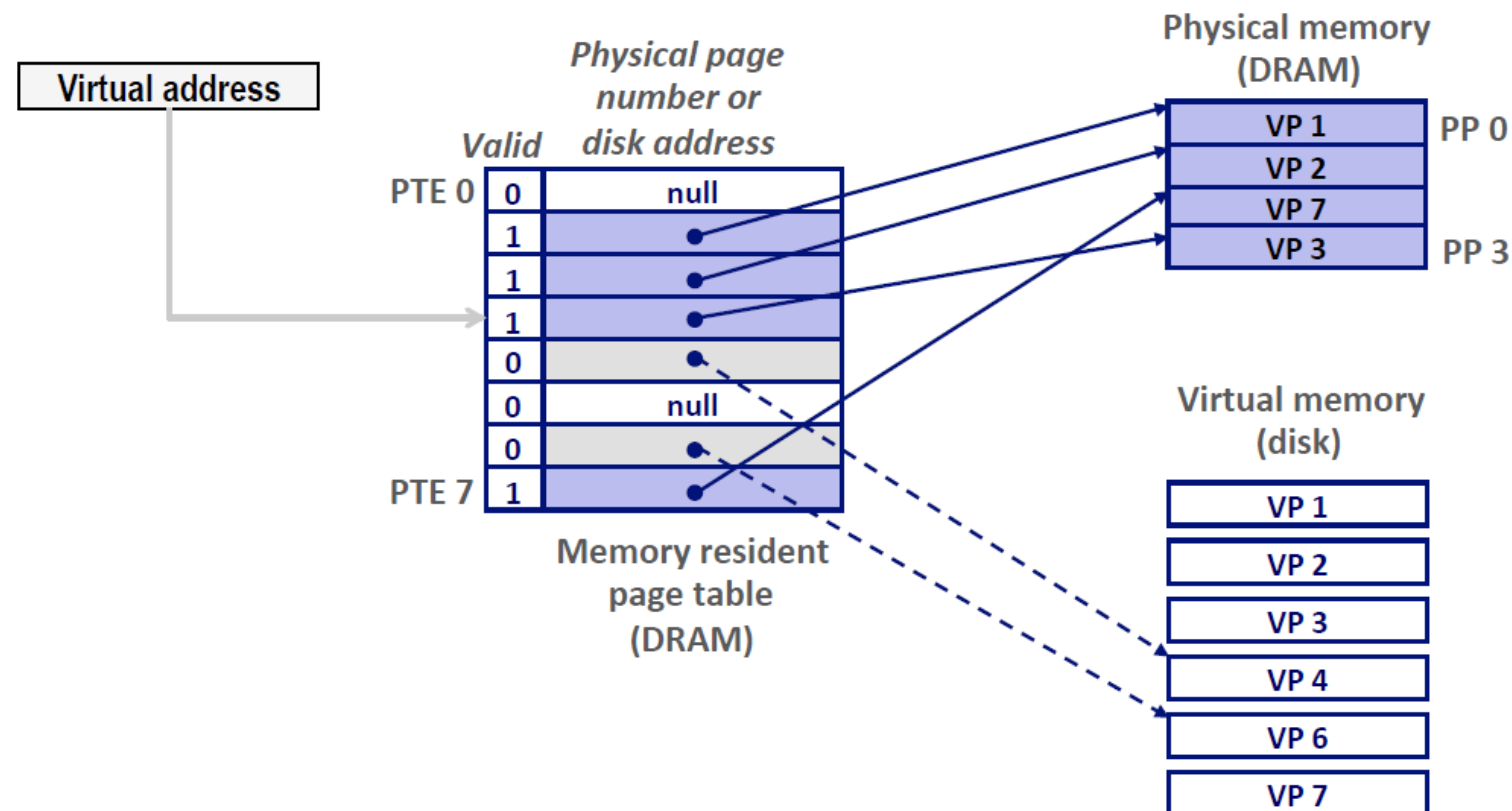
- Page fault causes an exception
- OS page fault handler selects a victim to be evicted (here VP 4)





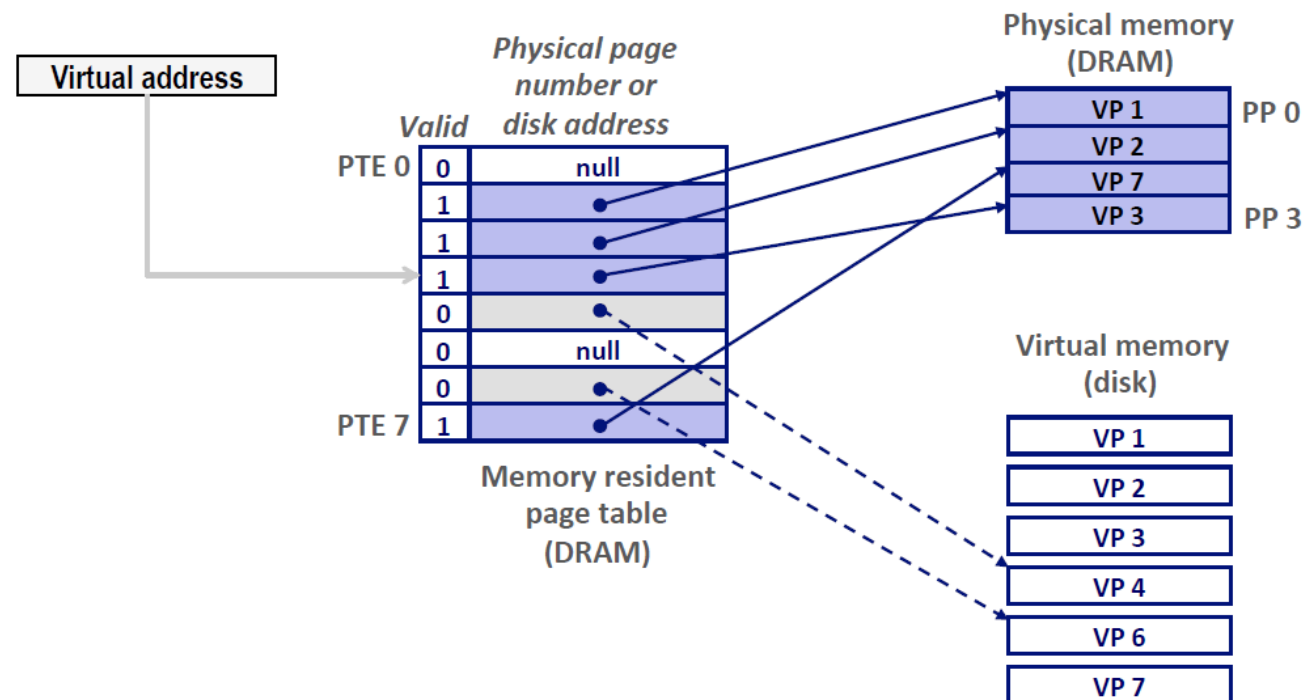
# Page Fault Example

- Page fault causes an exception
- OS page fault handler selects a victim to be evicted (here VP 4)
- OS reads in VP3 from disk to memory in physical memory PP3 and modifies page table



# Page Fault Example

- Page fault causes an exception
- OS page fault handler selects a victim to be evicted (here VP 4)
- OS reads in VP3 from disk to memory in physical memory PP3 and modifies page table
- Offending instruction is restarted: page hit!

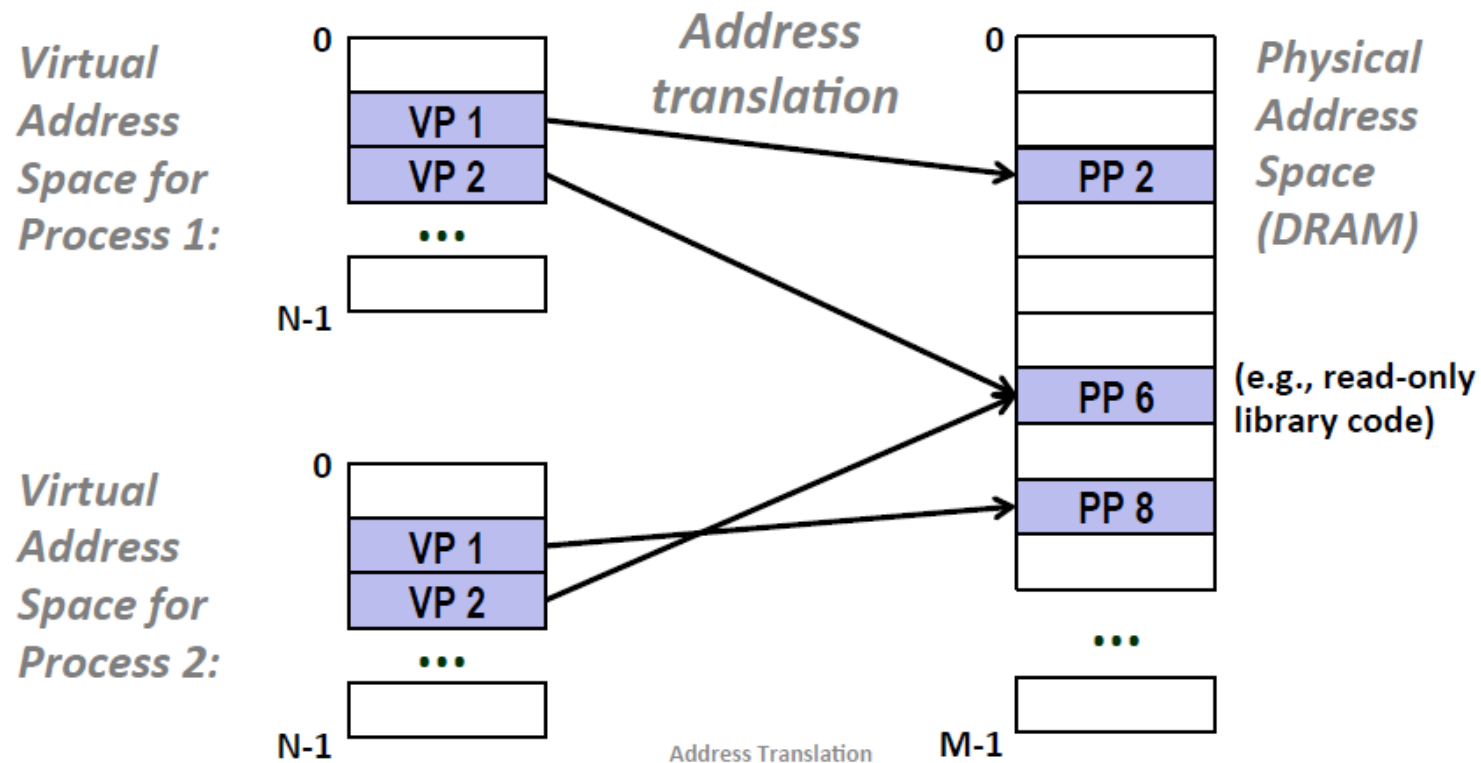


# Why does it work? Locality

- **Virtual memory works well because of locality**
  - **Spatial locality:** once an address is accessed, it is very likely that its neighbors will be accessed in the near future
  - **Temporal locality:** once an address is accessed, it is likely to be accessed again in the near future
- **The set of virtual pages that a program is “actively” accessing at any point in time is called its **working set****
- **If (working set size < main memory size)**
  - Good performance for the process after initial compulsory misses
- **If (working set size > main memory size)**
  - **Thrashing:** Performance meltdown where pages are swapped (copied) in and out continuously

# VM for Managing Multiple Processes

- **Key abstraction: Each process has its own virtual address space**
  - It can view memory as a simple linear array
- **With virtual memory, this simple linear virtual address space need not be contiguous in physical memory**

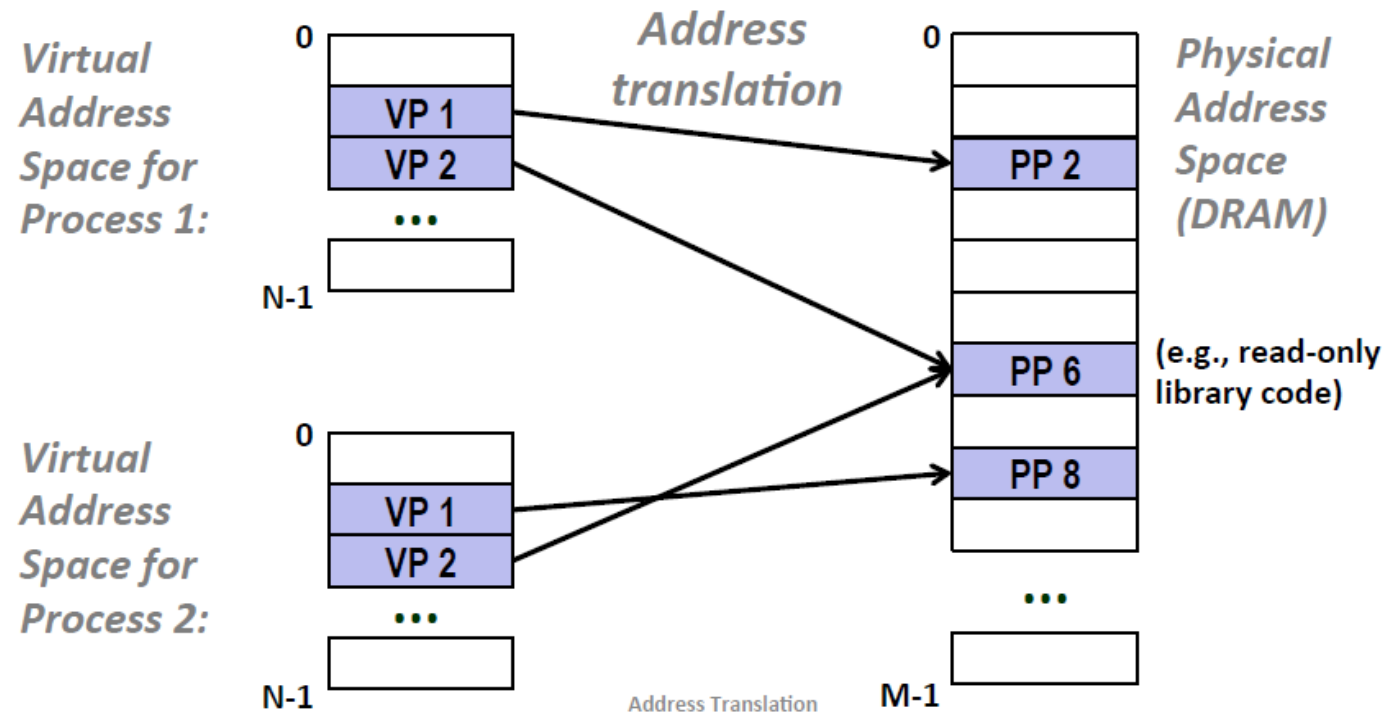


# Virtual Memory Advantages

- Virtual memory solves some of the issues with memory allocation in physical memory
- Virtual memory uses fixed page-size allocation: Internal Fragmentation if the process is not using the entire page
- Virtual memory does not require contiguous allocation in physical memory for a process: No external fragmentation

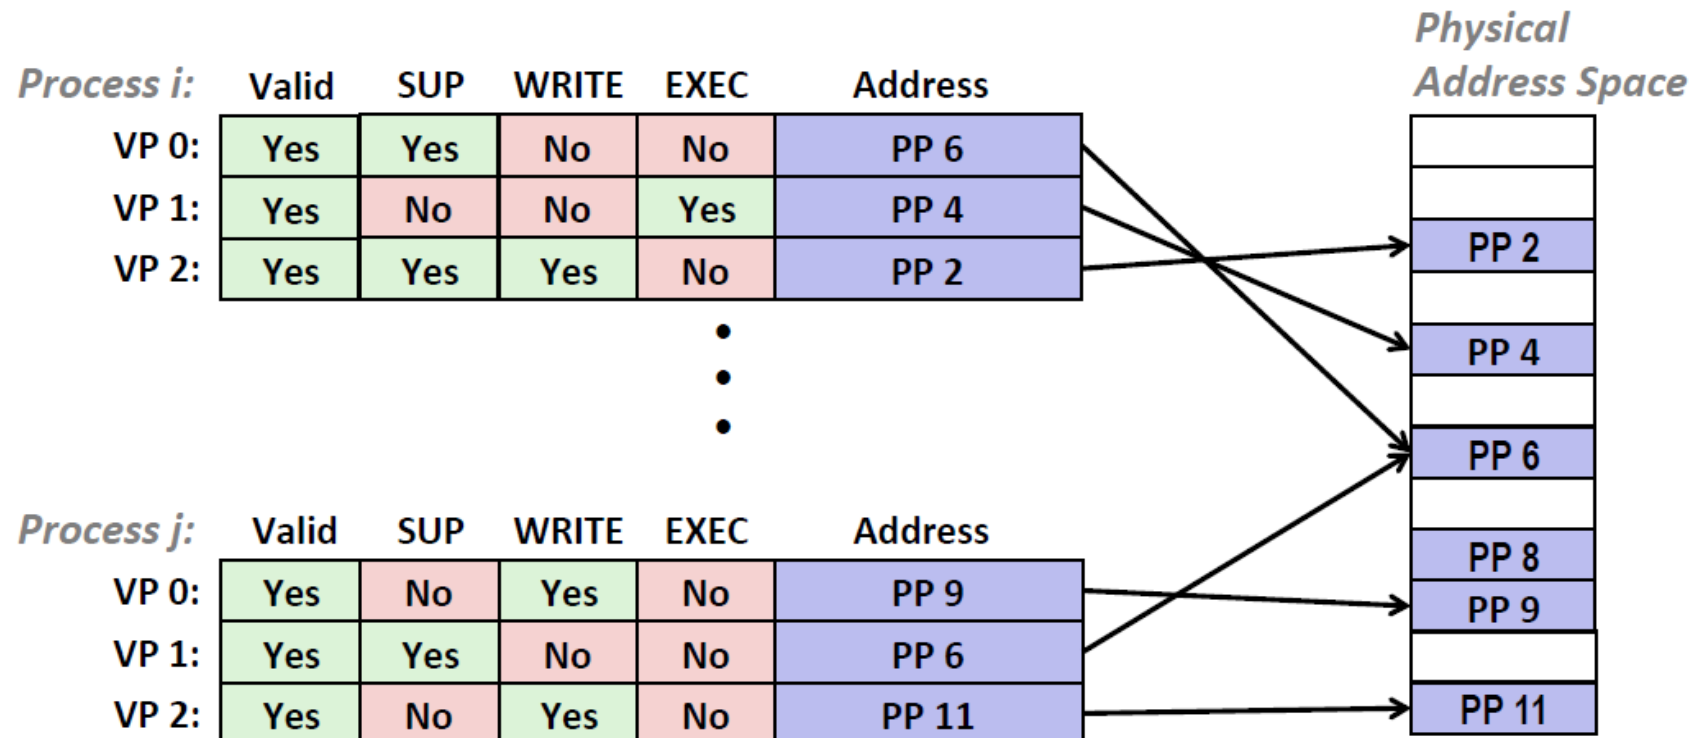
# VM for Protection and Sharing

- **Mapping of VPs to PPs provides a simple mechanism for protecting memory and for sharing memory between processes**
  - Sharing: Map virtual pages in separate address spaces to the same physical page (here PP6)
  - Protection: a process cannot access physical pages it does not have mapping for (here Process 2 cannot access PP2)



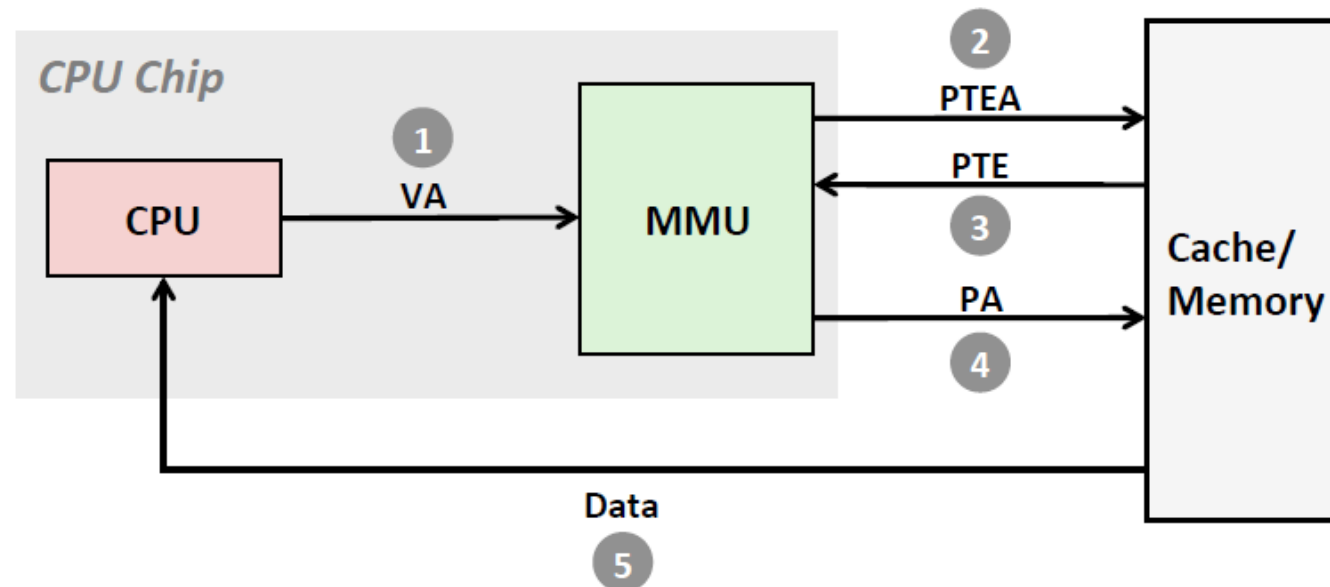
# Memory Protection within a Single Process

- **Extend page table entries with permission bits**
- **MMU checks these permission bits on every memory access**
  - If violated, raises an exception and OS sends segmentation fault signal to the process



# Address Translation: Page Hit

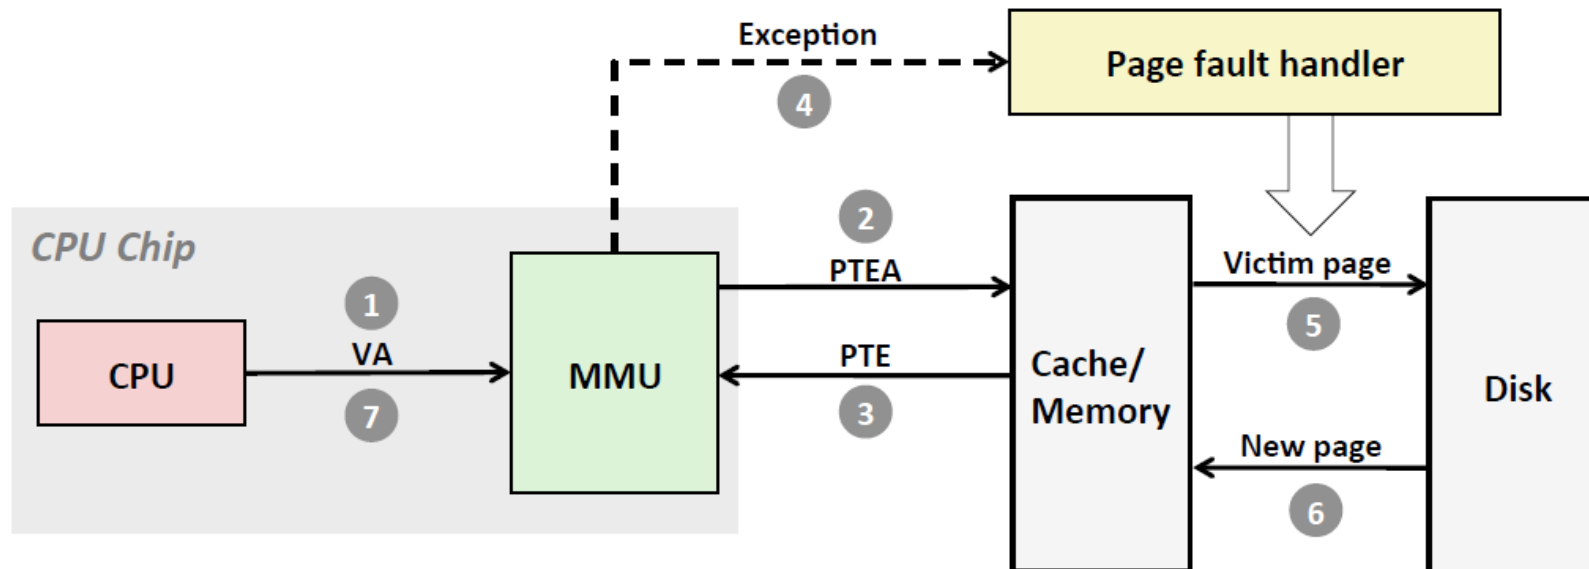
- (1) Processor sends virtual address to **MMU** (memory management unit)
- (2-3) **MMU** fetches page table entry (**PTE**) by accessing the page table stored in memory
- (4) **MMU** sends physical address to memory
- (5) **Memory** sends data word to processor





# Address Translation: Page Fault

- (1) Processor sends virtual address to **MMU**
- (2-3) **MMU** fetches **PTE** from page table
- (4) valid bit is zero, so **MMU** triggers page fault exception
- (5) Handler identifies victim (and if dirty, pages it out to disk)
- (6) Handler pages in new page and updates **PTE** in memory
- (7) Handler returns to original process, restarting faulty instr.



# Hmm...Translation Sounds Slow!!

- **MMU accesses memory *twice*: once to first get the PTE for translation and then again for the actual memory request from the CPU**
  - The PTE *may* be cached like any other memory word
  - But they may be evicted by other data references
  - And hit in the cache still requires 1-3 cycles
- **What can we do to make translation faster?**

# Speeding up Translation with TLB

- **Solution: Add another cache**
- **Translation Lookaside Buffer (TLB)**
  - Small hardware cache inside MMU
  - Maps virtual page numbers to physical page numbers
  - Contains page table entries for small number of virtual pages
  - Modern Intel processors: 128 or 256 entries in TLB
  - Works because of locality

# TLB Example

- TLB contains only a subset of page table entries
- Both Virtual Page number and Physical Page number are maintained
- Search all the TLB entries in parallel (associative search)

	PP	Valid
0	0x10	1
1	DISK	0
2	0x23	1
3	0x15	1
4	DISK	0
5	0x1B	1
6	0x2E	1
...	...	...
255	0x2A	1

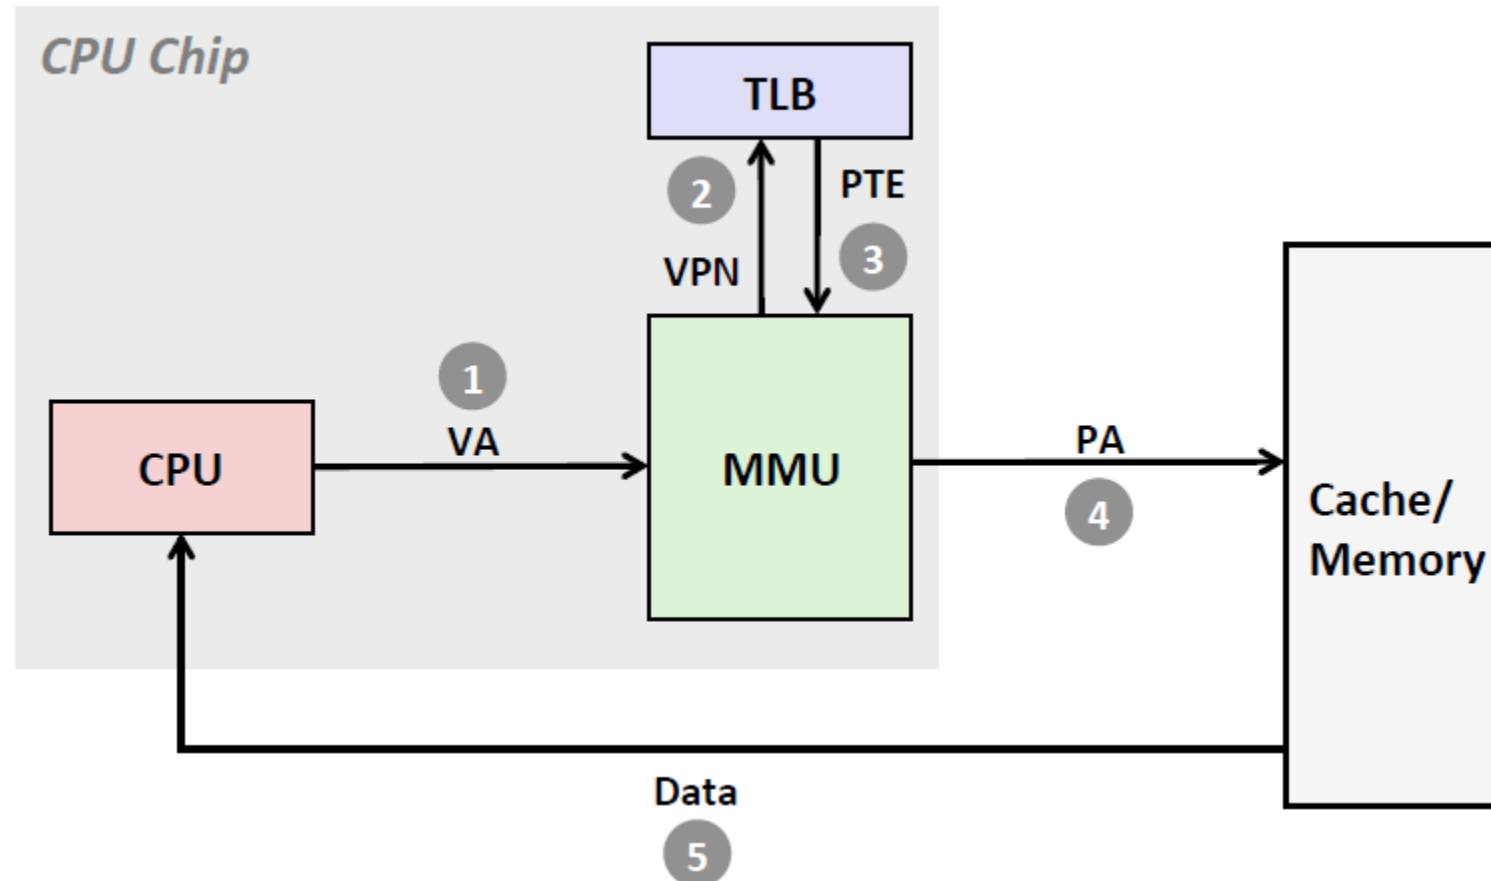
Full Page Table in DRAM

VP	PP	Valid
0x06	0x2E	1
-	-	0
0x02	0x23	1
-	-	0

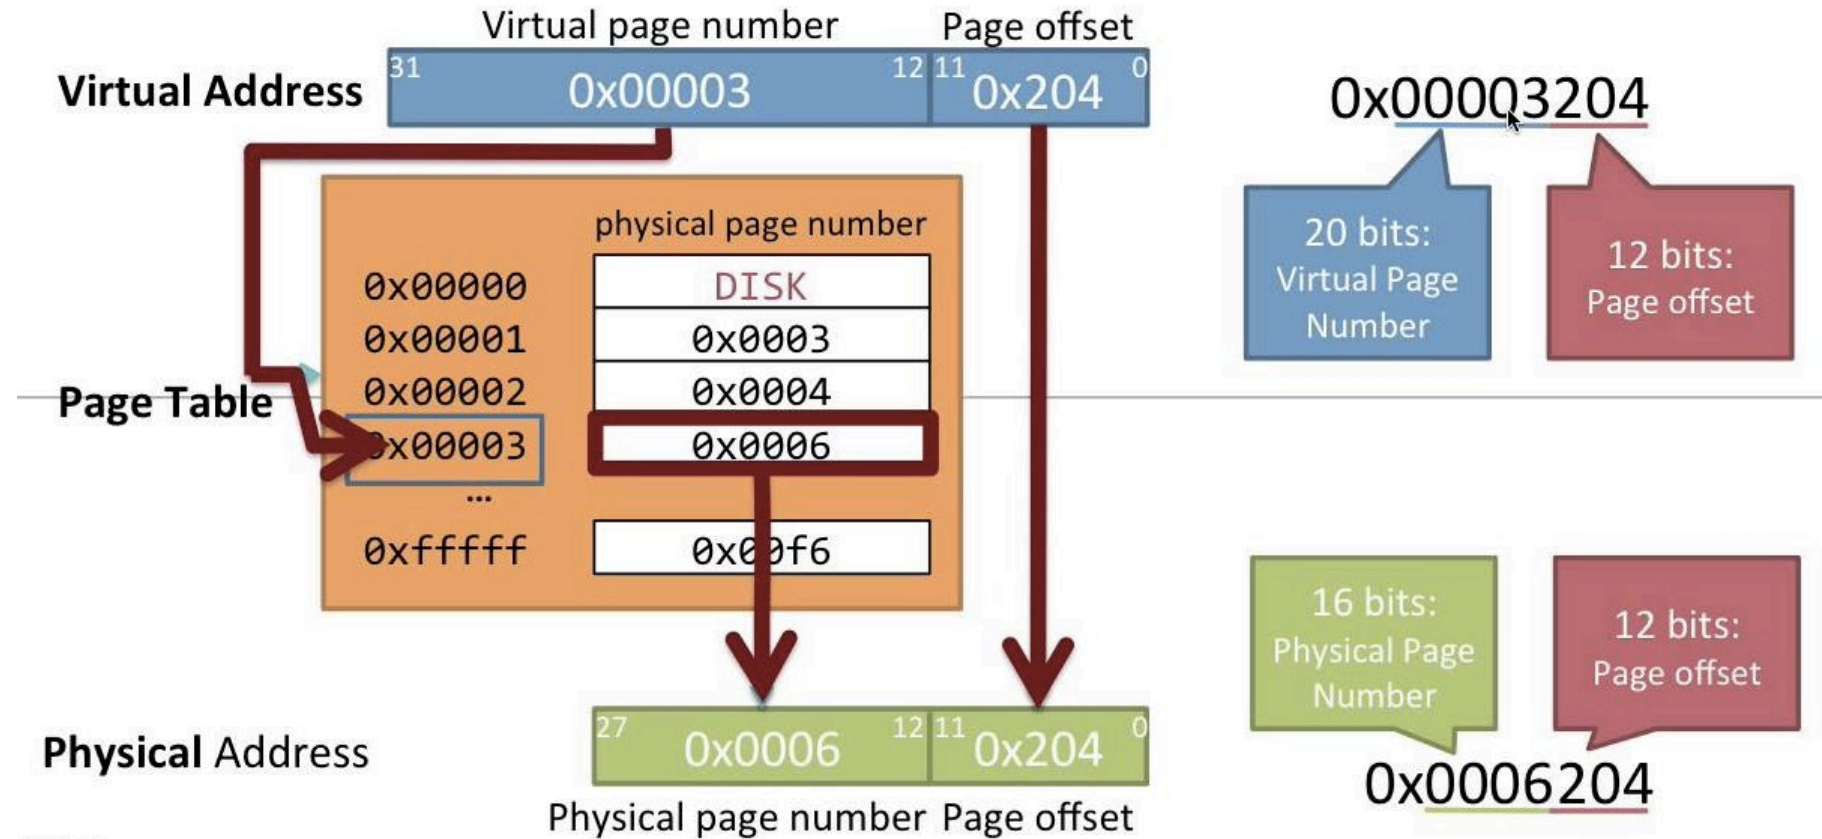
On-chip TLB

# TLB Hit

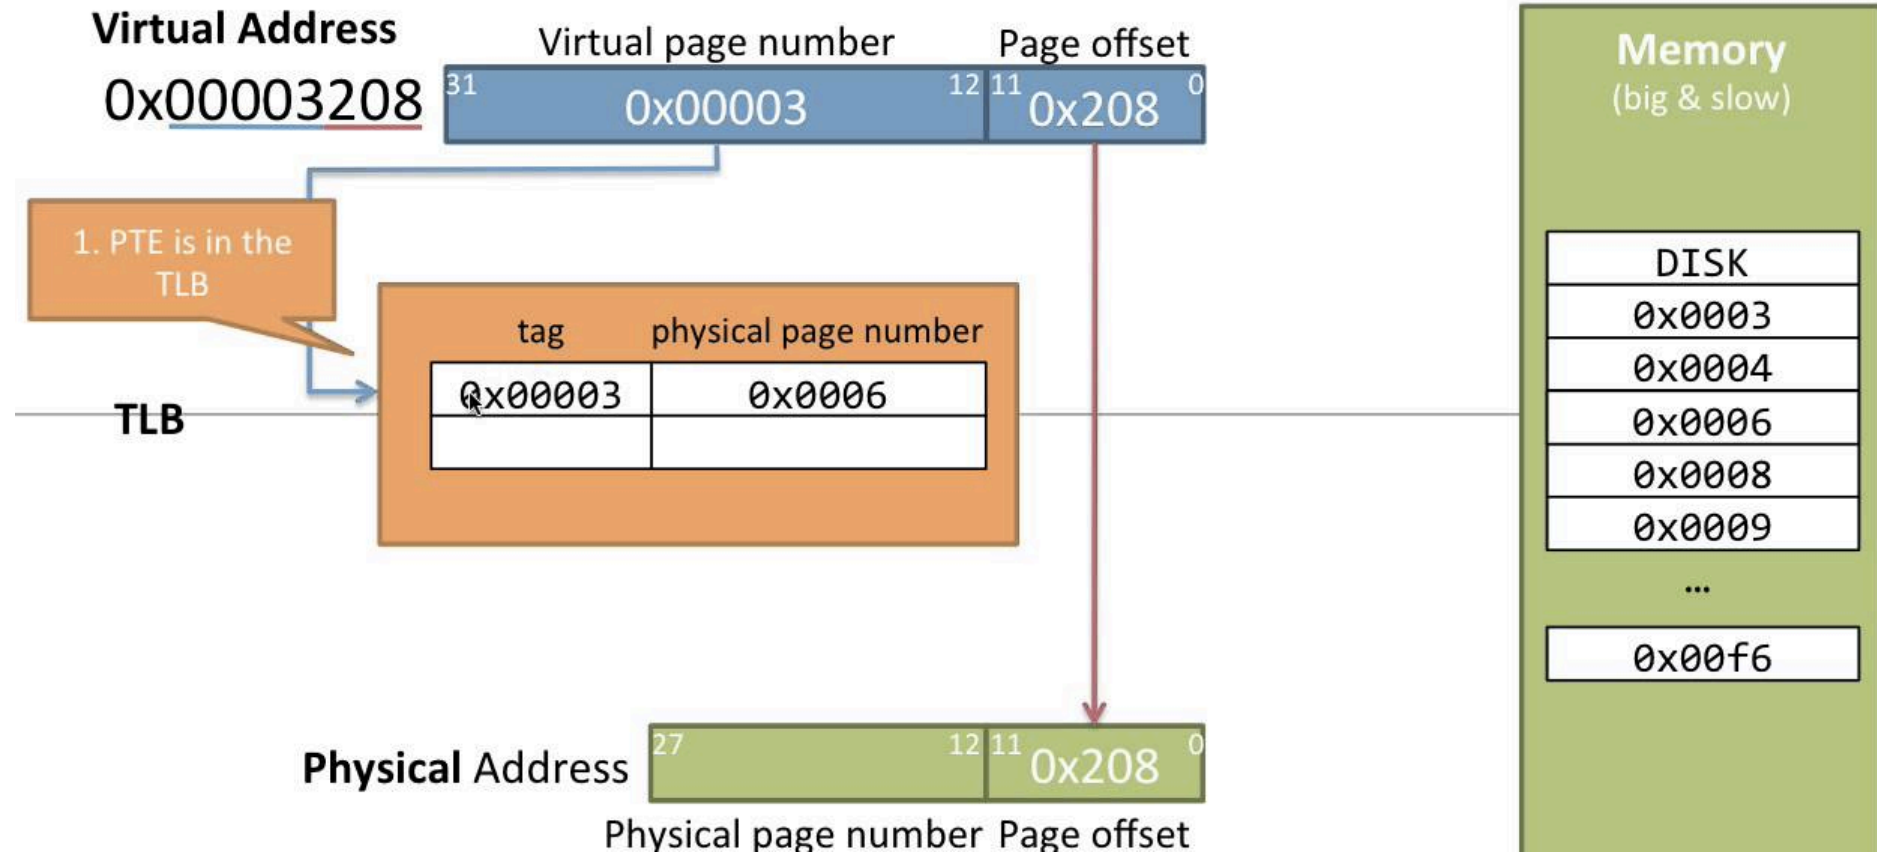
- A TLB hit eliminates costly memory access for page table entry



# Translation with Page Table in DRAM

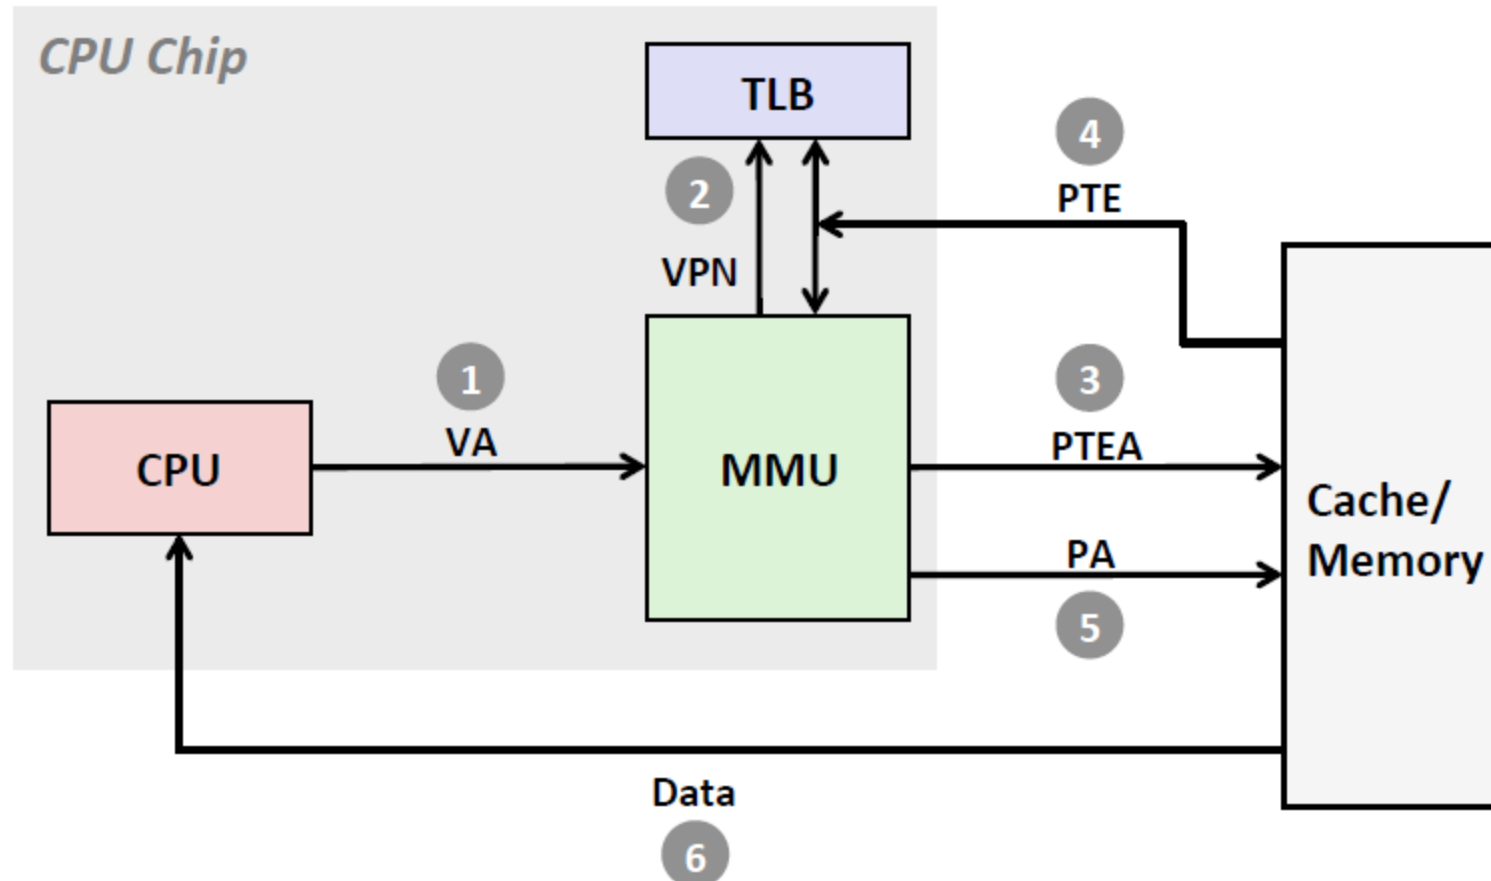


# Translation with TLB Hit



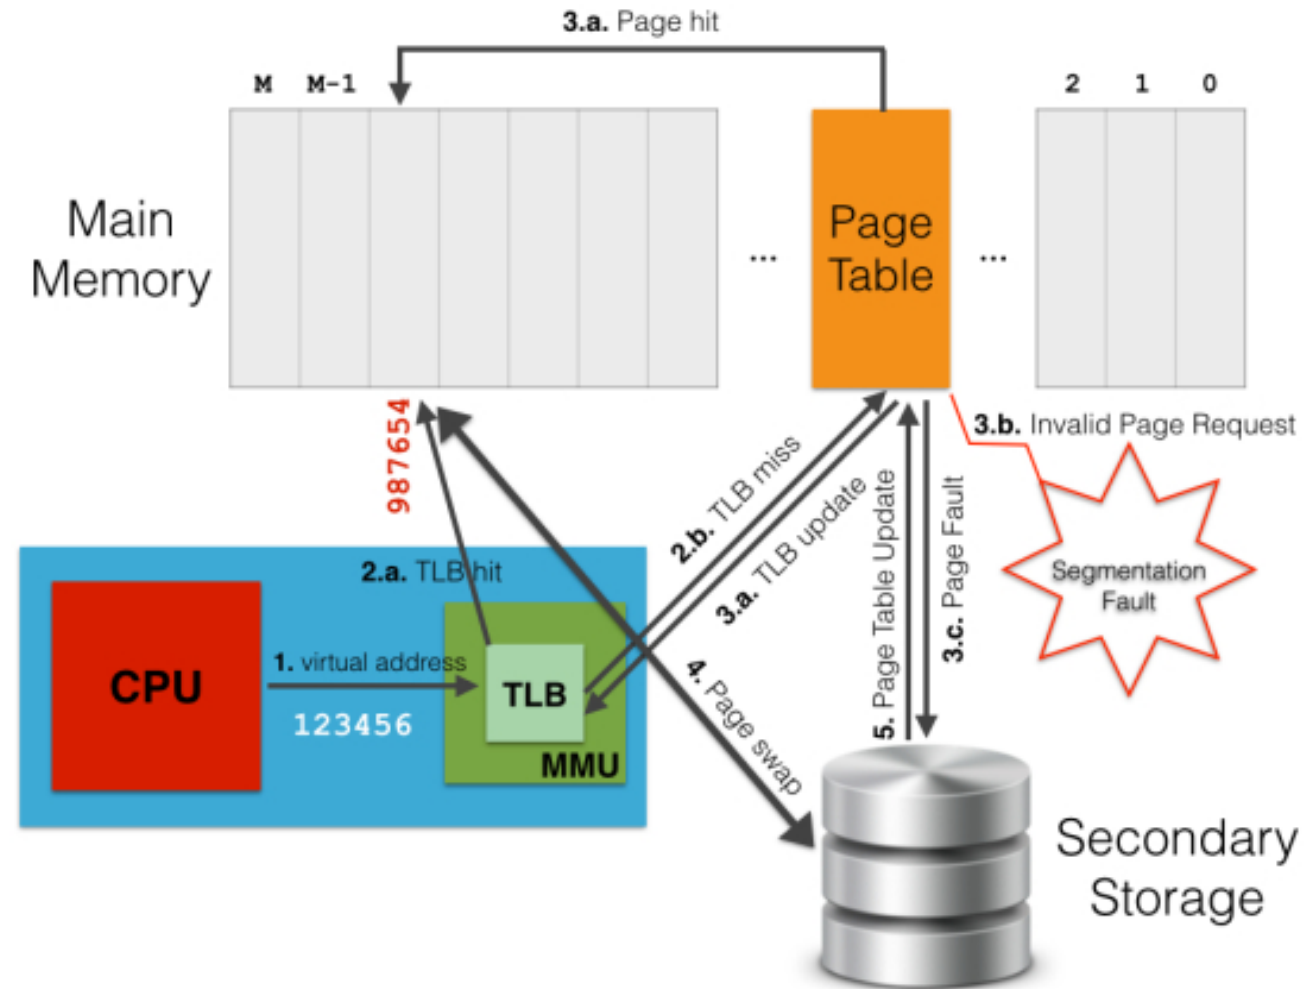
# TLB Miss

- **A TLB miss incurs additional memory access for the PTE**
  - Fortunately TLB misses are rare





# Putting it all together

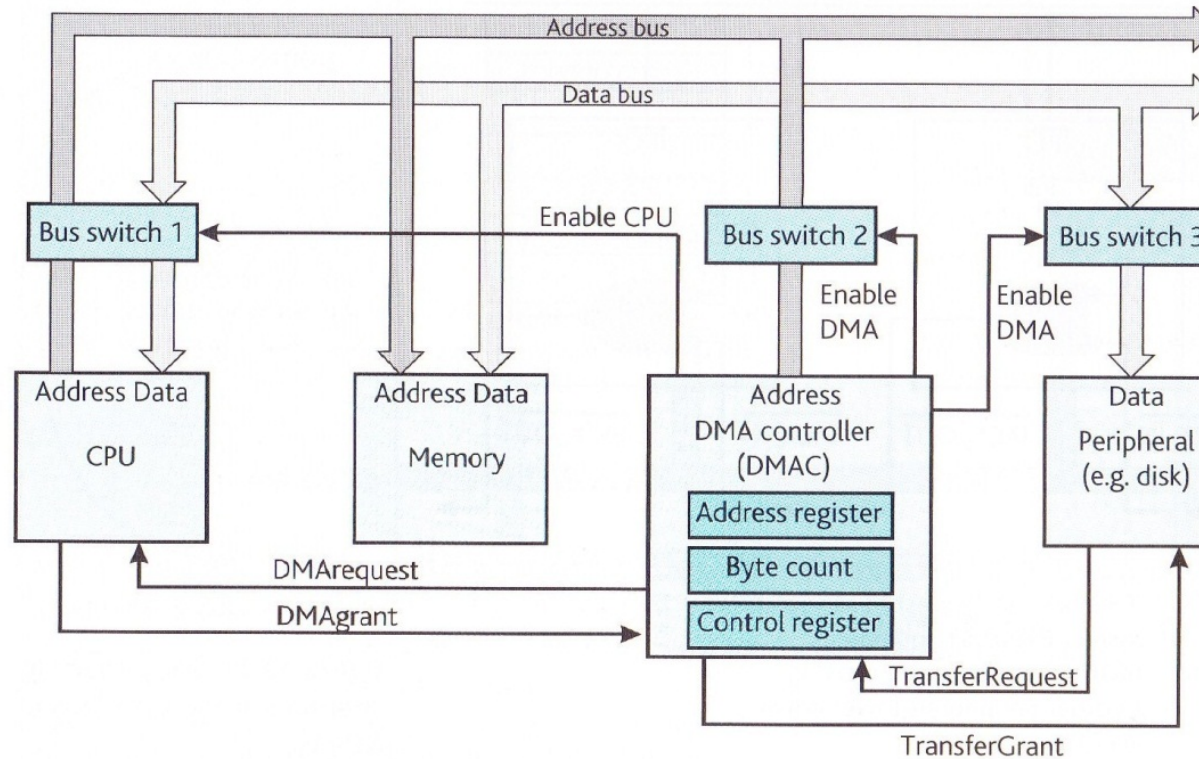


# Direct Memory Access

- Direct Memory Access (DMA) moves data between a peripheral and the CPU's memory without the direct intervention of the CPU itself.
- Provides the fastest possible means of transferring data between an interface and memory.
- It requires no CPU overhead and leaves the CPU free to do useful work.
- Helps to increase data throughput.

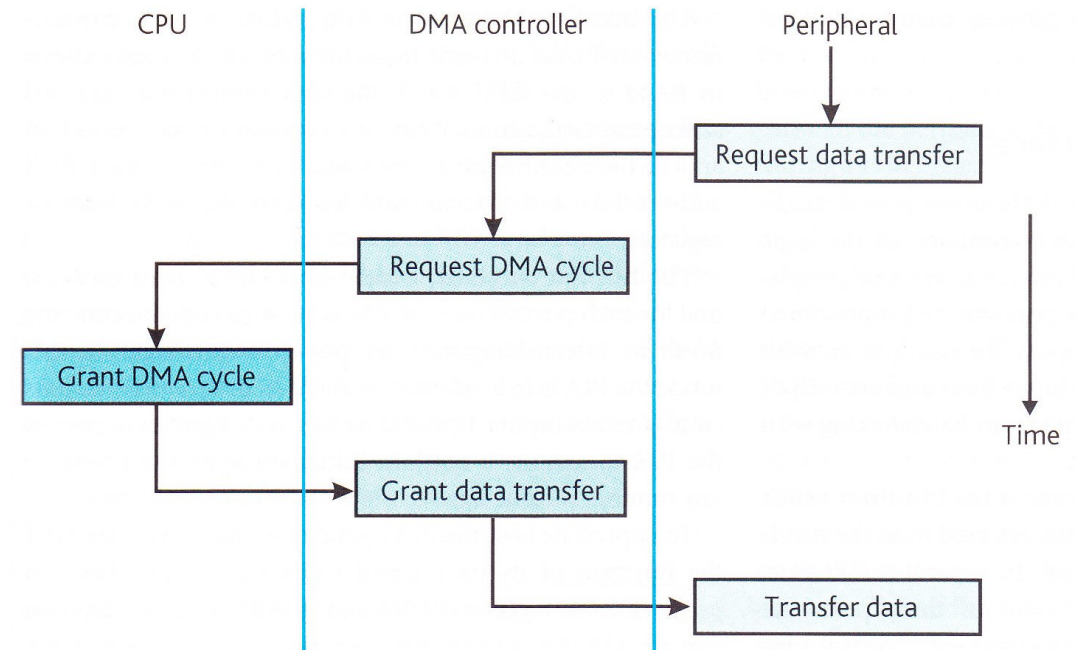
# Direct Memory Access

- During normal operation of the computer, bus switch 1 is closed and bus switches 2 and 3 are open.
- The CPU controls the buses, providing an address on the address bus and reading and writing data from memory via the data bus.



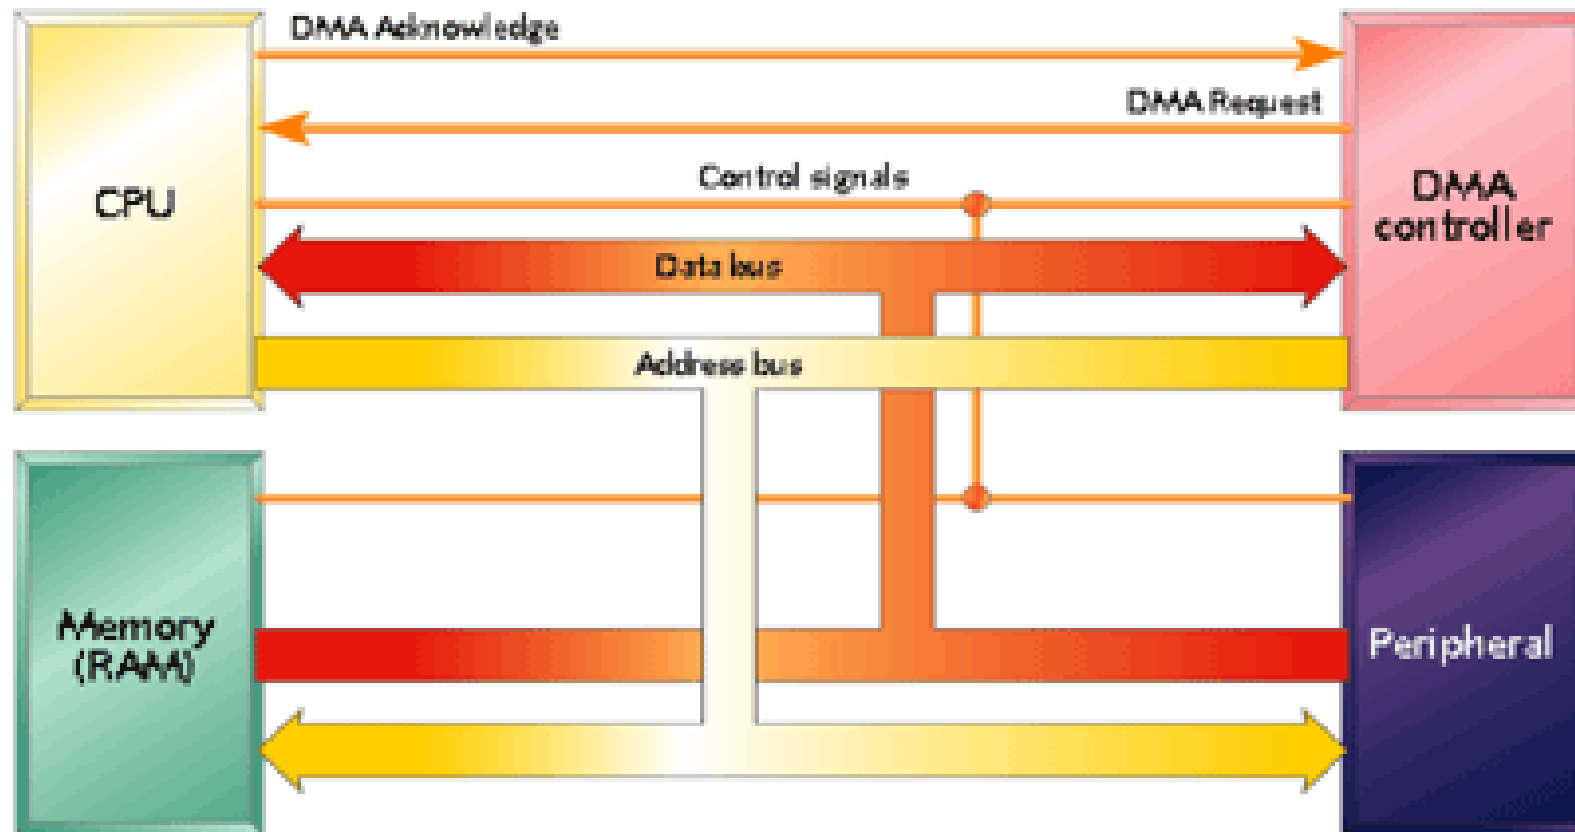
# Direct Memory Access

- When a peripheral wishes to take part in an I/O transaction it asserts the TransferRequest input of the DMA controller (DMAC).
- In turn, the DMAC controller asserts the DMArequest (BusReq) to request control of the buses from the CPU.
- When the CPU returns DMAgrant (BusAck) to the DMAC, a DMA transfer takes place.
- Bus switch 1 is open and switches 2 and 3 are closed.
- The DMAC provides TransferGrant signal to the peripheral.
- The peripheral is able to communicate with the memory using the DMA
- When DMA operation is completed, DMAC hands control of the bus back to the CPU.



# Direct Memory Access

- It is not necessary for the DMA to be initiated by the Peripheral.
- It can be initiated within the Processor itself.




# Direct Memory Access

- Many DMACs are able to handle several interfaces, which requires their registers to be duplicated.
- Each interface is referred to as a channel.
- DMA operates in one of two modes: burst mode or cycle stealing
- Burst Mode
  - DMA controller seizes the system bus for the duration of the data transfer operation.
  - Allows data to be moved at the fastest possible rate within the memory/bus/interface timing requirements
  - CPU is effectively halted in the burst mode because it cannot use its buses.

# Direct Memory Access

- Cycle Steal Mode
  - DMA operations are interleaved with the processor's normal memory accesses.
  - DMA can take place when CPU is busy.
  - Transparent because transfer is invisible to the processor and no processing time is lost.



# Summary

- **Programmer's view of virtual memory**
  - Each process has its own private linear address space
  - Cannot be corrupted by other processes
  
- **System view of virtual memory**
  - Uses memory efficiently by caching virtual memory pages
  - Efficient only because of locality
  - Simplifies memory management and sharing
  - Simplifies protection by providing a convenient intermediate point to check permission
  
- **DMA**
  - A means to free up the CPU to allow memory access to be handled by the HW



# The End!

- We have completed the module! Hooray! 😊
- You have obtained extensive experience and knowledge in both Embedded Systems Development together with key RTOS concepts.
- If you are keen to be a TA for this module next semester, please apply and let me know!
- We appreciate your feedback, which helps us improve
- We look forward to not just feedback for this module, but also feedback for the CEG programme.
- Thank You and Good Luck for the Exams! 😊